

## Module 4F10: DEEP LEARNING AND STRUCTURED DATA

**Examples Paper 2**

*Straightforward questions are marked †*

*Tripos standard (but not necessarily Tripos length) questions are marked \**

*Deep Learning*

1. Residual networks are often used with very deep networks. For a single layer, layer  $l$  with  $N_l$  nodes, of a residual network the activation function can be expressed as (note no bias is used here)

$$\mathbf{y}^{(l)} = \mathbf{x}^{(l)} + \phi(\mathbf{W}\mathbf{x}^{(l)}); \quad \phi(\mathbf{z}) = \begin{bmatrix} 1/(1 + \exp(-z_1)) \\ \vdots \\ 1/(1 + \exp(-z_{N_l})) \end{bmatrix}$$

where  $\mathbf{x}^{(l)}$  and  $\mathbf{y}^{(l)}$  are the input and output to layer  $l$ ,  $\phi()$  is a sigmoid activation function that operates on each element of the vector. Thus the standard activation function,  $\phi()$ , has a residual connection added.

- (a) For this network configuration derive an expression for the derivative  $\frac{\partial \mathbf{y}^{(l)}}{\partial \mathbf{x}^{(l)}}$ .
  - (b) Based on the answer in (a), comment on how the use of a residual connection can help in reducing the vanishing gradient problem.
  - (c) Give the relationship between the input and output for this layer when the residual connection is replaced by a highway connection. Do you expect this form of connection to also help address the vanishing gradient problem?
2. A simplified recurrent neural network is to be trained. For this simplified network a linear activation function is used such that the recurrence relationship has the form

$$\mathbf{h}_t = \mathbf{W}\mathbf{h}_{t-1}$$

Show that the activation function output at time  $t$  can be expressed as

$$\mathbf{h}_t \approx \lambda^t \mathbf{q}\mathbf{v}^\top \mathbf{h}_0$$

when  $t$  gets large. Clearly define scalar  $\lambda$ , and vectors  $\mathbf{q}$  and  $\mathbf{v}$ . What is the implication of this expression for training recurrent neural networks with long sequences?

3. Figure 3 shows a Gated Recurrent Unit (GRU) that is to be used for a  $K$  class classification problem. The output of the GRU at time  $t$ ,  $\mathbf{h}_t$ , is used to predict a probability mass function (PMF) over the  $K$  classes at time  $t$ ,  $\mathbf{y}_t$ , as well as the input to the next time instance.

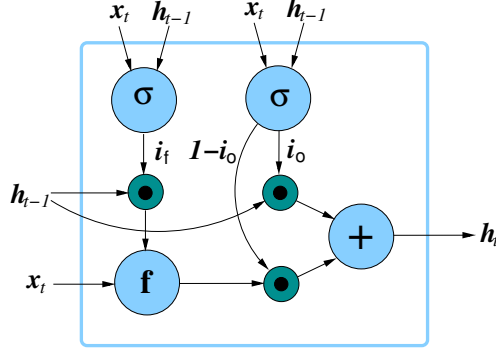


Figure 1: Gated Recurrent Unit

- What is an appropriate form for the gating activation function  $\sigma(\mathbf{x}_t, \mathbf{h}_{t-1})$ ? You should justify your answer.
  - Give the overall expression for how the inputs,  $\mathbf{x}_t$  and  $\mathbf{h}_{t-1}$ , are transformed to yield the output  $\mathbf{y}_t$  and new history vector  $\mathbf{h}_t$ . You should clearly state the form of activation functions being used.
  - An additional GRU is added to the network. Briefly describe the revised network configuration.
4. Layer normalisation (layer norm) is a popular form of normalisation for deep learning. This acts on the the activation function outputs for particular layer. For layer  $k$  and input vector  $\mathbf{x}_p$  the  $n_k$  nodes associated with the layer yields the vector of activation function outputs  $\mathbf{y}_p^{(k)}$ . Layer norm is applied to yield the normalised vector  $\tilde{\mathbf{y}}_p^{(k)}$ , which is passed to the next layer, where

$$\tilde{\mathbf{y}}_p^{(k)} = \frac{1}{\tilde{\sigma}_p^{(k)}}(\mathbf{y}_p^{(k)} - \tilde{\mu}_p^{(k)}\mathbf{1}); \quad \tilde{\mu}_p^{(k)} = \frac{1}{n_k} \sum_{j=1}^{n_k} y_{pj}^{(k)}; \quad \tilde{\sigma}_p^{(k)2} = \frac{1}{n_k} \sum_{j=1}^{n_k} (y_{pj}^{(k)} - \tilde{\mu}_p^{(k)})^2$$

and  $\mathbf{1}$  is the vector of ones length  $n_k$ .

- Compare layer and batch normalisation for training networks. You should consider the form of normalisation applied and any difficulties/sensitivities in applying the transform during training or inference
- How does layer norm alter the derivative  $\partial E(\boldsymbol{\theta})/\partial \mathbf{y}^{(k)}$  when  $\tilde{\sigma}_p^{(k)}$  is ignored in the normalisation (so  $\tilde{\sigma}_p^{(k)} = 1$  for all inputs)?

5. (optional) The following form of expression is often used in variational optimisation with training examples  $\mathbf{x}_1, \dots, \mathbf{x}_n$  to find the model parameters  $\boldsymbol{\lambda}$

$$\mathcal{L}(\boldsymbol{\lambda}) = \sum_{i=1}^n \log(\mathbf{x}_i; \boldsymbol{\lambda}) \geq \sum_{i=1}^n \int \log \left( \frac{p(\mathbf{x}_i, \mathbf{z}; \boldsymbol{\lambda})}{q_1(\mathbf{z}; \tilde{\boldsymbol{\lambda}}_1)} \right) q_2(\mathbf{z}; \tilde{\boldsymbol{\lambda}}_2) d\mathbf{z} = \sum_{i=1}^n \left\langle \log \left( \frac{p(\mathbf{x}_i, \mathbf{z}; \boldsymbol{\lambda})}{q_1(\mathbf{z}; \tilde{\boldsymbol{\lambda}}_1)} \right) \right\rangle_{q_2(\mathbf{z}; \tilde{\boldsymbol{\lambda}}_2)}$$

- (a) Show that if  $q_1(\mathbf{z}; \tilde{\boldsymbol{\lambda}}_1) = p(\mathbf{z}|\mathbf{x}; \boldsymbol{\lambda})$  for all values of  $\mathbf{x}$  and  $\mathbf{z}$  then the left and right-hand sides of this expression are equal for all functions  $q_2(\mathbf{z}; \tilde{\boldsymbol{\lambda}}_2)$ .
- (b) Show that if  $q_1(\mathbf{z}; \tilde{\boldsymbol{\lambda}}_1) = q_2(\mathbf{z}; \tilde{\boldsymbol{\lambda}}_2)$  then equality only occurs when  $q_1(\mathbf{z}; \tilde{\boldsymbol{\lambda}}_1) = p(\mathbf{z}|\mathbf{x}; \boldsymbol{\lambda})$  for all values of  $\mathbf{x}$  and  $\mathbf{z}$ .
- (c) Briefly discuss why this form of expression may be useful.

### *Support Vector Machines*

- 6. † A binary classifier is to be trained. What are the limitations of linear decision classifiers and why do non-linear mappings of the feature space allow improved discrimination? Under what conditions is it guaranteed that a non-linear mapping will allow perfect classification of the data?
- 7. In the XOR classification from the lectures, if we use the kernel  $k(\mathbf{x}_n, \mathbf{x}_m) = (1 + \mathbf{x}_n^T \mathbf{x}_m)^2$ , what is the resulting Gram matrix? The solution to the dual problem with such kernel is  $a_1 = a_2 = a_3 = a_4 = 1/8$ . Show that this solution satisfies the constraints  $t_n y(\mathbf{x}_n) \geq 1$  for  $n = 1, \dots, 4$ , where  $y(\mathbf{x})$  is the output of the classifier for input  $\mathbf{x} \in \mathbb{R}^2$ . What is the equation of the final decision boundary?
- 8. The following data is to be used for training an SVM

$$\text{Positive class } (t_n = +1): \quad \mathbf{x}_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix}, \quad \mathbf{x}_2 = \begin{bmatrix} 2 \\ 2 \end{bmatrix}, \quad \mathbf{x}_3 = \begin{bmatrix} 2 \\ 0 \end{bmatrix}.$$

$$\text{Negative class } (t_n = -1): \quad \mathbf{x}_4 = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \quad \mathbf{x}_5 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \quad \mathbf{x}_6 = \begin{bmatrix} 0 \\ 1 \end{bmatrix}.$$

- (a) Plot the training points and, by inspection, draw the optimal, maximum margin, decision boundary.
- (b) What are the support vectors? Let  $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$  be the classifier's output for input  $\mathbf{x}$ . What are  $\mathbf{w}$  and  $b$  so that  $t_n y(\mathbf{x}_s) = 1$  for any support vector  $\mathbf{x}_s$ ?
- (c) Express  $y(\mathbf{x})$  in terms of the Lagrange multipliers,  $a_n$  and show that  $\mathbf{w}$ ,  $b$  and the  $a_n$  satisfy the KKT conditions.

9. \* A Support Vector Machine (SVM) is to be used for classifying sequences represented by 1-dimensional feature-vectors of variable length. To solve the classification problem, each input sequence  $\mathbf{x}_{1:T} = \{x_1, \dots, x_T\}$  is transformed using the feature mapping

$$\phi(\mathbf{x}_{1:T}) = \begin{bmatrix} \frac{\partial}{\partial \mu_1} \log p(\mathbf{x}_{1:T}) \\ \vdots \\ \frac{\partial}{\partial \mu_M} \log p(\mathbf{x}_{1:T}) \\ \frac{\partial^2}{\partial \mu_1^2} \log p(\mathbf{x}_{1:T}) \\ \vdots \\ \frac{\partial^2}{\partial \mu_1 \partial \mu_M} \log p(\mathbf{x}_{1:T}) \\ \vdots \\ \frac{\partial^2}{\partial \mu_M^2} \log p(\mathbf{x}_{1:T}) \end{bmatrix},$$

where  $p(\mathbf{x}_{1:T})$  is specified by a generative model given by an  $M$ -component Gaussian Mixture Model (GMM):

$$p(\mathbf{x}_{1:T}) = \prod_{t=1}^T \sum_{m=1}^M c_m \mathcal{N}(x_t; \mu_m, \sigma_m^2)$$

- (a) Why is this form of feature-space suitable for use with SVMs when classifying variable-length data-sequences? Why is an SVM a suitable form of classifier as  $M$  (the number of components) gets large? What is the dimensionality of the feature-space in this case?
- (b) Derive an expression for  $\frac{\partial}{\partial \mu_i} \log p(\mathbf{x}_{1:T})$ . This should be expressed in terms of  $P(i|x_t)$ , the posterior probability that the  $i$ -th Gaussian component generated the observation.
- (c) Hence show that

$$\frac{\partial^2}{\partial \mu_j \partial \mu_i} \log p(\mathbf{x}_{1:T}) = - \sum_{t=1}^T P(i|x_t) P(j|x_t) \frac{(x_t - \mu_j)(x_t - \mu_i)}{\sigma_i^2 \sigma_j^2}.$$

Do you expect these second-order derivative terms to help in classification?