

# Image Structure

**Feature Detection and Matching**

**4F12: Computer Vision**

**Combined slides from lecture and Q & A**

**Instructor: Samuel Albanie**

*Based on course material authored by Roberto Cipolla*



Tag = Very Examinable

Tag = Non Examinable

# Image Structure 1: Q&A

## Feature Detection and Matching

### 4F12: Computer Vision

We will start at **11:05 am** (to allow everyone to join from previous in-person lectures/sports fixtures/poetry classes etc.)

If you are reading this before 11:05 am, you may wish to apply the computer vision researcher algorithm:

**If 5 mins spare:** start coding up latest idea from the "my amazing research ideas" list and make hot drink. There's still time!

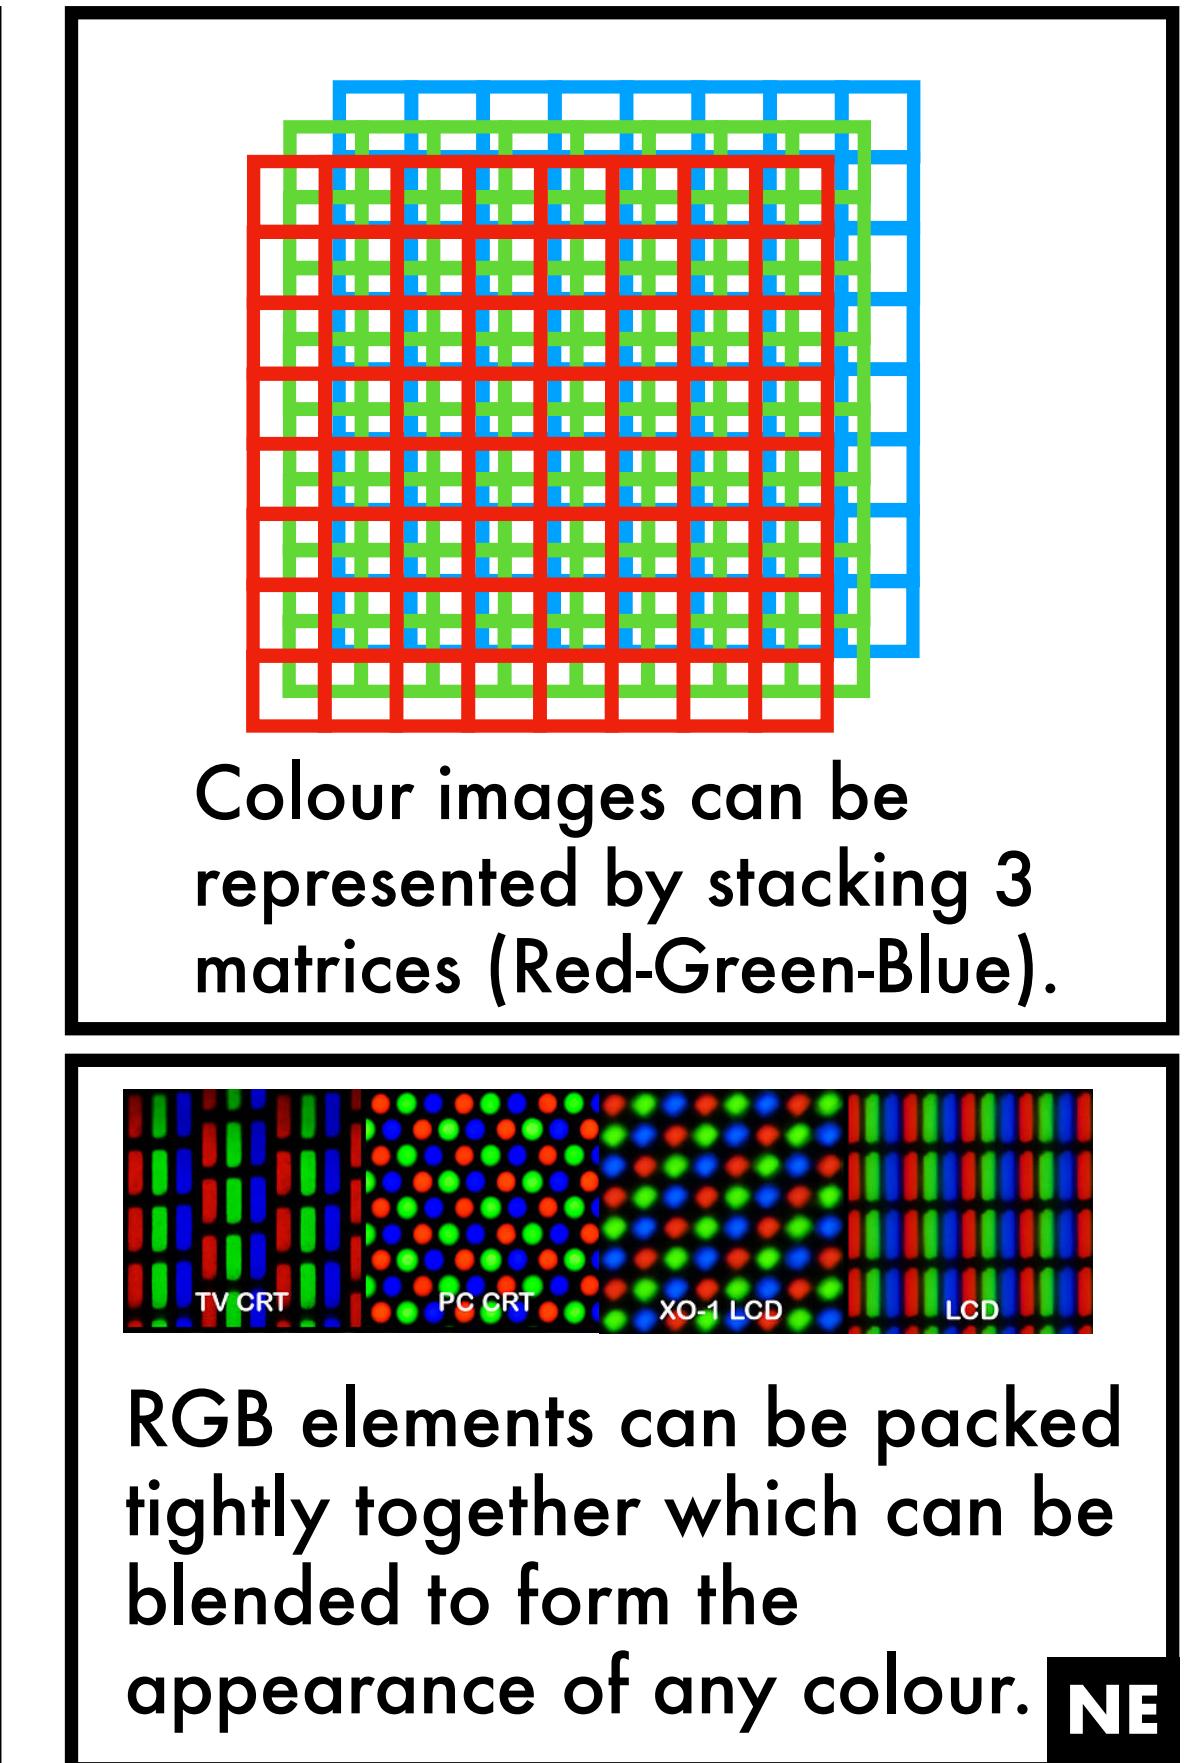
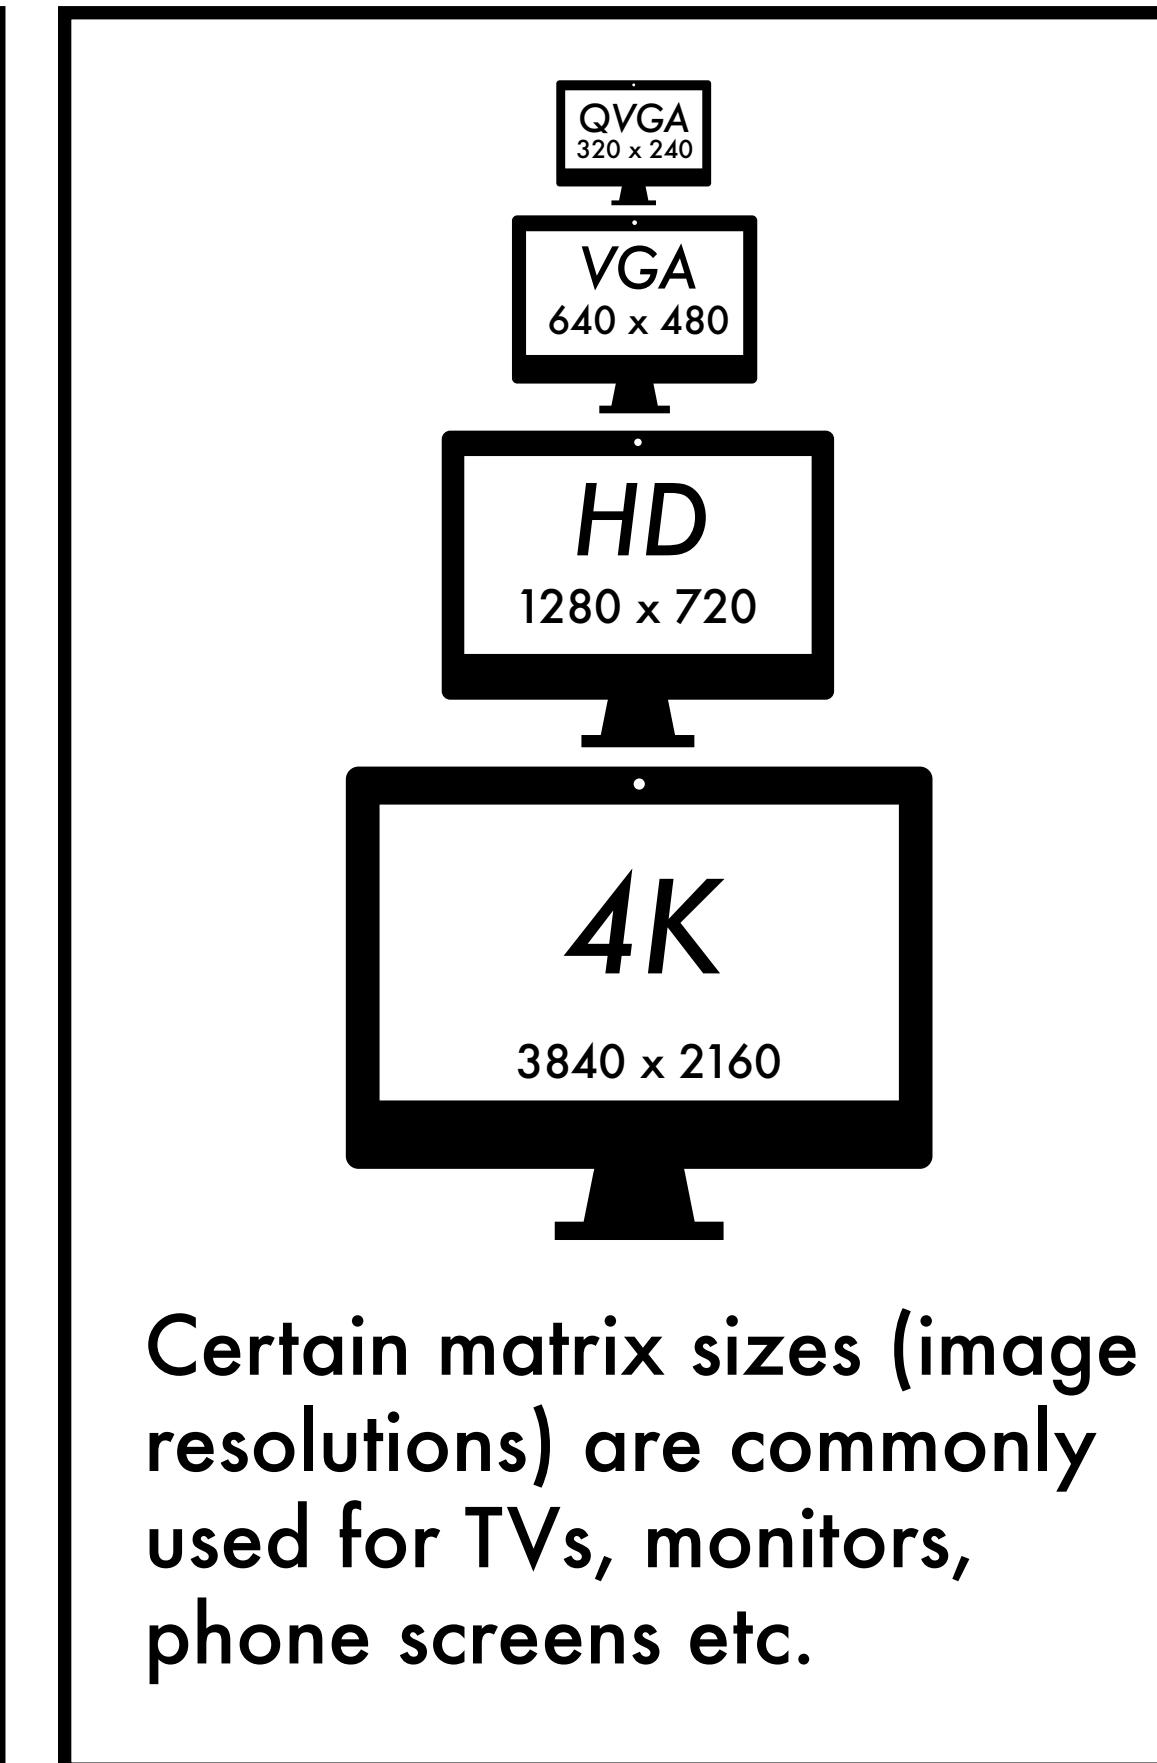
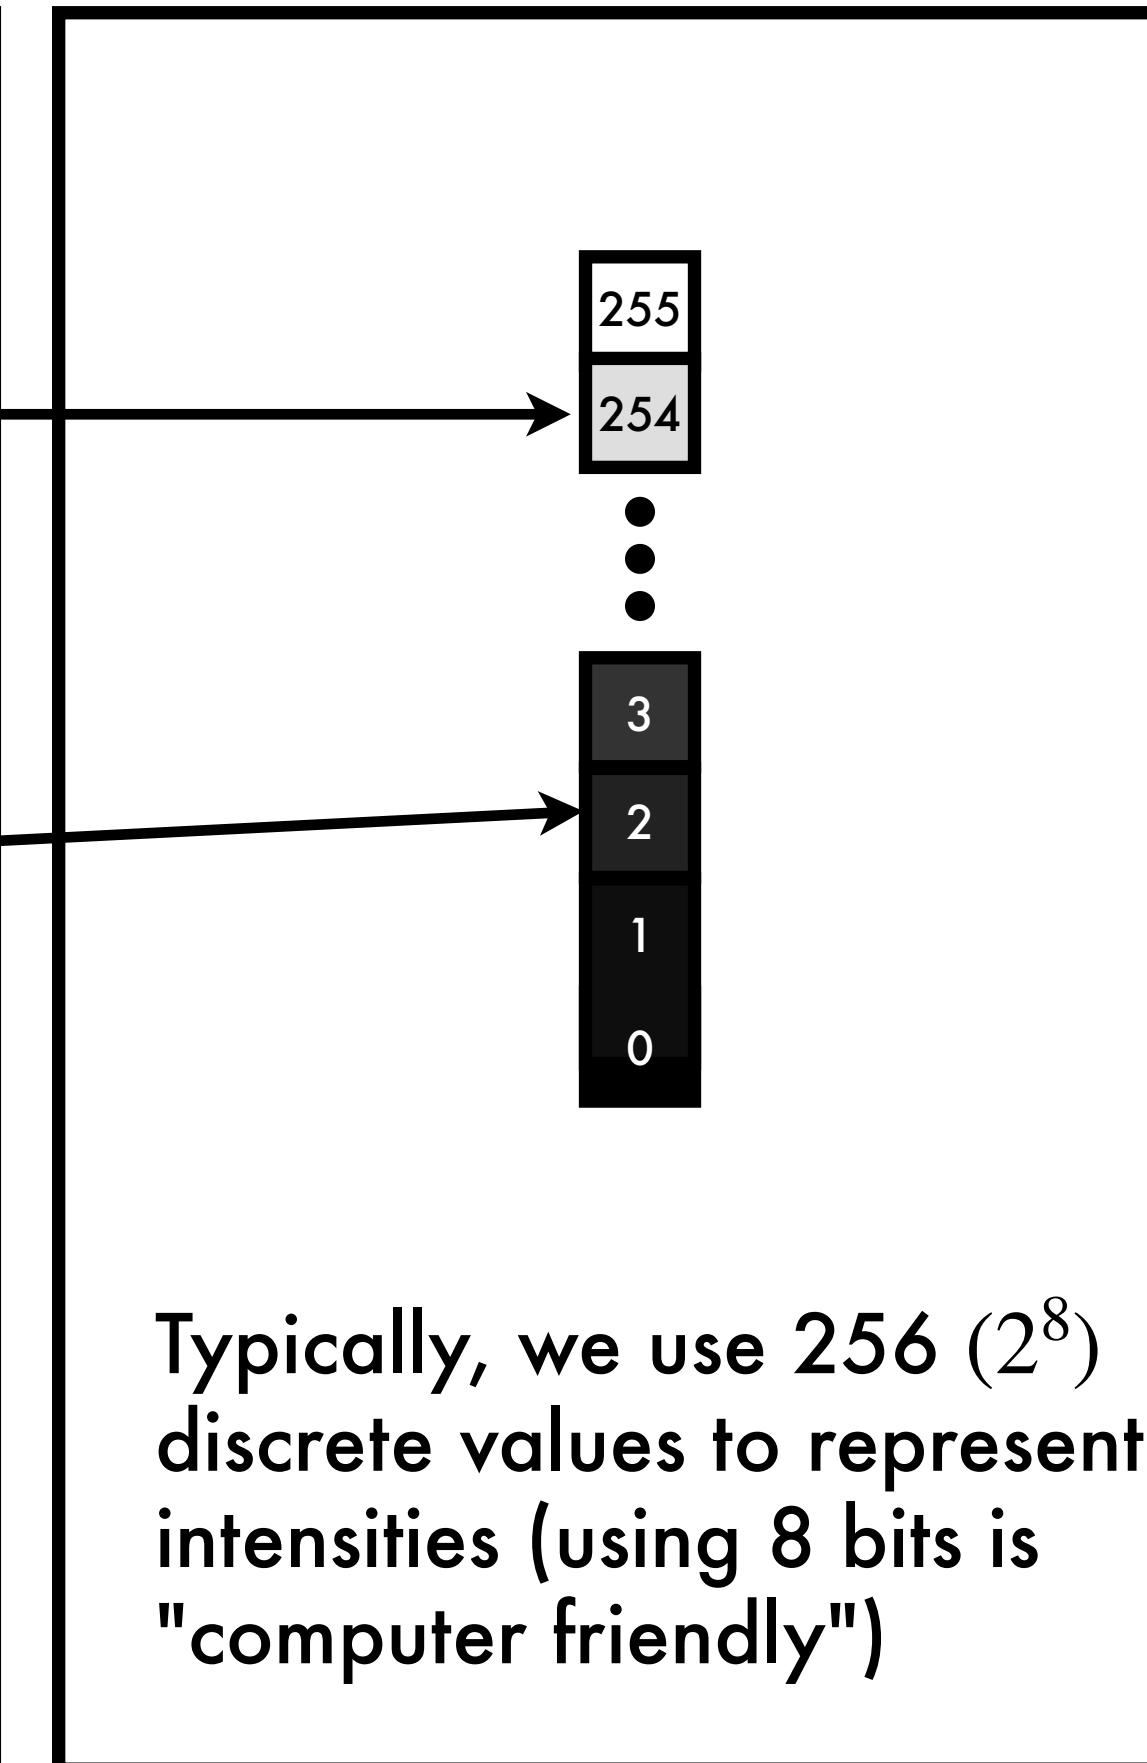
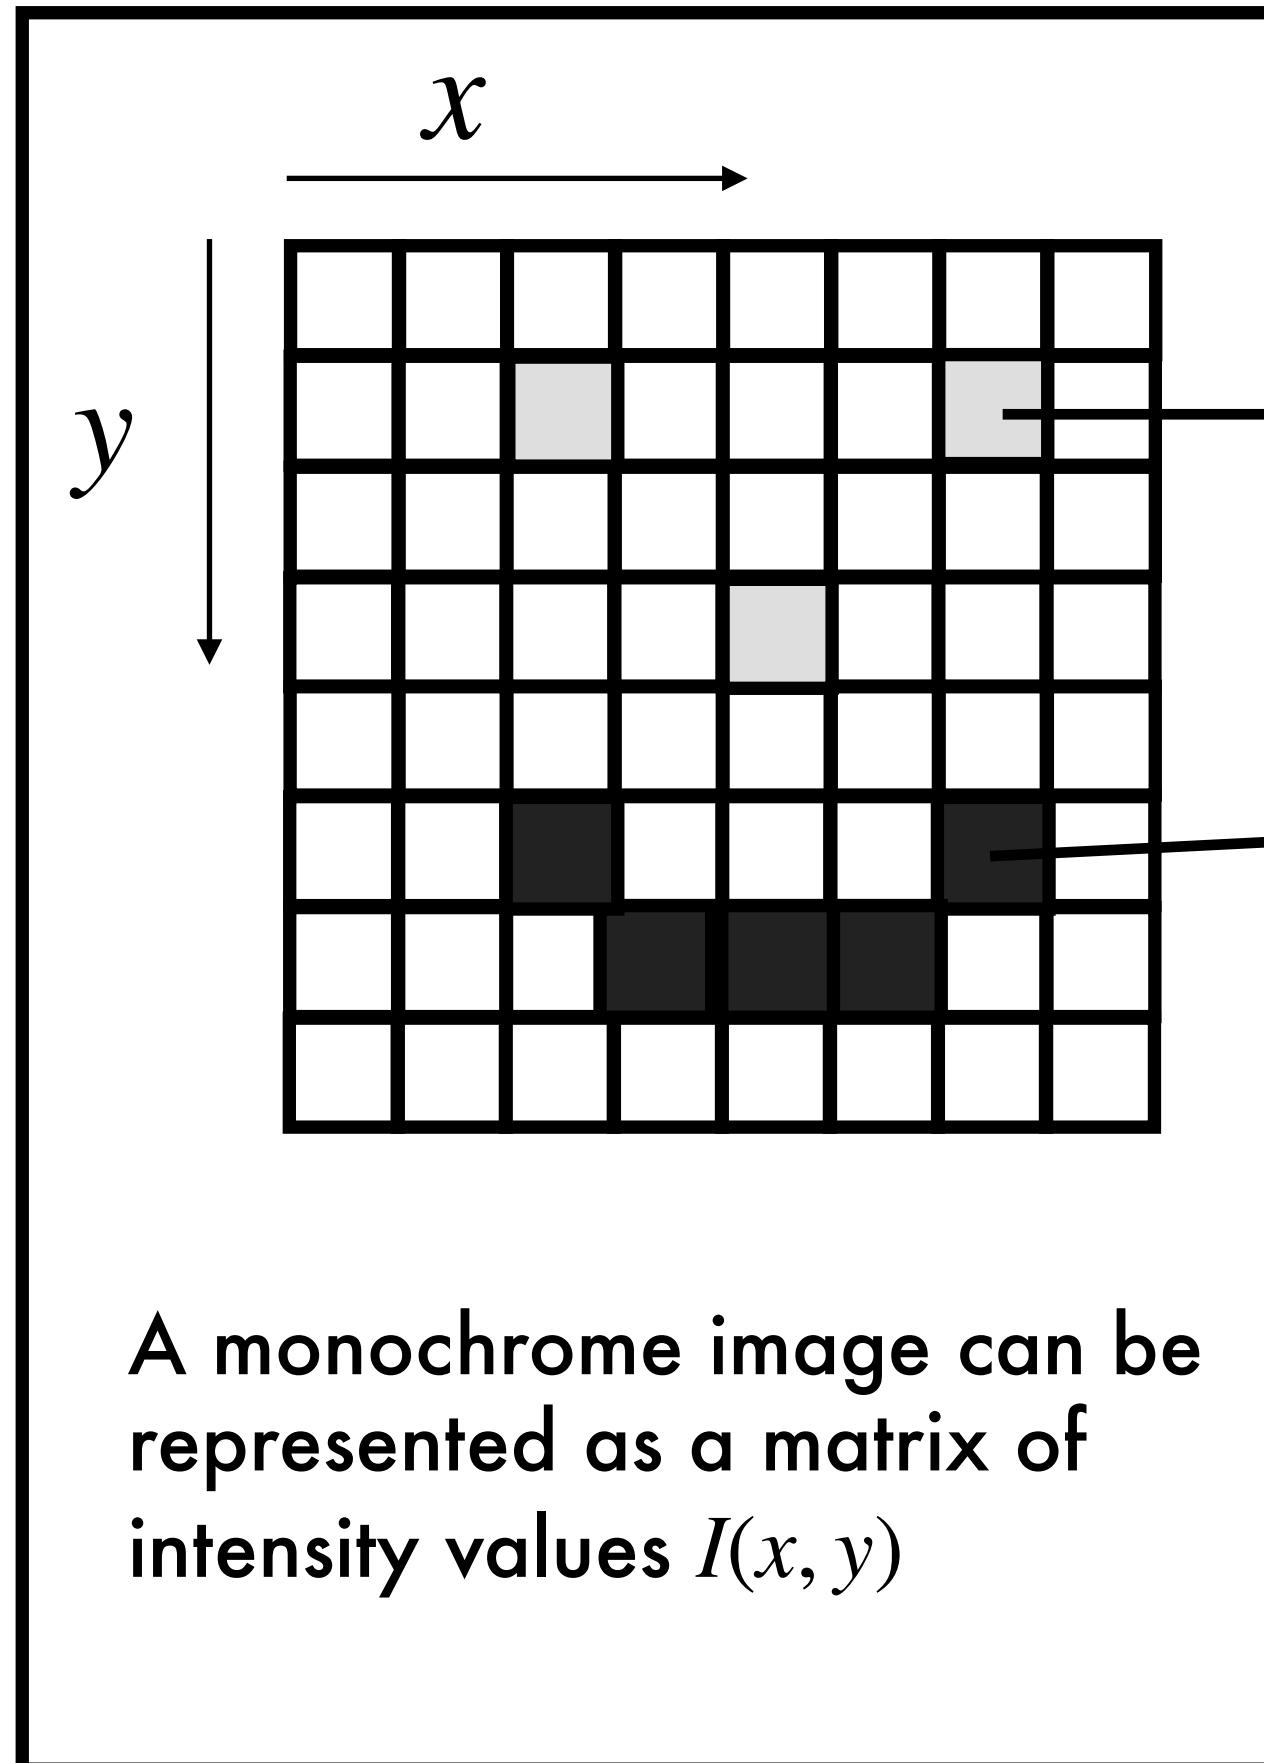
**If 3 mins spare:** check arxiv and make hot drink

**If 1 min spare:** check failed experiments, sigh, relaunch, and make hot drink

**Instructor: Samuel Albanie**

Based on course material authored by Roberto Cipolla

# Representing Images via Intensities



# Challenge: Nuisance factors in image data



If a point on an object (in this case, a *torta di mele*) is visible in view, the intensity  $I(x, y)$  is a function of many **geometric** and **photometric** variables:

- The position and orientation of the camera
- The geometry of the scene (3D shapes and layout)
- The nature and distribution of light sources
- Reflectance properties of the surfaces: specular-Lambertian, albedo 0 (black) - 1 (white)
- The properties of the camera lens and sensor array

In practice, the point may only be partially visible, or its appearance may also be affected by occlusion.

# Challenge: Data reduction

With current computers, it is necessary to discard most of the data from the camera before any attempt can be made at real-time interpretation.

Raw image data e.g. from GoPro 10 (CMOS) cameras:  
5.3K video at 60 fps



$O(10 \text{ Gbits/sec})$

iSight CCD camera:  
640x480 video at 30 fps



$O(100 \text{ Mbits/sec})$

How can we achieve such a dramatic **data reduction?**

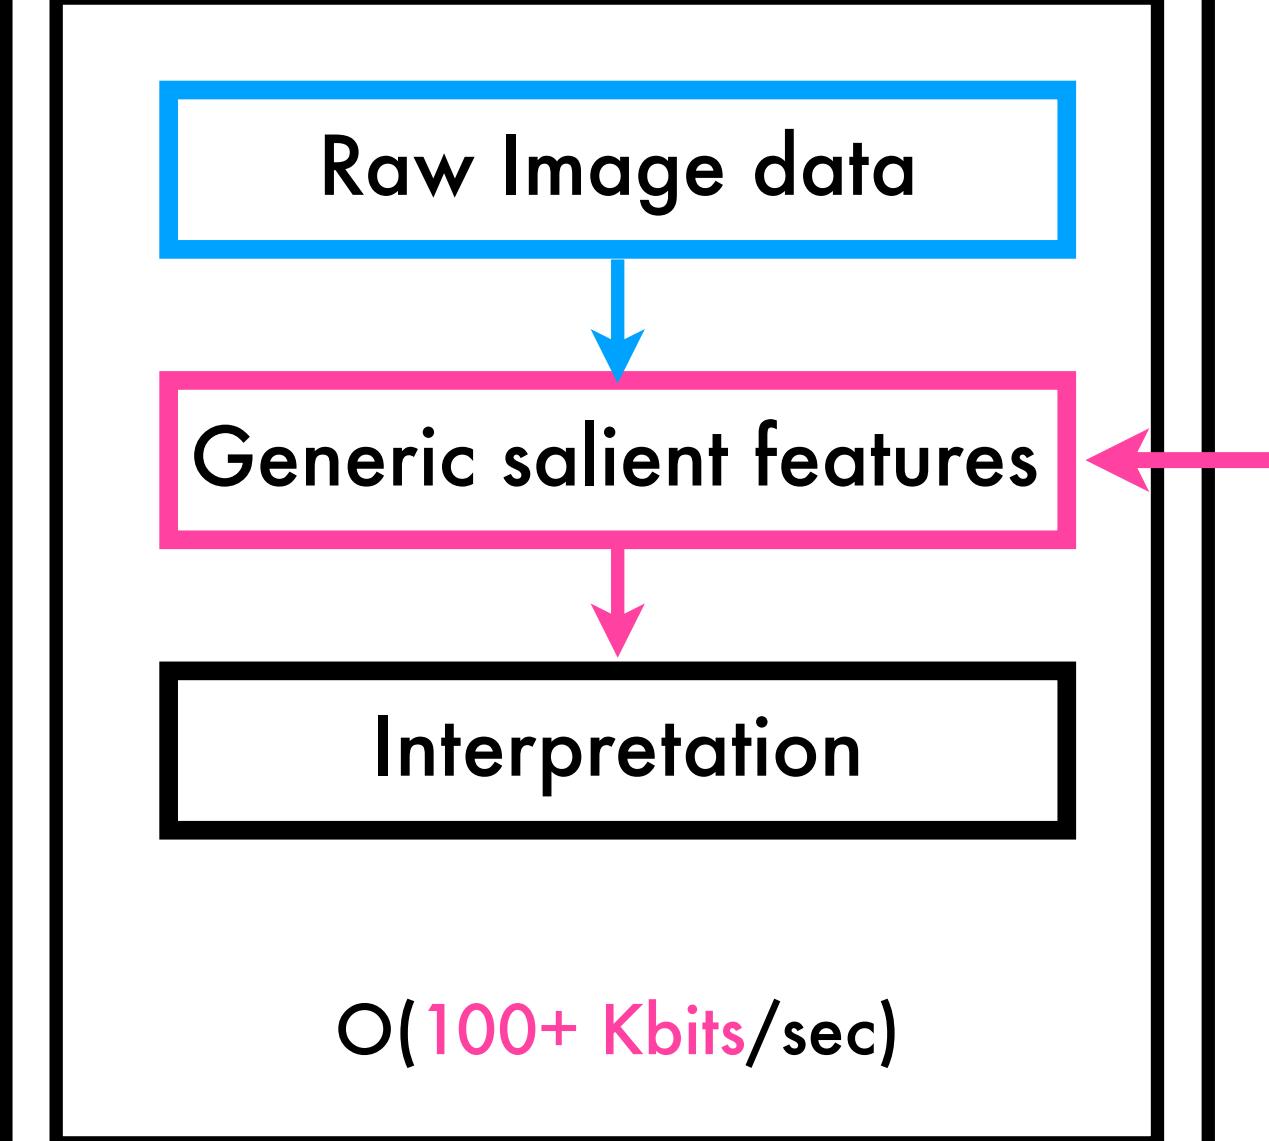


Image credits: <https://pixabay.com/photos/camera-go-pro-people-hand-mount-2590899/>  
[https://en.wikipedia.org/wiki/File:Apple\\_iSight\\_FireWire\\_Camera.jpg](https://en.wikipedia.org/wiki/File:Apple_iSight_FireWire_Camera.jpg)

Boyle and Smith won the Nobel Prize in Physics in 2009 for their invention of CCD technology

Research Trivia

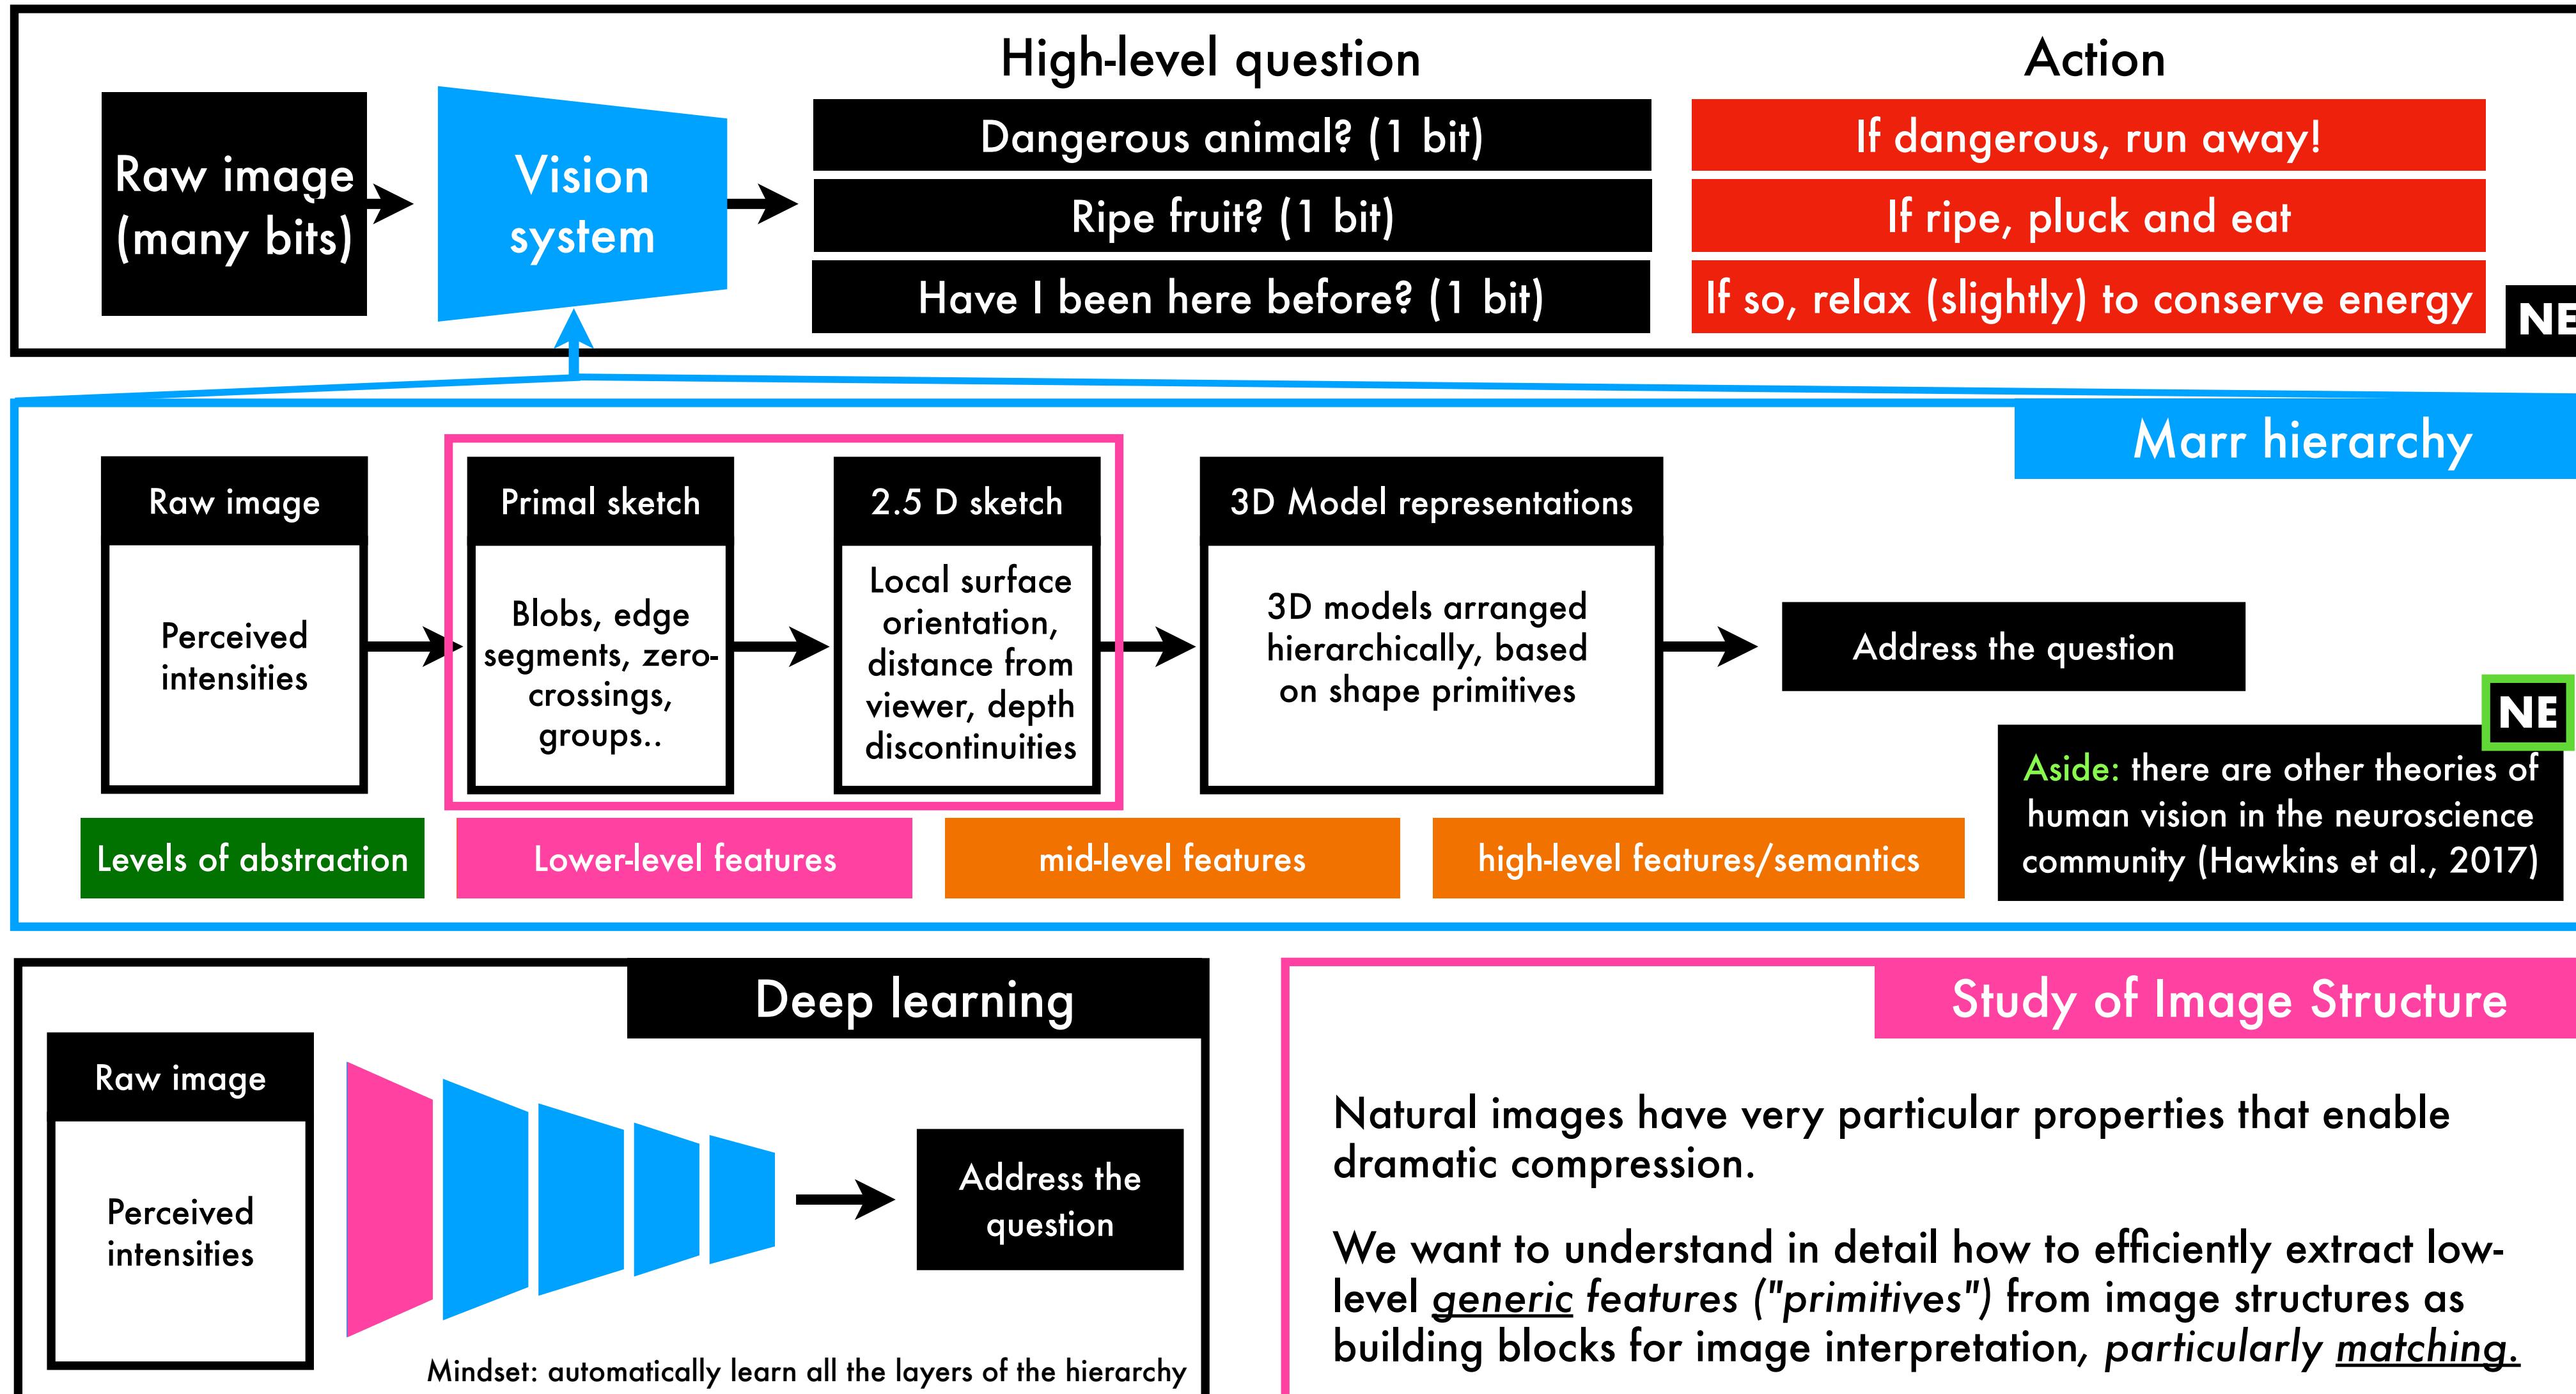
NE

OpenAI's powerful 2021 CLIP model (trained on thousands of GPUs) still only uses 224x224 pixel images

Research Trivia

NE

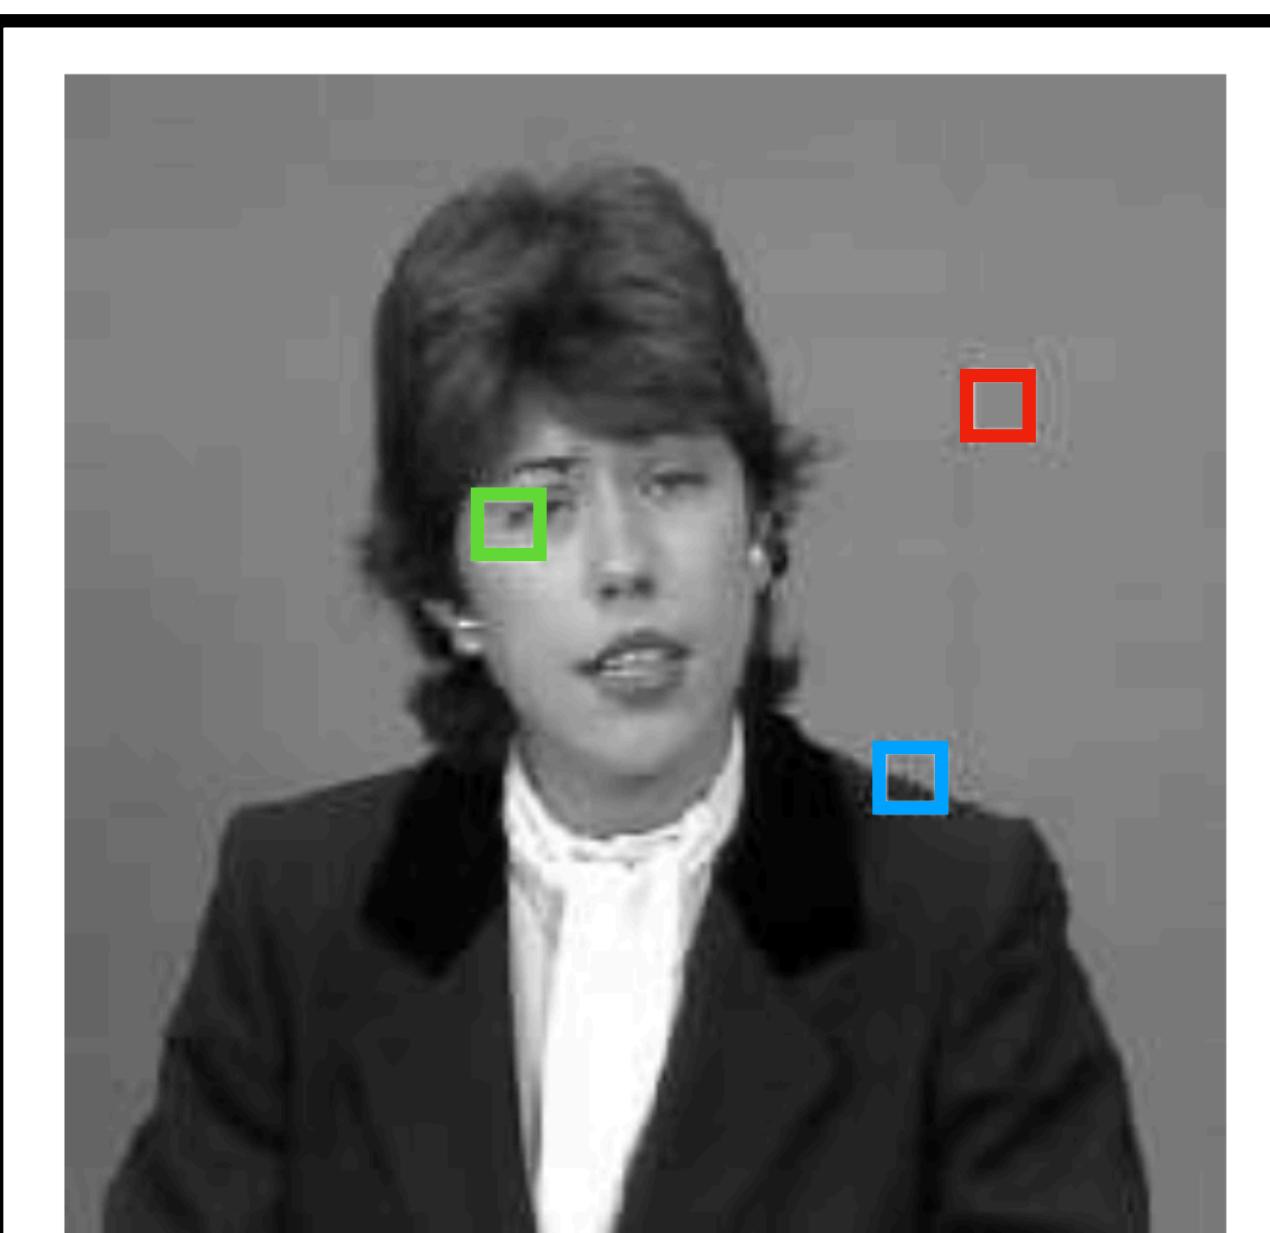
# Computer vision as hierarchical processing



Reference: D. Marr, "Vision: A computational investigation into the human representation and processing of visual information" (1982)

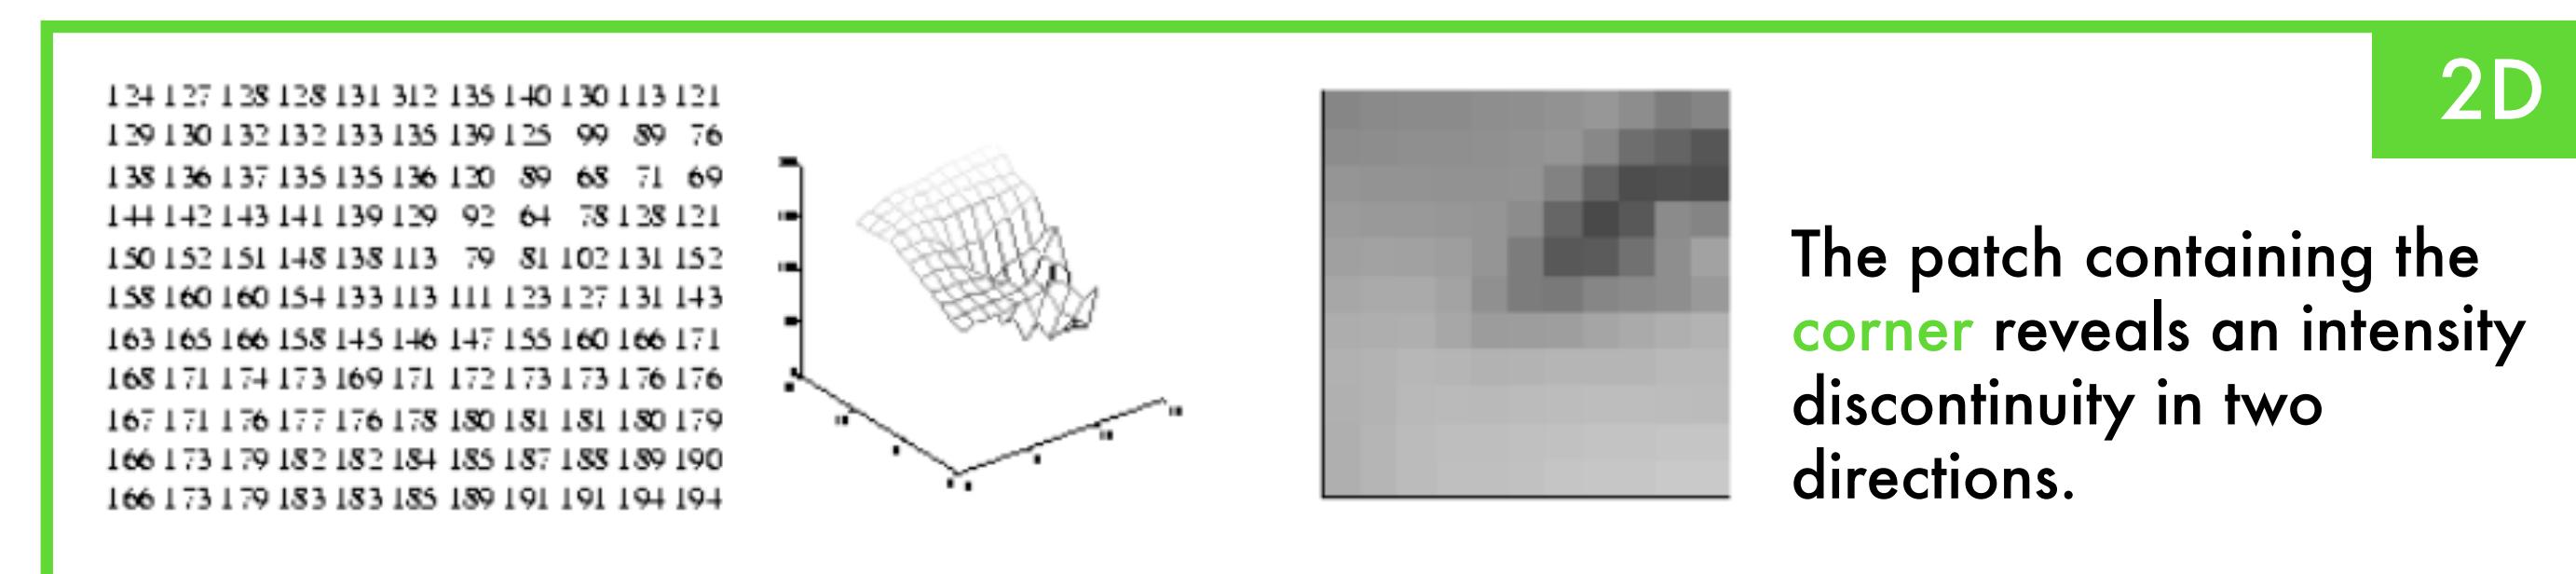
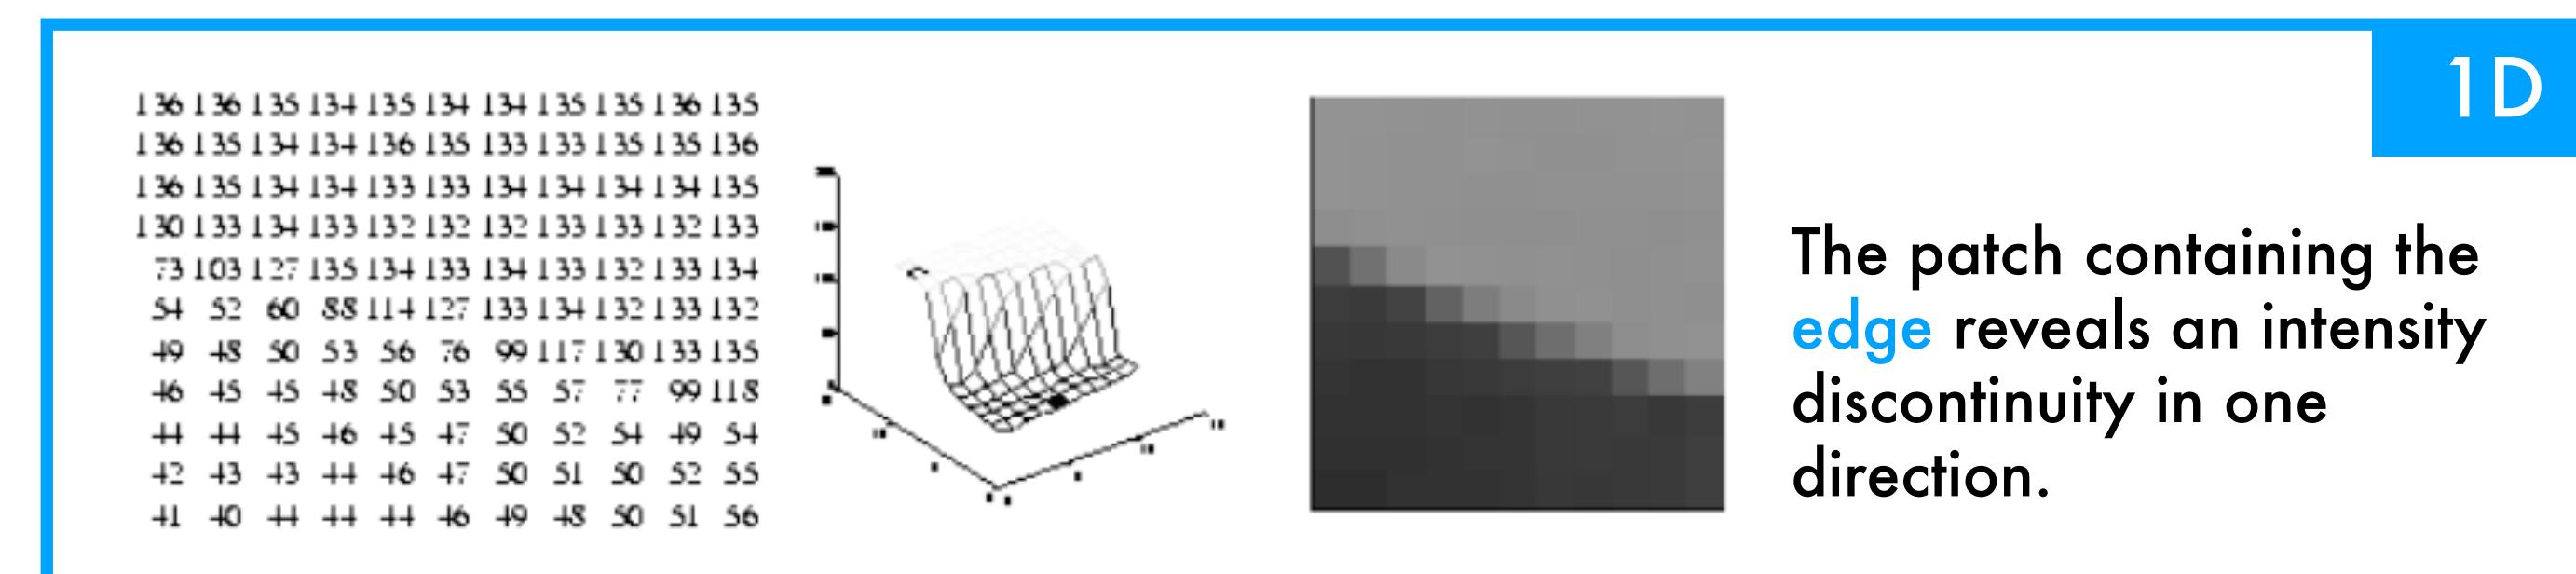
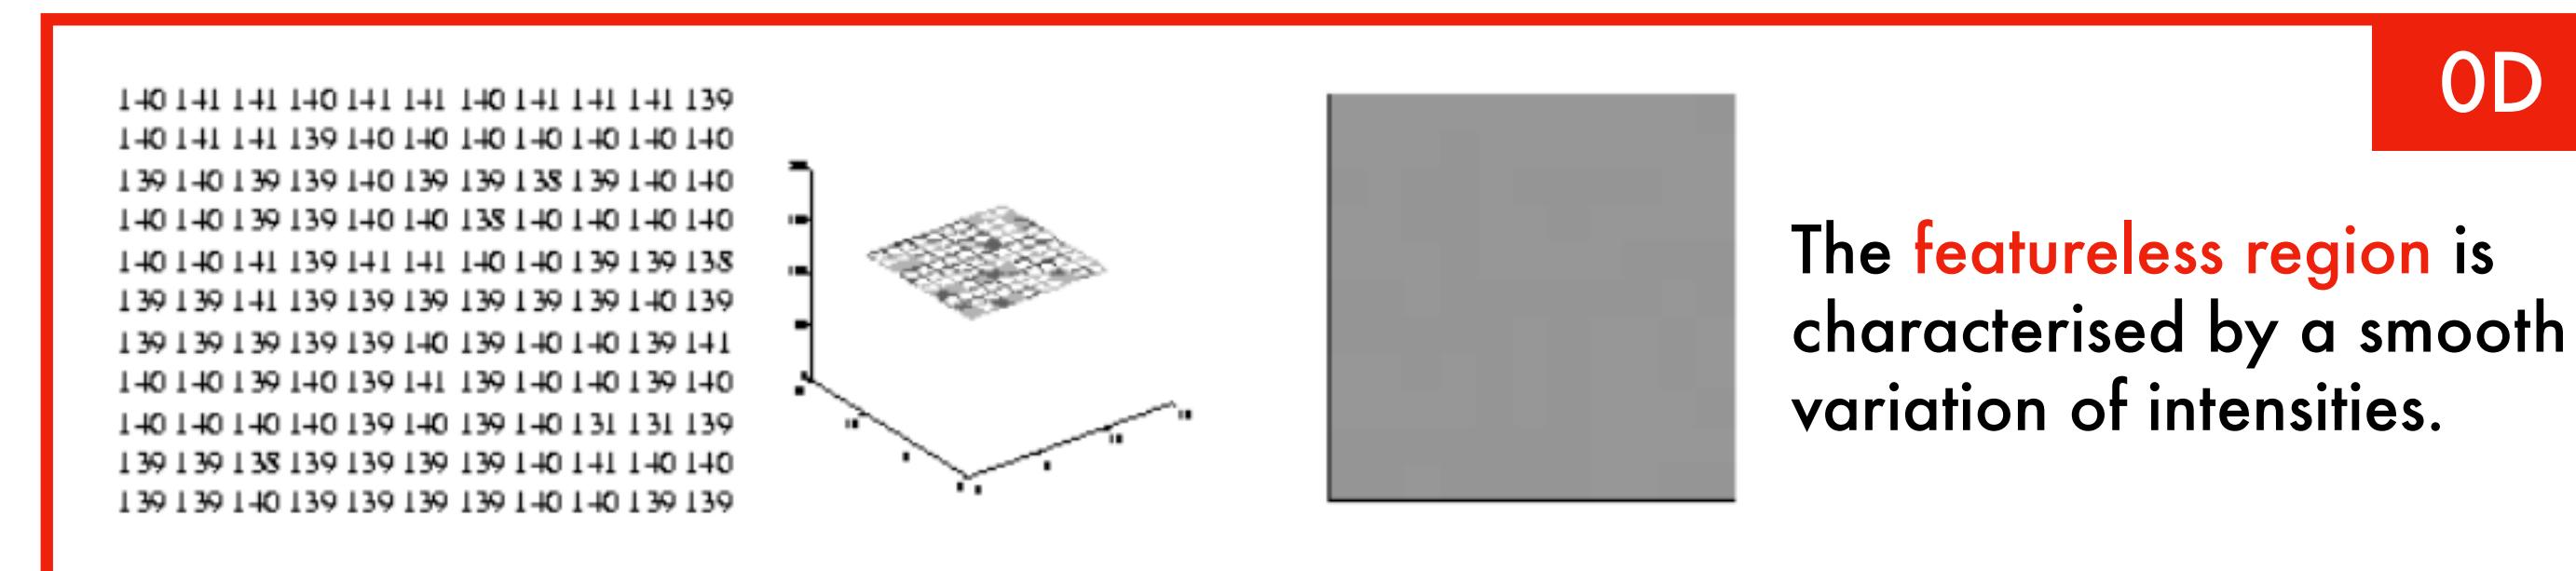
J. Hawkins, S. Ahmad, and Y. Cui, "A theory of how columns in the neocortex enable learning the structure of the world." *Frontiers in neural circuits* (2017)

# Image structure



**Let's examine pixel values in three patches in this photo of Claire:**

- A featureless region
  - An edge
  - A corner



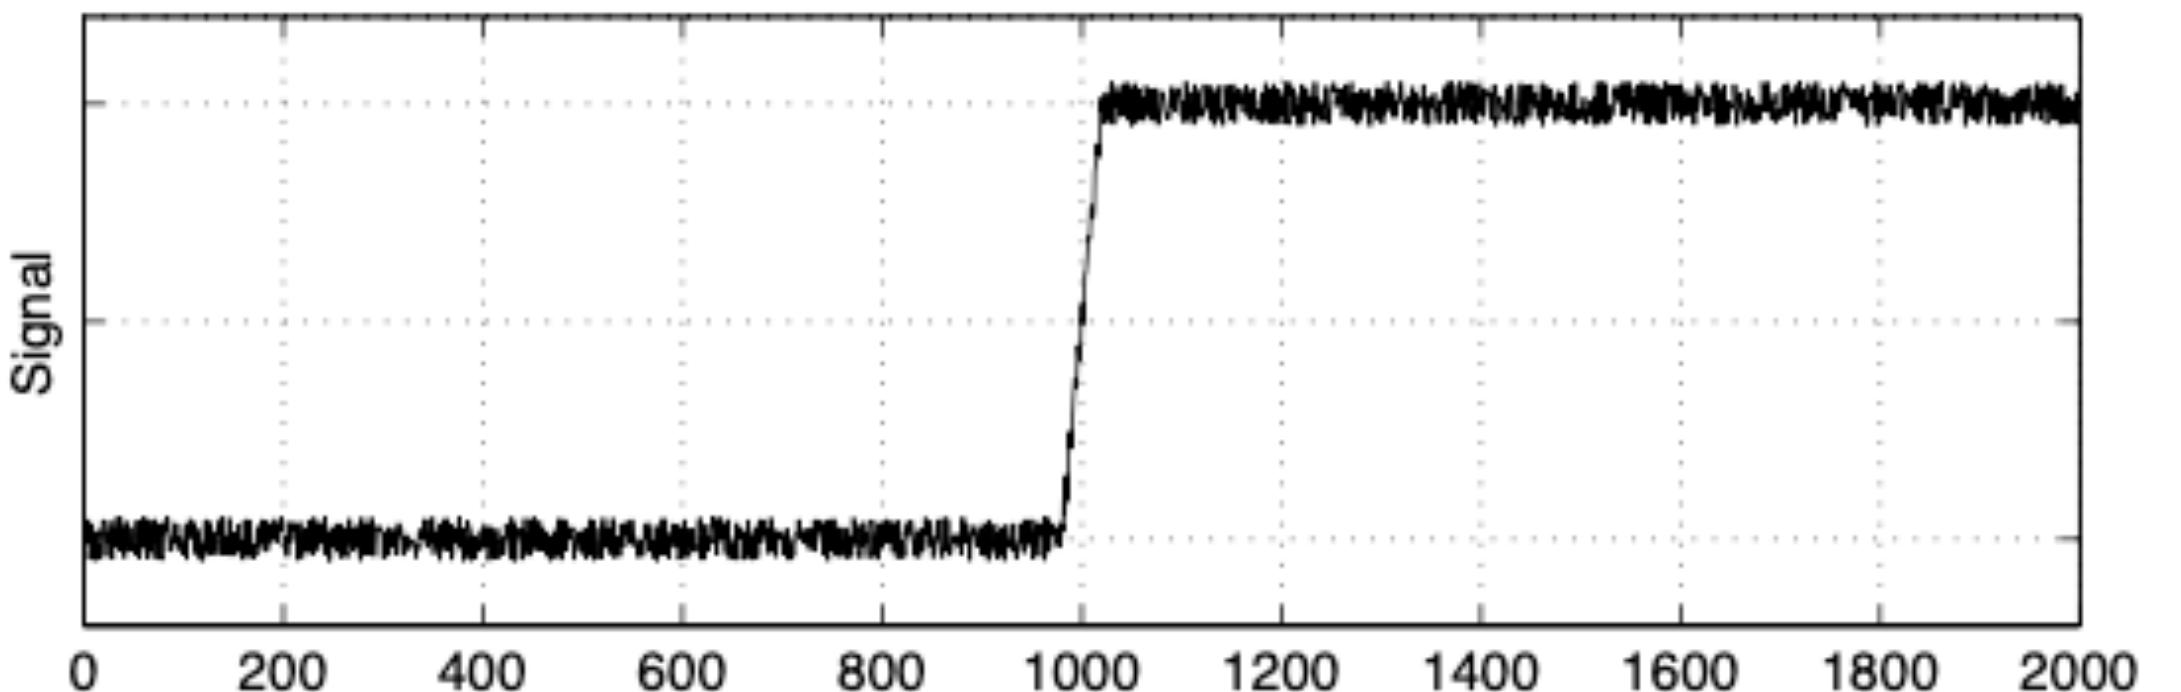
Note that an **edge** or **corner** representation imparts a desirable invariance to lighting: **the intensity discontinuities are likely to be prominent, whatever the lighting conditions.**

**Computational question: *How can we find these structures efficiently?***

# 1D Edge Detection

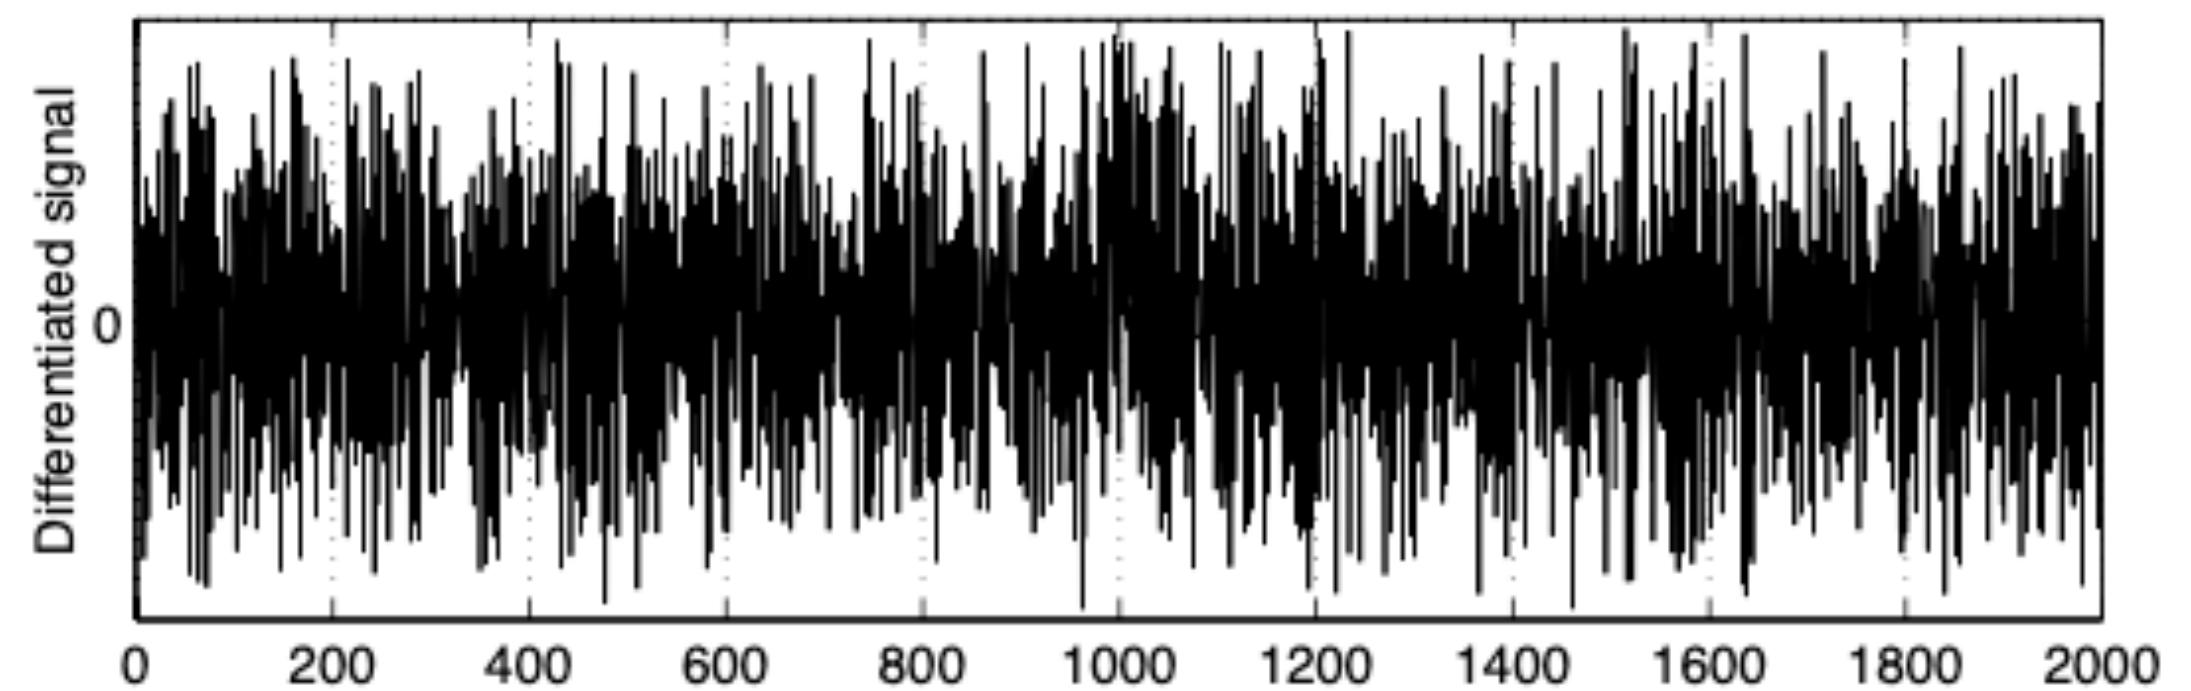
When developing an edge detection algorithm, it is important to bear in mind the invariable presence of **image noise**.

Consider this signal  $I(x)$  with an obvious edge:



An intuitive approach to edge detection might be to look for **maxima** and **minima** in  $I'(x)$ .

The derivative of the signal looks like this:



Oh dear: it's hard to spot the edge in this signal!

Our simple strategy was defeated by high-frequency **noise** (which is amplified by differentiation).

For this reason, all edge detectors start by **smoothing** the signal to suppress noise.

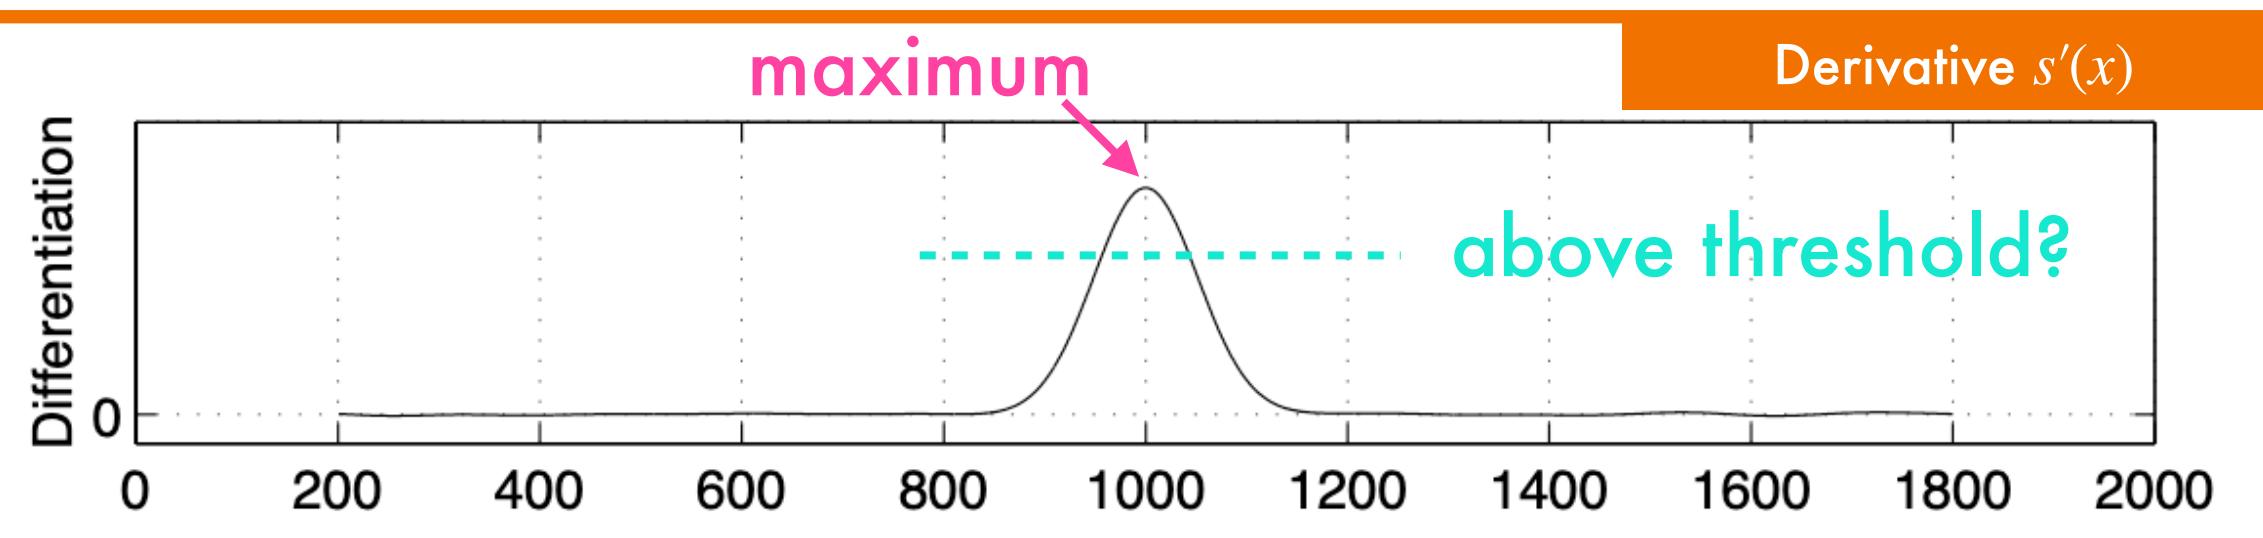
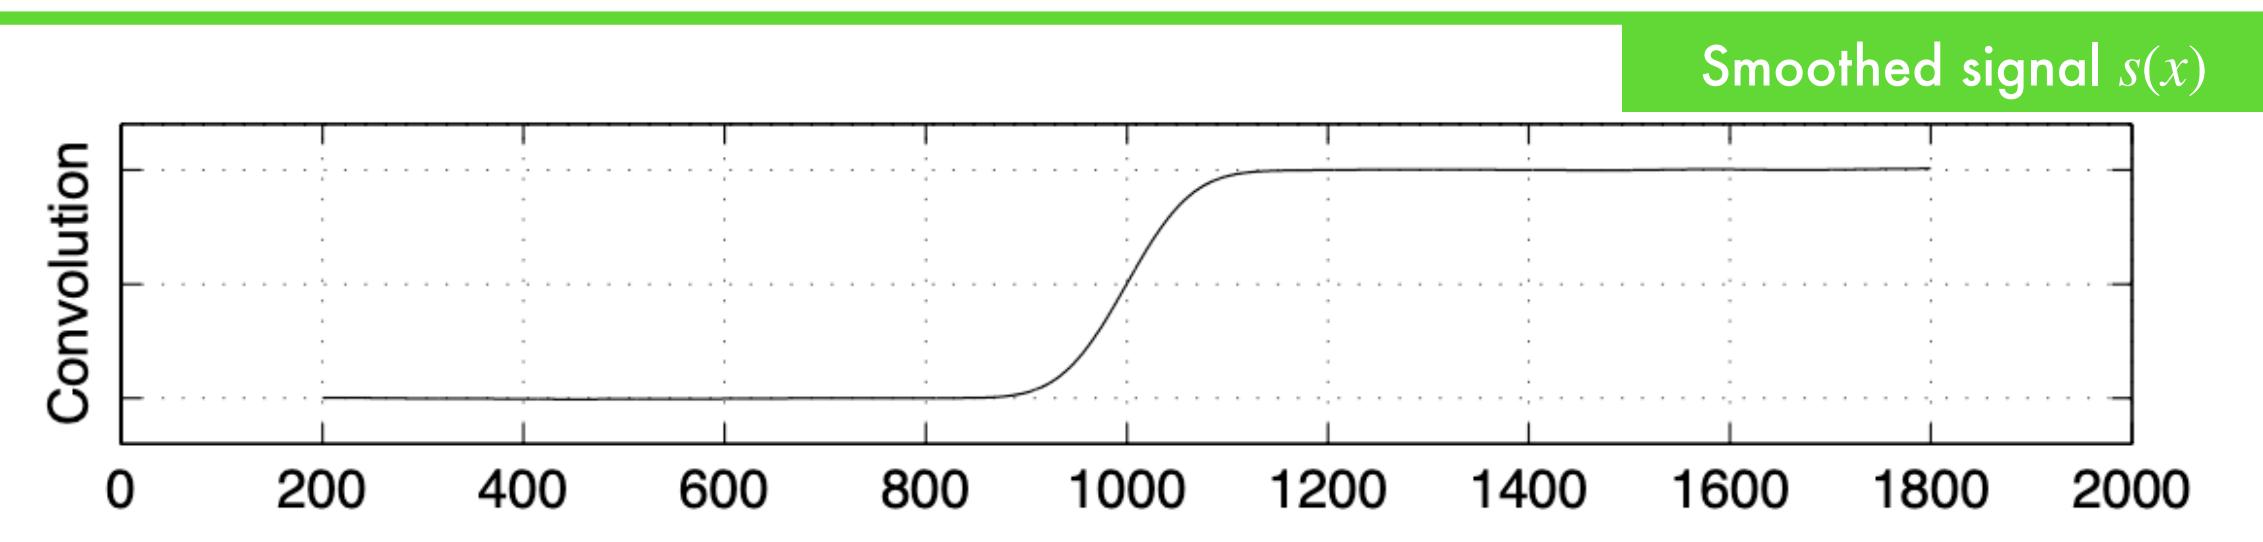
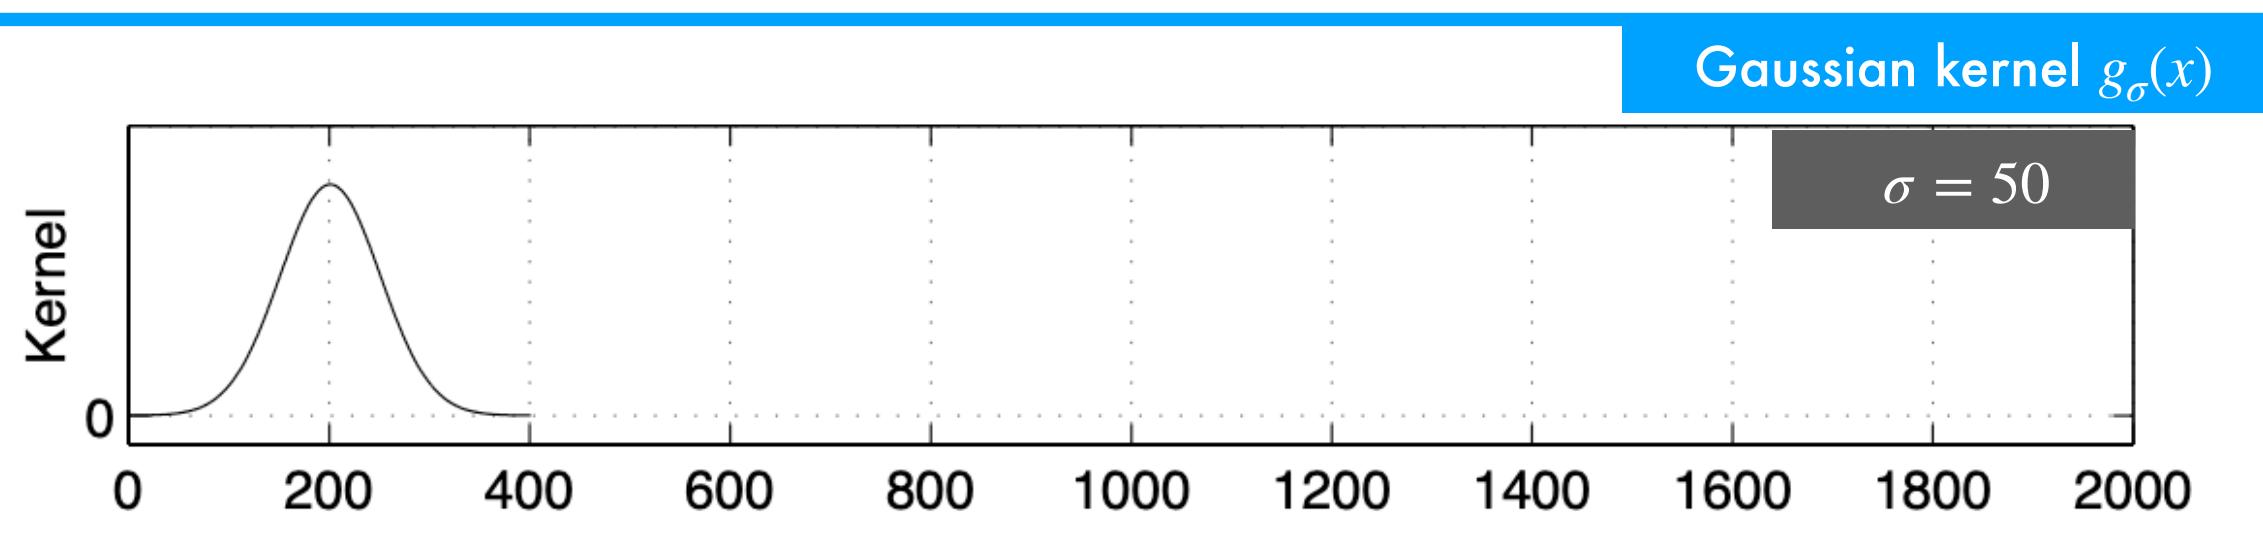
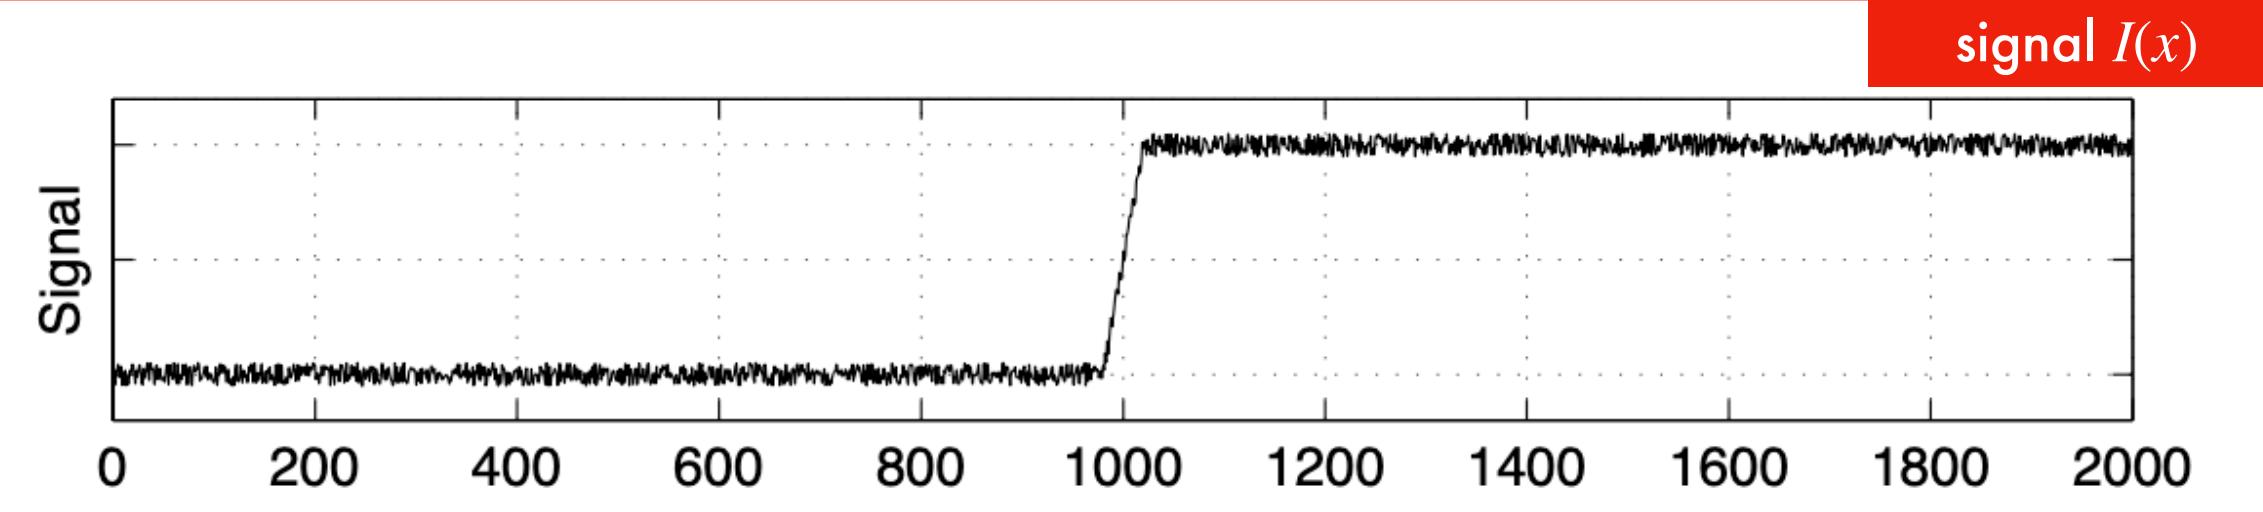
The most common approach is to use a Gaussian filter (a low-pass filter that suppresses high frequencies).

# 1D Edge Detection (with smoothing)

## 1D Edge Detection algorithm

A broad overview of 1D edge detection is:

1. First, convolve the **signal  $I(x)$**  with a **Gaussian kernel**  $g_\sigma(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{x^2}{2\sigma^2}\right)$  to produce smooth  $s(x)$ .
2. Compute  $s'(x)$ , the derivative of  $s(x)$ .
3. Find **maxima** and **minima** of  $s'(x)$ .
4. Use **thresholding** on the magnitude of the extrema to mark edges.



# 1D Edge Detection: a computational trick

The smoothing in step 1 was performed by a 1D convolution with a Gaussian:

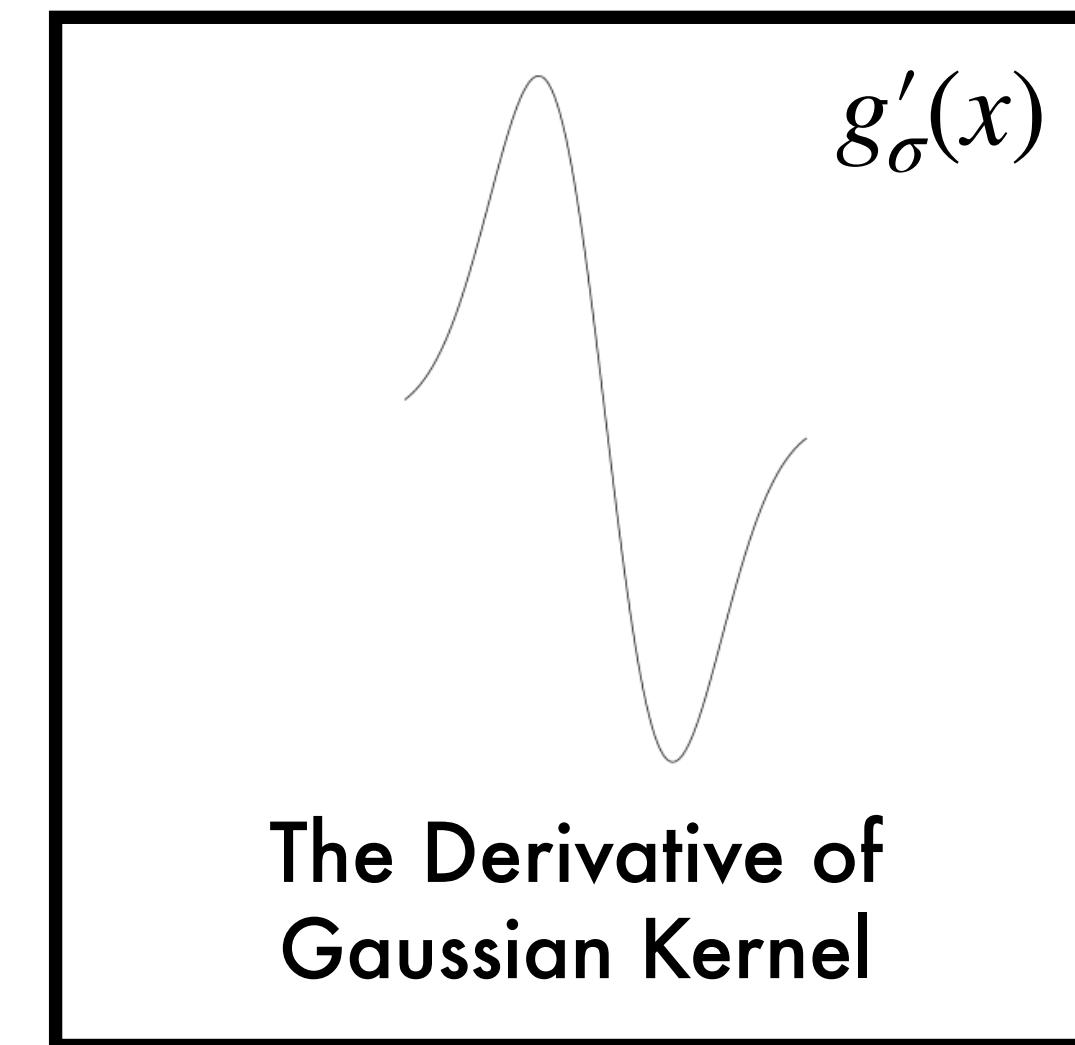
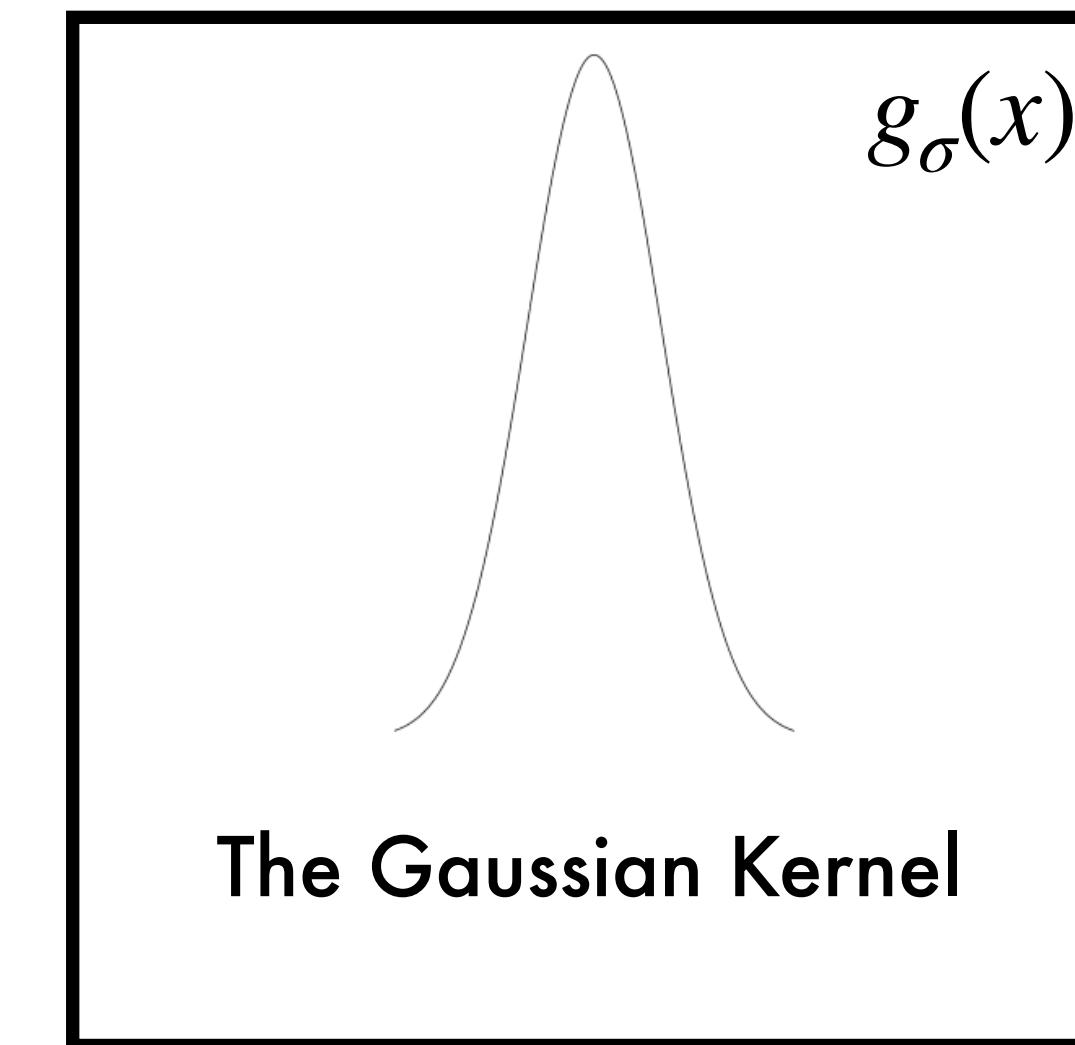
$$s(x) \circledast g_\sigma(x) = \int_{-\infty}^{\infty} g_\sigma(u)I(x-u)du = \int_{-\infty}^{\infty} g_\sigma(x-u)I(u)du$$

The differentiation in step 2 is also performed by a 1D convolution. So it seems that edge detection requires two computationally expensive convolutions.

However, the **derivative theorem of convolution** comes to the rescue!

$$s'(x) = \frac{d}{dx}[g_\sigma(x) \circledast I(x)] = g'_\sigma(x) \circledast I(x)$$

So we can compute  $s'(x)$  with just a single convolution - a major saving!



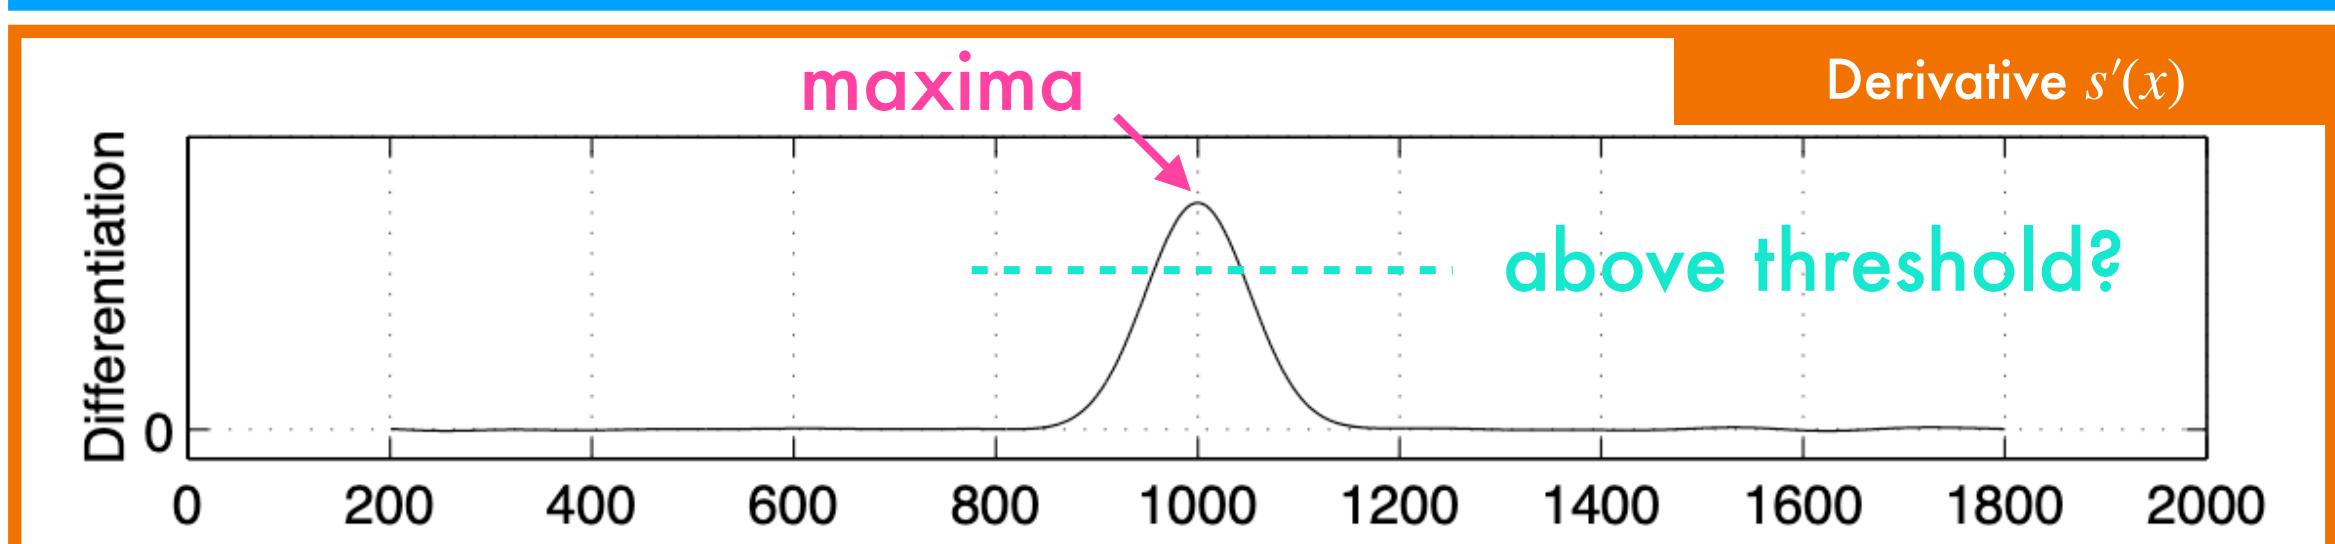
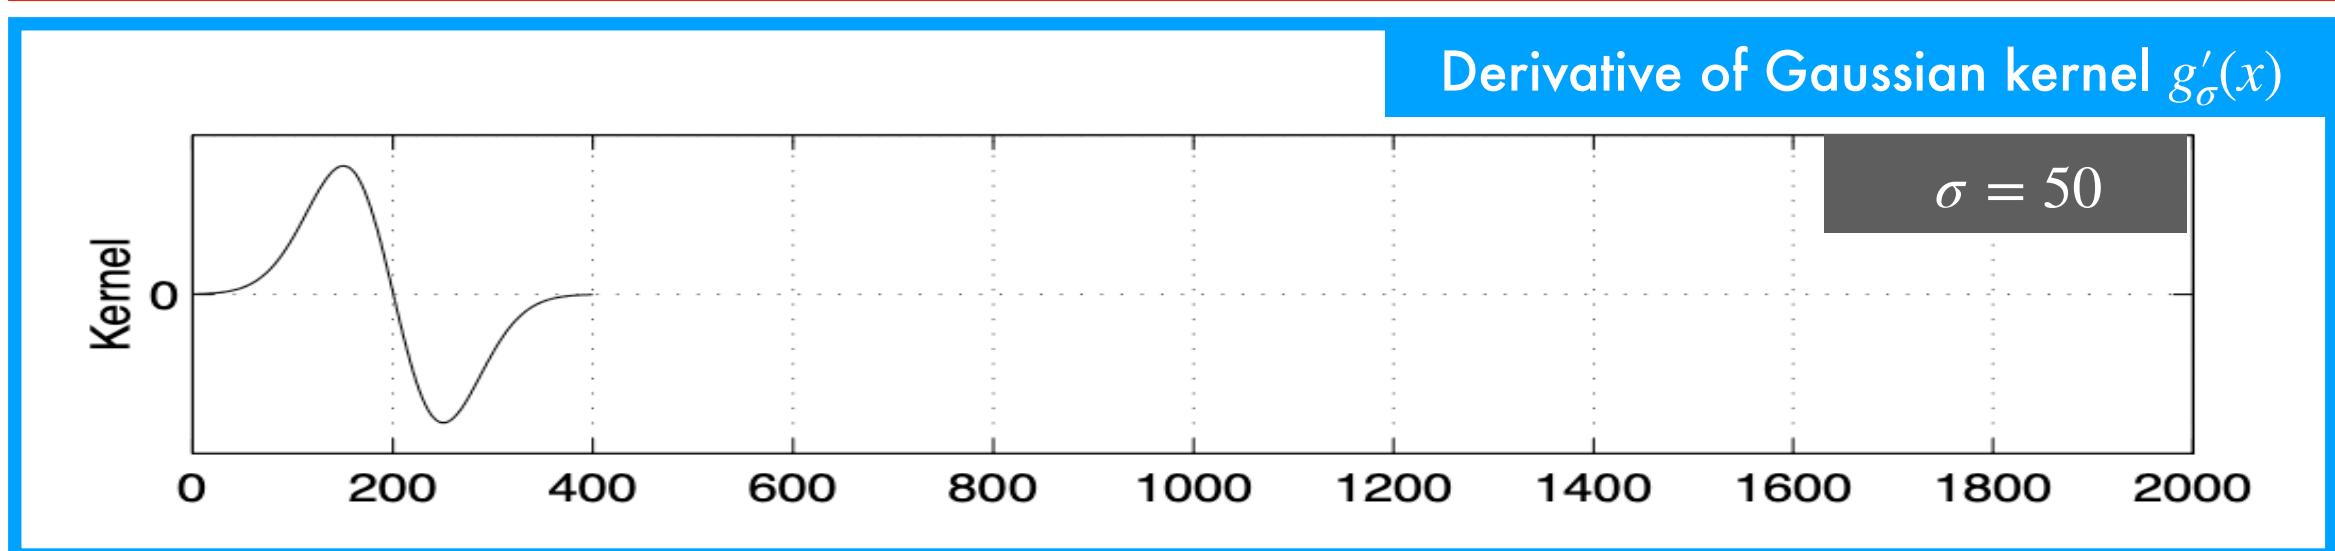
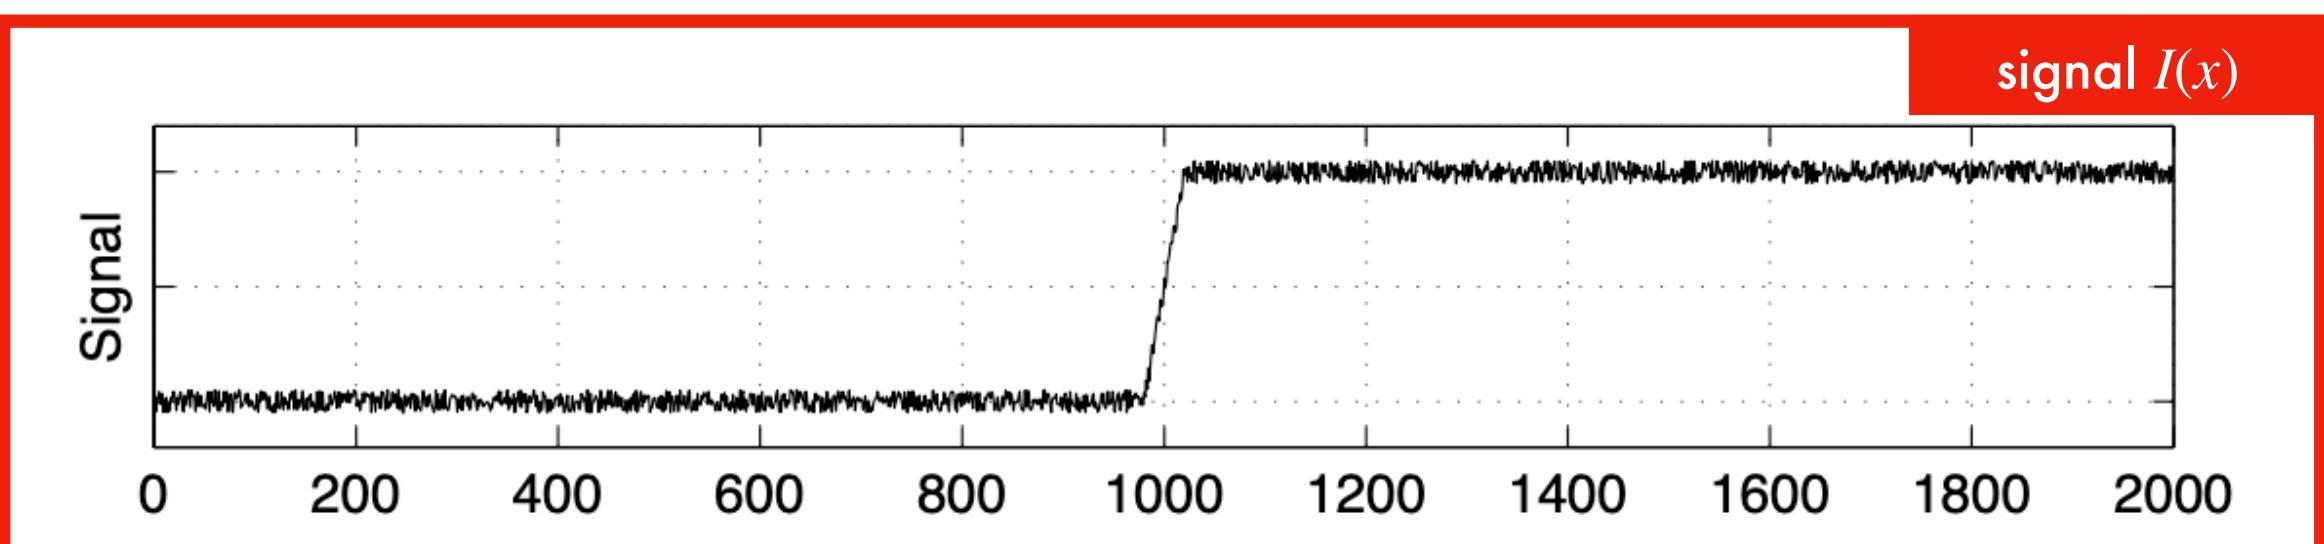
# 1D Edge Detection (faster)

## Faster 1D Edge Detection algorithm

The fast variant of the edge detection algorithm becomes

1. Convolve the **signal  $I(x)$**  with a **derivative of Gaussian Kernel  $g'_\sigma(x)$**  to produce  **$s'(x)$**  directly.
2. Find **maxima** and **minima** of  **$s'(x)$** .
3. Use **thresholding** on the magnitude of the extrema to mark edges.

VE



# 1D Edge Detection: zero-crossings

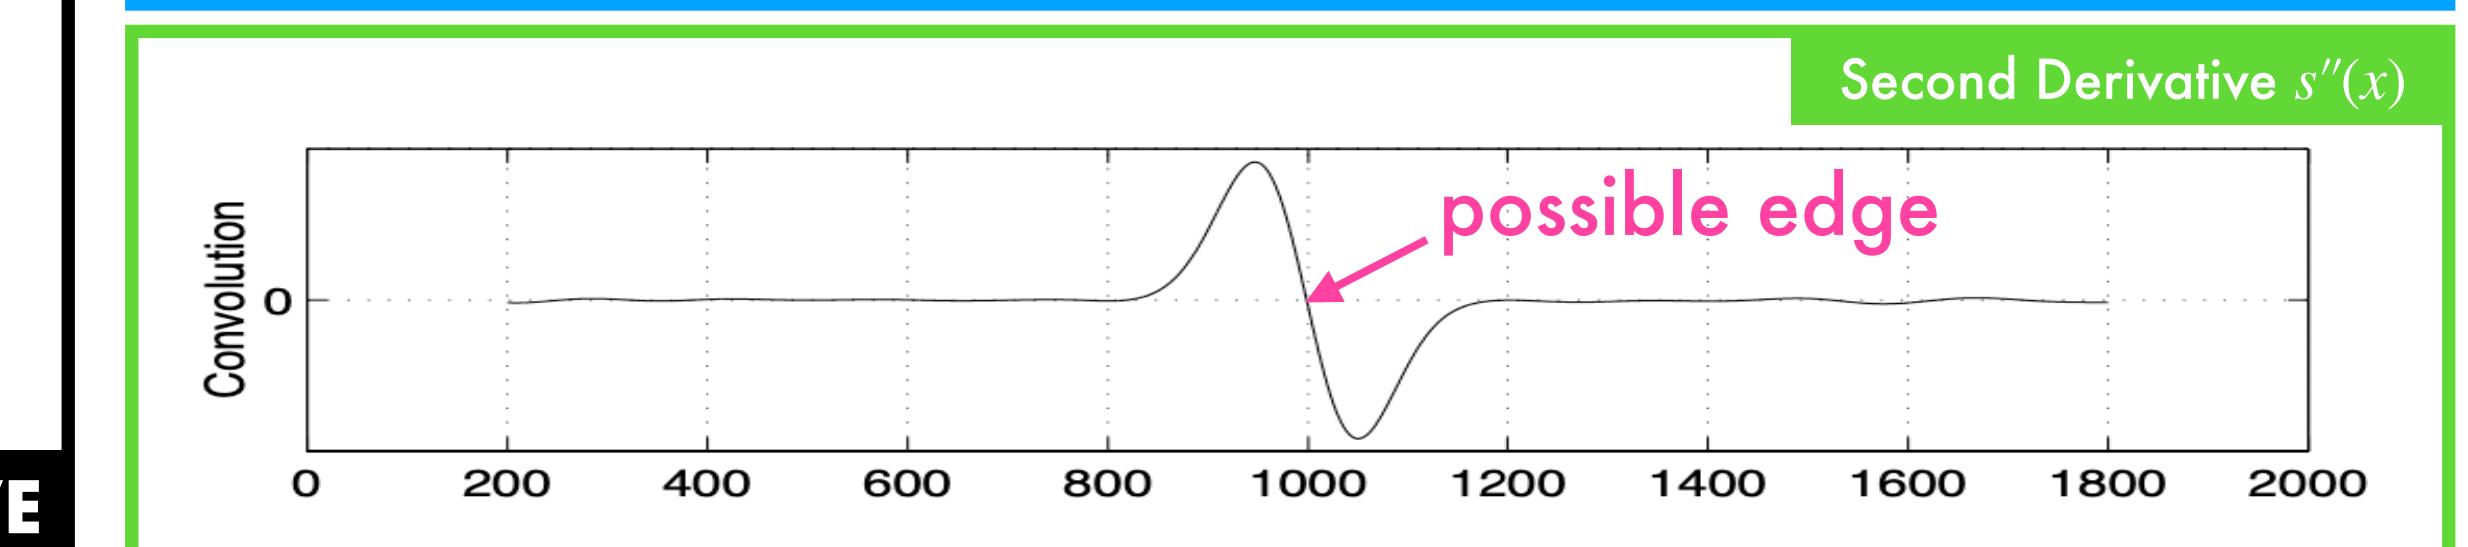
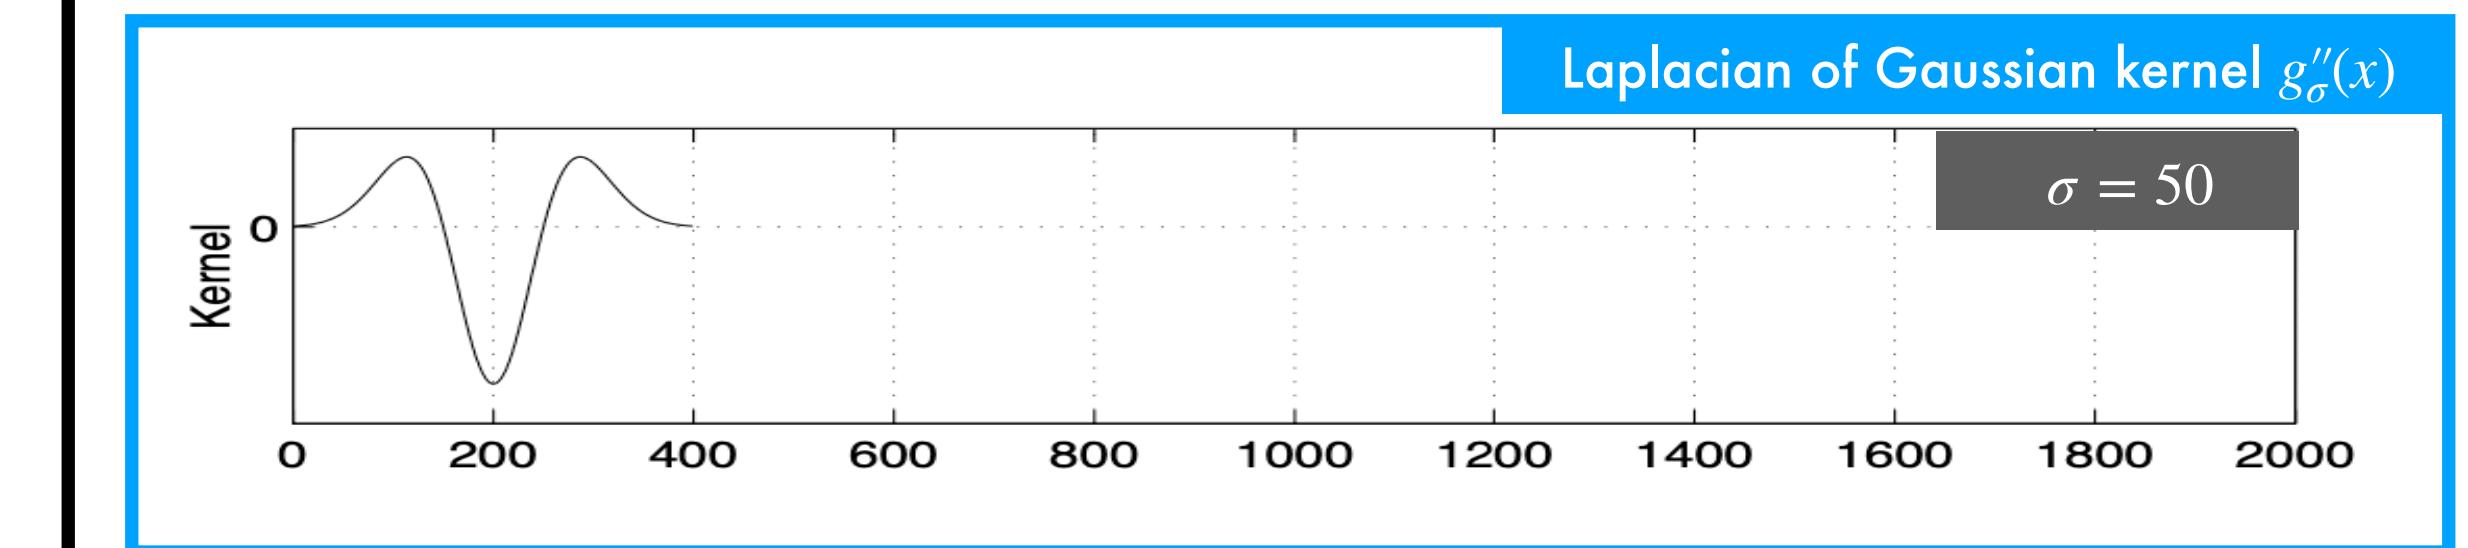
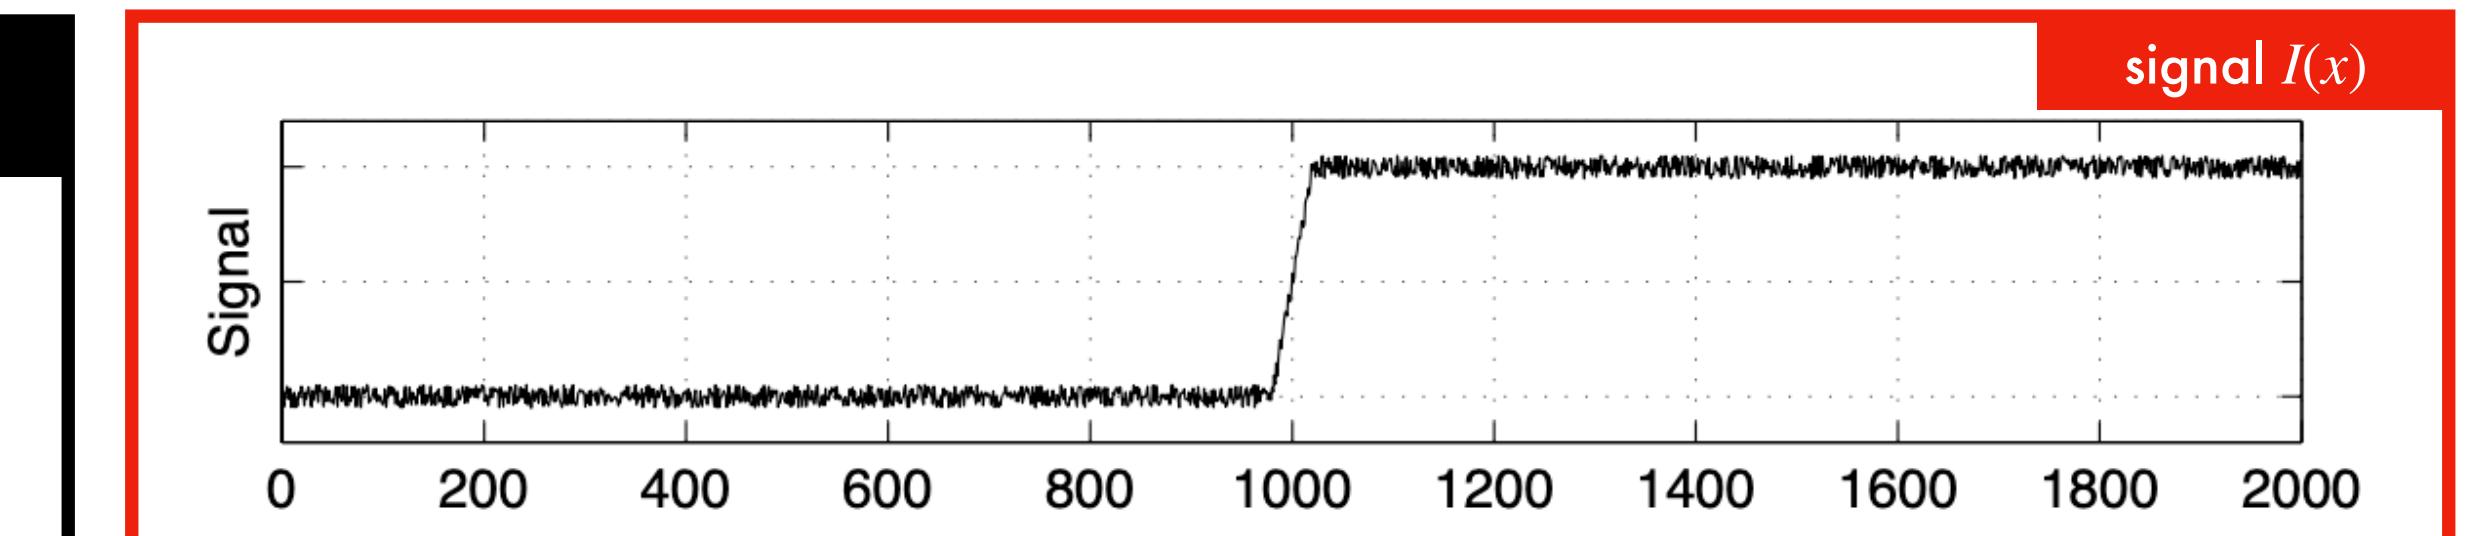
Fastest 1D Edge Detection algorithm

Finding **maxima** and **minima** of  $s'(x)$  is the same as looking for zero-crossings of  $s''(x)$ !

In many implementations of edge detection algorithms, the **signal** is convolved with the Laplacian of a Gaussian ("LoG" kernel),  $g''_\sigma(x)$ , by applying the derivative theorem of convolution a second time:

$$s''(x) = g''_\sigma(x) \circledast I(x)$$

The zero crossings of  $s''(x)$  mark **possible edges**.

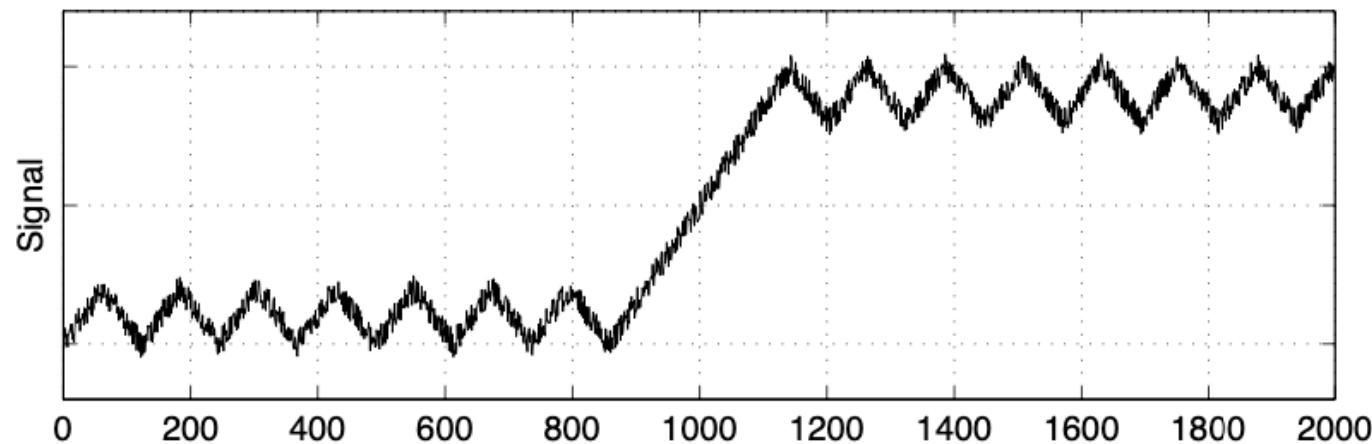


VE

# 1D Edge Detection: scale

We have not yet addressed the critical issue of what value of  $\sigma$  to use.

Consider this signal:

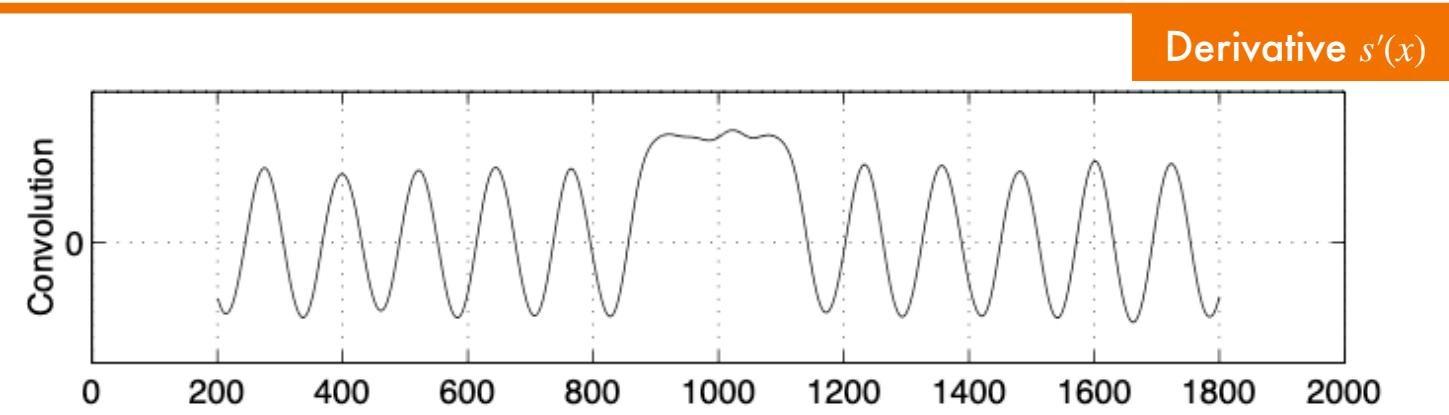
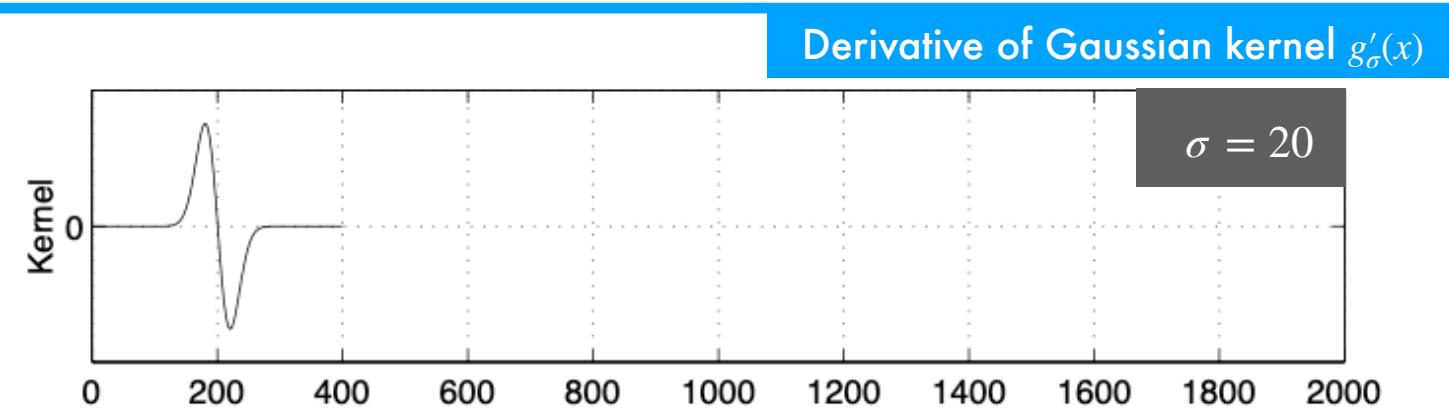
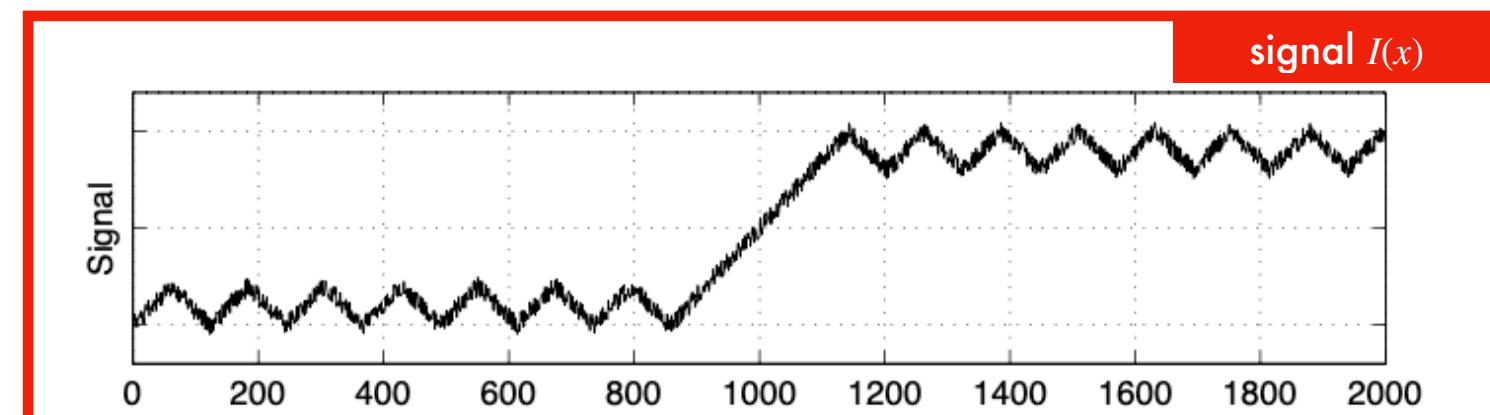


Does the signal have one "positive" edge or a number of "positive" and "negative" edges?

It's up to you to choose!

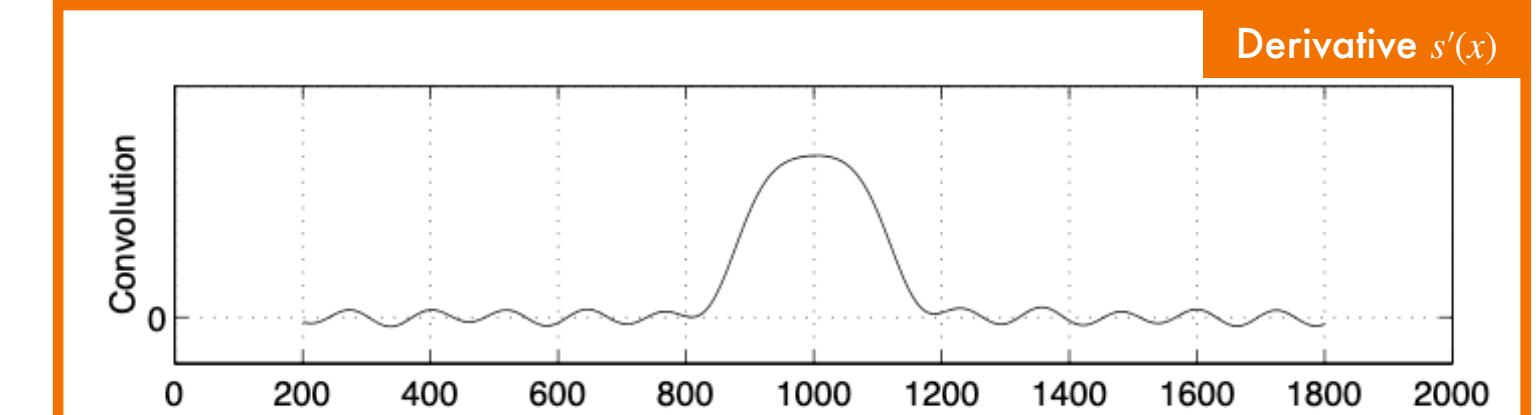
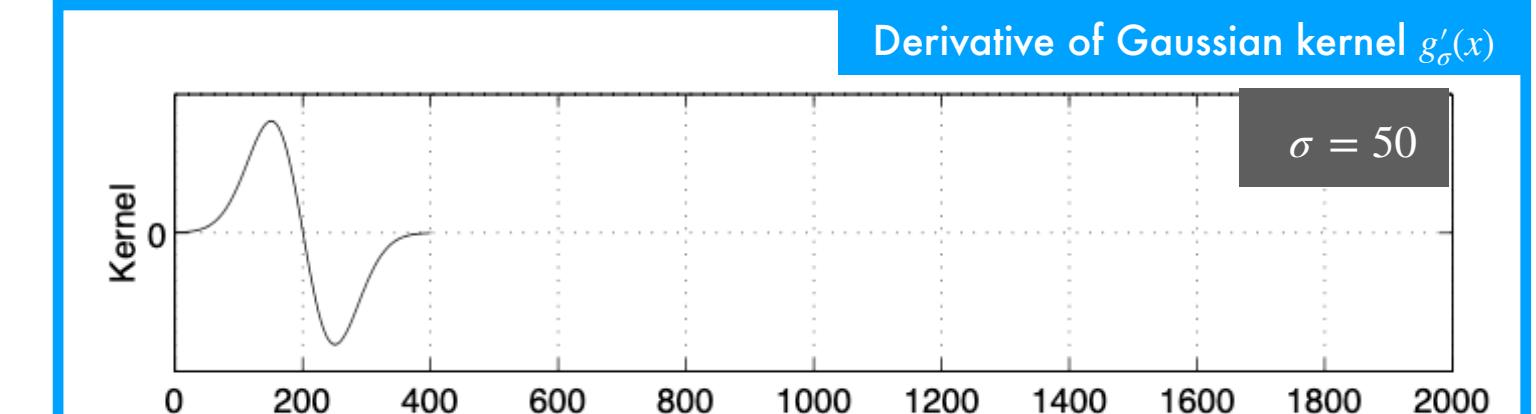
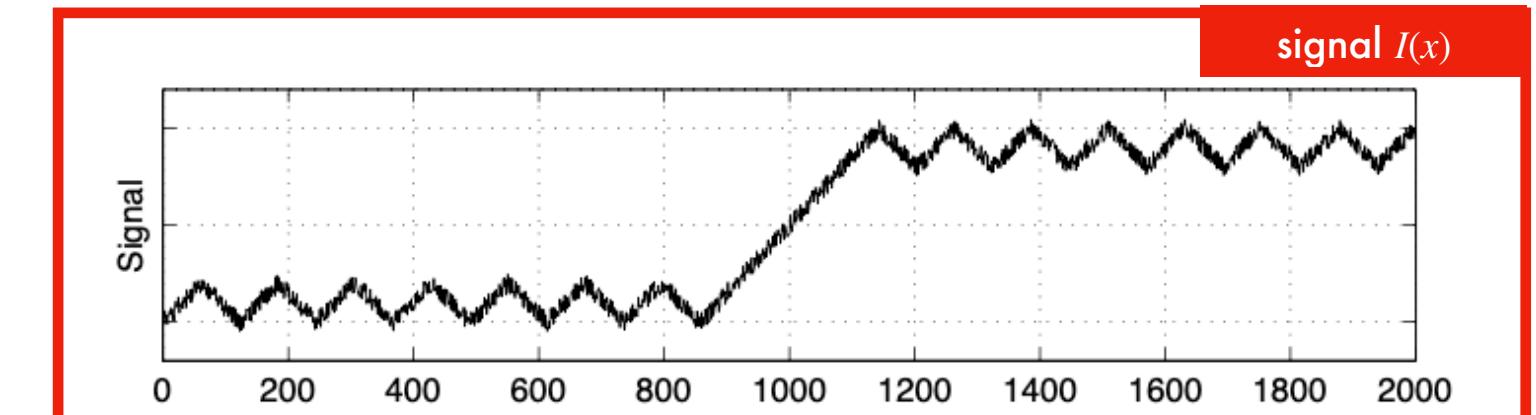
Fine detail edges

Using a small  $\sigma$  brings out all the edges.



Coarse detail edges

As  $\sigma$  increases, the signal is smoothed more and more, and only the central edge survives.

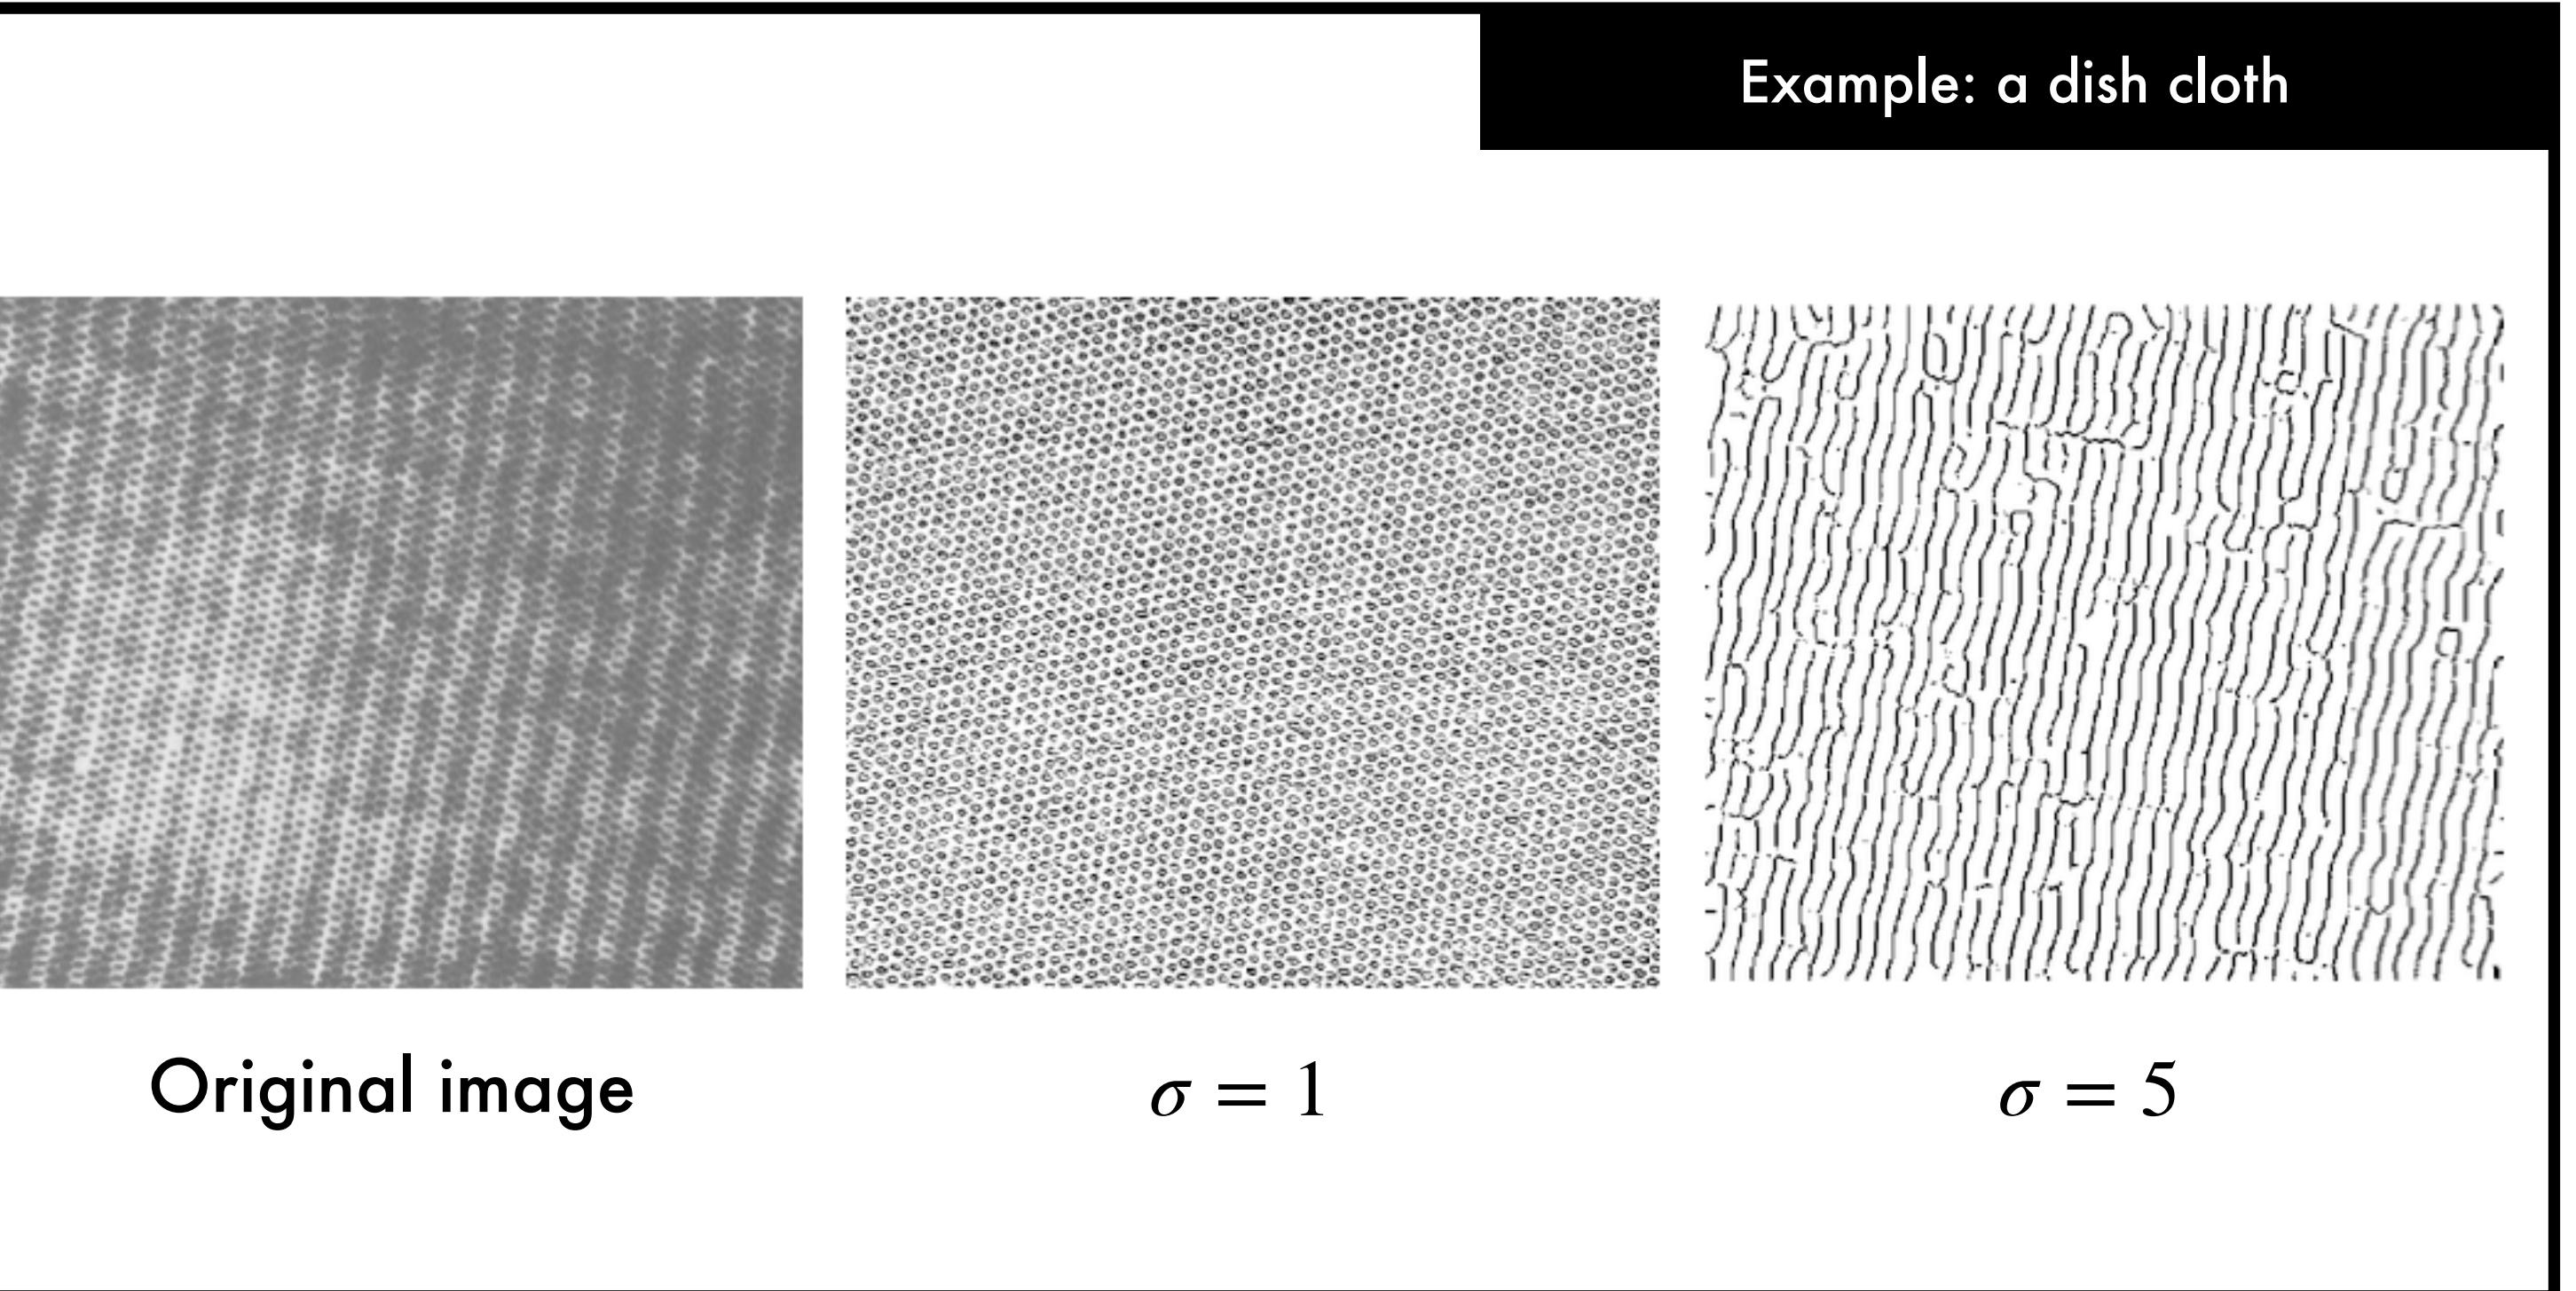


# 1D Edge Detection: multi-scale

## The link between **smoothing** and **scale**

The amount of **smoothing** controls the **scale** at which we analyse the image. There is no right or wrong size for the Gaussian kernel: it all depends on the scale we're interested in.

Modest smoothing (a Gaussian kernel with **small  $\sigma$** ) brings out edges at a **fine scale**. More smoothing (**larger  $\sigma$** ) identifies edges at **larger scales**, suppressing the finer detail.



Note: Fine scale edge detection is particularly sensitive to noise (less of an issue when analysing images at coarse scales).

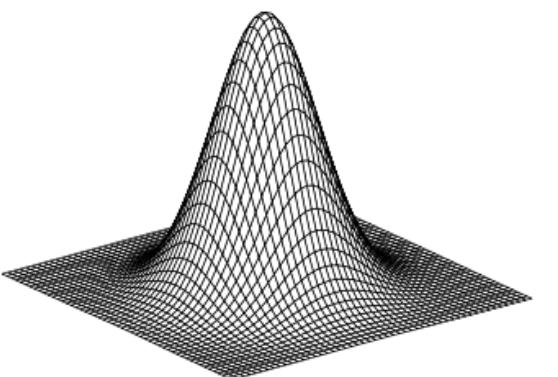
# 2D Edge Detection (step 1: smoothing)

## Step 1: smoothing

The 1D edge detection scheme can be extended to work in two dimensions.

First we smooth the image  $I(x, y)$  by convolving with a 2D Gaussian  $G_\sigma(x, y)$ :

$$G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$



$$S(x, y) = G_\sigma(x, y) \circledast I(x, y)$$

$$= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} G_\sigma(u, v) I(x - u, y - v) du dv$$

## Effects of Gaussian smoothing



Original image



$\sigma = 3$  pixels



$\sigma = 4$  pixels

# 2D Edge Detection (step 2: gradients)

## Step 2: gradients

The second step is to compute the gradient of the smoothed image  $S(x, y)$  at every pixel:

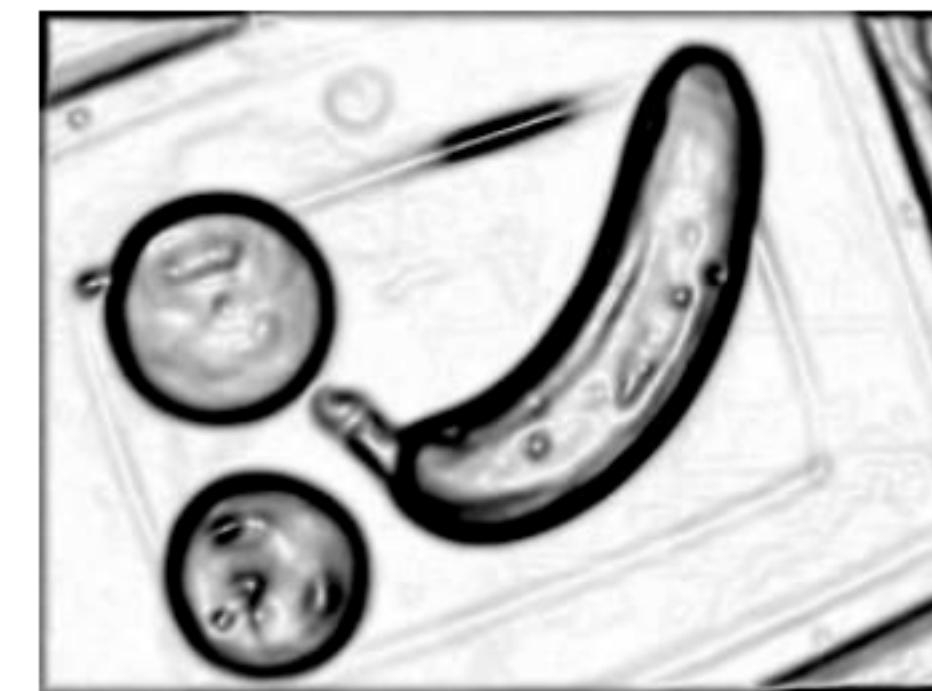
$$\nabla S = \nabla(G_\sigma \circledast I)$$

$$= \begin{bmatrix} \frac{\partial(G_\sigma \circledast I)}{\partial x} \\ \frac{\partial(G_\sigma \circledast I)}{\partial y} \end{bmatrix} = \begin{bmatrix} \frac{\partial G_\sigma}{\partial x} \circledast I \\ \frac{\partial G_\sigma}{\partial y} \circledast I \end{bmatrix}$$

## Example: fruity gradients



Original image



Edge strength  $|\nabla S|$

# 2D Edge Detection (step 3: NMS & step 4: thresholding)

## Step 3: Non-Maxima Suppression

The third stage of the edge detection algorithm is **Non-Maxima Suppression (NMS)**.

Edge elements, or **edgels**, are placed at locations where  $|\nabla S|$  is greater than local values of  $|\nabla S|$  in the directions  $\pm \nabla S$ . This aims to ensure that all **edgels** are located at ridge-points of the surface  $|\nabla S|$ .



Edge strength  $|\nabla S|$  after NMS

## Step 4: Thresholding

In the fourth and final step, the **edgels** are thresholded, so that only those with  $|\nabla S|$  above a certain value are retained.



Edge strength  $|\nabla S|$  after NMS and thresholding

# 2D Edge Detection (variations)

## Canny Edge Detection

The edge detection algorithm we have been describing is due to Canny (1986).

The output is a list of *edgel* positions, each with a strength  $|\nabla S|$  and an orientation  $\nabla S / |\nabla S|$ .

The Canny detector is a directional edge finder (both the gradient **magnitude** and **direction** are computed)

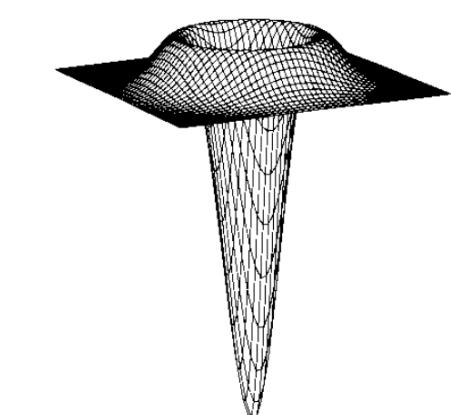
## Marr-Hildreth Edge Detection

An alternative approach to edge detection was developed by Marr and Hildreth (1980).

Unlike the directional Canny edge detector, the Marr-Hildreth operator is **isotropic**.

It finds zero-crossings of  $\nabla^2 G_\sigma \otimes I$ , where  $\nabla^2 G_\sigma$  is the Laplacian of  $G_\sigma$  (recall the Laplacian operator

$$\nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}.$$



$\nabla^2 G_\sigma$

## References:

Canny, J. A computational approach to edge detection. *TPAMI* 1986

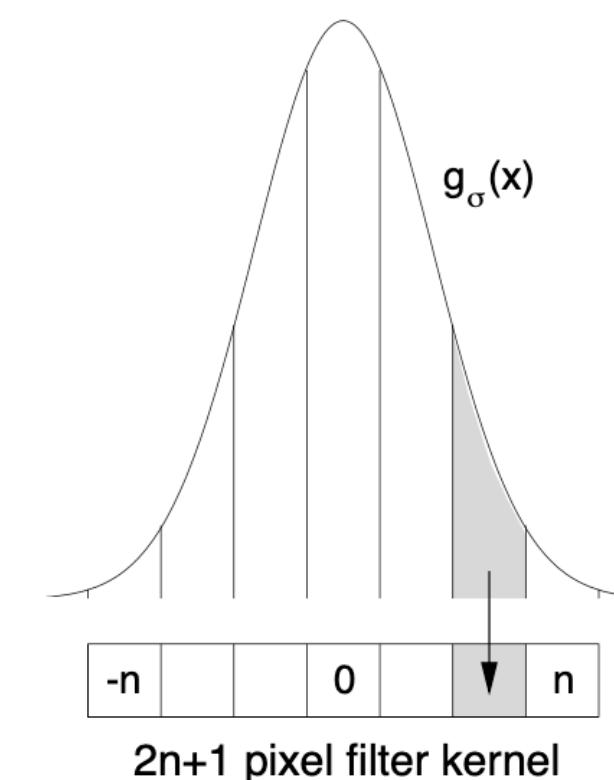
D. Marr and E. Hildreth, "Theory of edge detection." *Proc. of the Royal Society of London. Series B. Biological Sciences*, 1980

# Edge Detection: Implementation details

## Truncated summations

In practice, the image and filter kernels are discrete quantities and the convolutions are performed as **truncated summations**:

$$S(x, y) = \sum_{u=-n}^n \sum_{v=-n}^n G_\sigma(u, v) I(x - u, y - v)$$



## Truncation: how much is acceptable?

For acceptable accuracy, kernels are generally truncated so that the discarded samples are less than **1/1000** of the peak value.

$\sigma$	1.0	1.5	3	6
$2n + 1$	7	11	23	45

## Another computational trick!

The 2D convolutions would appear to be computationally expensive.

However, they can be decomposed into two 1D convolutions:

$$G_\sigma(x, y) \circledast I(x, y) = g_\sigma(x) \circledast [g_\sigma(y) \circledast I(x, y)]$$

The computational saving is:

$$\frac{(2n + 1)^2}{2(2n + 1)}$$

# Edge Detection: Implementation details

## Differentiation via convolution

Differentiation of the smoothed image is also implemented with a **discrete convolution**.

By considering the Taylor-series expansion of  $S(x, y)$  one can show that a simple finite-difference approximation to the first-order spatial derivative of  $S(x, y)$  with respect to  $x$  is given by:

$$\frac{\partial S}{\partial x} = \frac{S(x+1, y) - S(x-1, y)}{2}$$

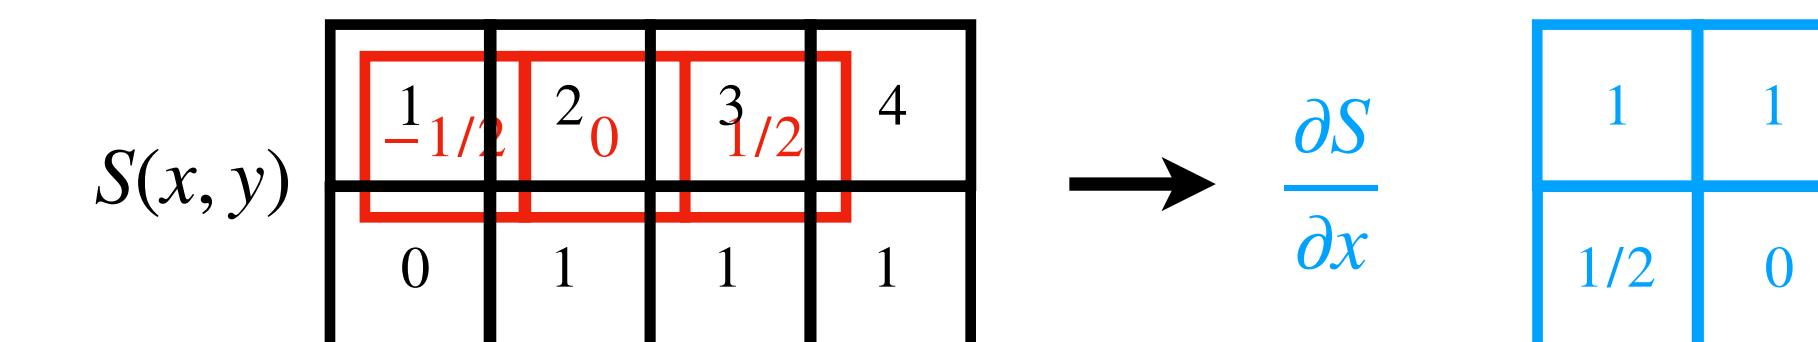
VE

## Implementing first order derivatives

We can calculate the finite-difference approximation to  $\partial S / \partial x$  by **convolving** the rows of smoothed image samples,  $S(x, y)$ , with the 3-element kernel:

$$\begin{array}{|c|c|c|} \hline 1/2 & 0 & -1/2 \\ \hline \end{array}$$

Recall that when convolving, we flip the kernel and sum the element-wise products under each kernel position:



**Note:** this is often called "valid" convolution (the kernel is not allowed to run off the edges of  $S(x, y)$ ).

**End of Image Structures Lecture 1:  
time for questions**

# Express 2D Gaussian Convolution as 1D Gaussian convolutions

Recall the familiar 2D Gaussian:  $G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$

**Goal:** show that convolution with this 2D Gaussian can be performed via two 1D convolutions:  $G_\sigma(x, y) \circledast I(x, y) = g_\sigma(x) \circledast [g_\sigma(y) \circledast I(x, y)]$

**Key idea:**  $G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{x^2}{2\sigma^2}\right) \cdot \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{y^2}{2\sigma^2}\right)$

We can plug this in to the LHS of our **goal** equation:

$$\begin{aligned} G_\sigma(x, y) \circledast I(x, y) &= \frac{1}{2\pi\sigma^2} \int_u \int_v I(x - u, y - v) \cdot \exp\left(-\frac{u^2 + v^2}{2\sigma^2}\right) du dv \\ &= \frac{1}{\sqrt{2\pi\sigma^2}} \int_u \exp\left(-\frac{u^2}{2\sigma^2}\right) \frac{1}{\sqrt{2\pi\sigma^2}} \int_v I(x - u, y - v) \cdot \exp\left(-\frac{v^2}{2\sigma^2}\right) dv du \quad (\text{using the key idea}) \\ &= \frac{1}{\sqrt{2\pi\sigma^2}} \int_u \exp\left(-\frac{u^2}{2\sigma^2}\right) [g_\sigma(y) \circledast I(x - u, y)] du \quad = g_\sigma(x) \circledast [g_\sigma(y) \circledast I(x, y)] \end{aligned}$$

Note: "separable" convolutions are very popular in deep learning architectures to save compute

# Show that Marr-Hildreth is isotropic

To show that the Marr-Hildreth edge detector is isotropic (uniform in all directions), we can use a **change of coordinates**.

The edge detector is the **Laplacian of Gaussian**:  $\nabla^2 G_\sigma = \left( \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) \left[ \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right) \right]$

First derivative

$$\frac{\partial}{\partial x} \left[ \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right) \right] = \frac{1}{2\pi\sigma^2} \left( \frac{-x}{\sigma^2} \right) \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)$$

Second derivative

$$\frac{\partial^2}{\partial x^2} \left[ \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right) \right] = \frac{1}{2\pi\sigma^2} \left( \frac{-1}{\sigma^2} \right) \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right) + \frac{1}{2\pi\sigma^2} \left( \frac{x^2}{\sigma^4} \right) \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)$$

$$\nabla^2 G_\sigma = \frac{1}{2\pi\sigma^2} \left( \frac{x^2}{\sigma^4} + \frac{y^2}{\sigma^4} - \frac{2}{\sigma^2} \right) \cdot \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right) \text{ (by plugging in the second derivatives for both } x \text{ and } y)$$

$$= \frac{1}{2\pi\sigma^2} \left( \frac{r^2}{\sigma^4} - \frac{2}{\sigma^2} \right) \cdot \exp\left(-\frac{r^2}{2\sigma^2}\right) \text{ by converting to } \underline{\text{polar coordinates}} \text{ with } x = r \cos \theta, y = r \sin \theta \text{ (with } r^2 = x^2 + y^2)$$

We can see that  $\nabla^2 G_\sigma$  has no dependence on direction  $\theta$ , so it must be isotropic!