

Deep Learning 4

Lecture 15: Neural Network Design, Scaling Laws and Transformers
4F12: Computer Vision

Instructor: Samuel Albanie

Outline for final lectures

Outline

For the last two lectures, we will focus on:

- Strategies for Neural Network Design This lecture
- Scaling phenomena
- Transformers

- Self-supervised learning Next lecture
- Pseudo-labelling

Qualitatively, these lectures:

- Are more *research* focused
- Prioritise *breadth* over *depth*

Strategies for Neural Network Design

Background

Modern deep learning stems from the **connectionist** approach, in which the **wiring of computational networks** plays an important role in building intelligent machines.

Conceptually, it can be helpful to categorise the structures that define the wiring between neural network units into two categories:

- **Network architecture** - connections between units that are (typically) fixed throughout training (e.g. operation types)
- **Network parameters** - connections between units that are updated during training (e.g. kernel weights learned via backpropagation)

Neural Network Design focuses principally on the finding good **network architectures** (although the distinction between the architecture and the parameters can be somewhat blurry).

Goals

We exist (probably) in a resource-limited environment. We have *limited supplies* of:

Energy

Computation

Memory

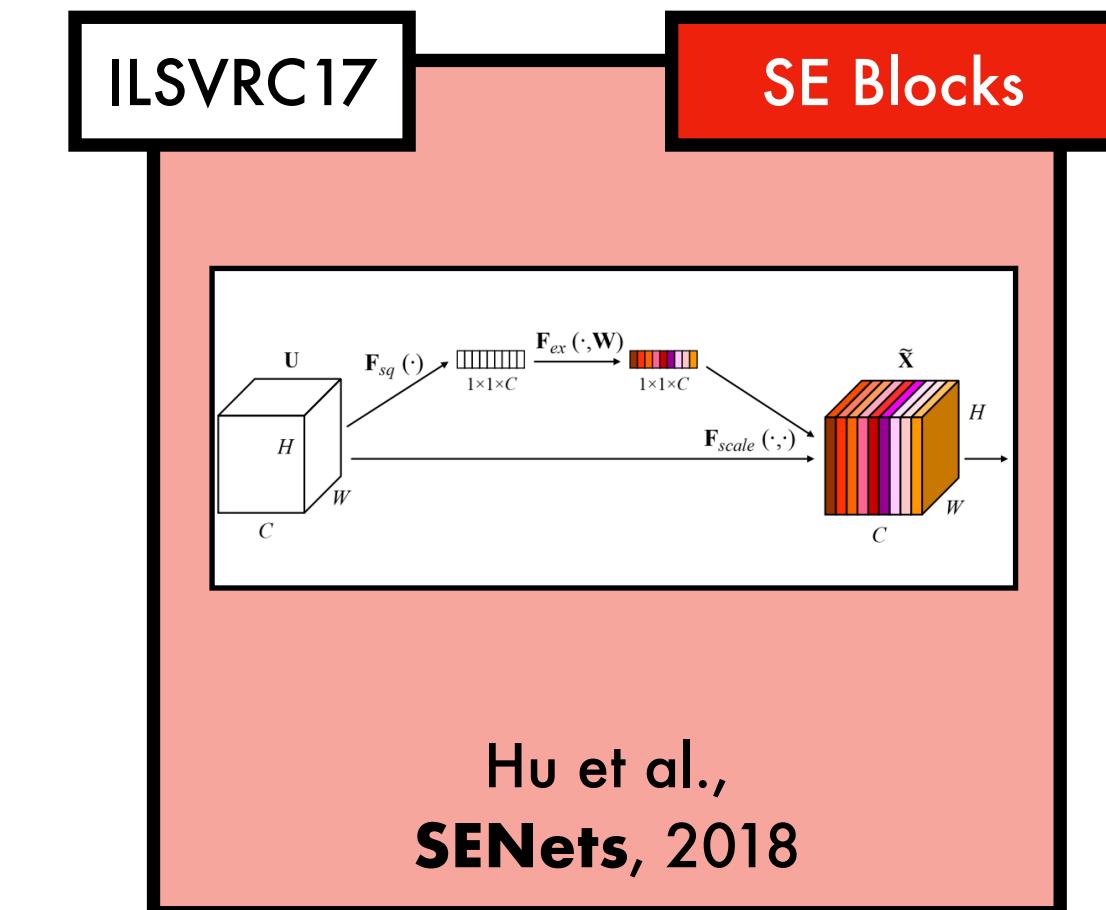
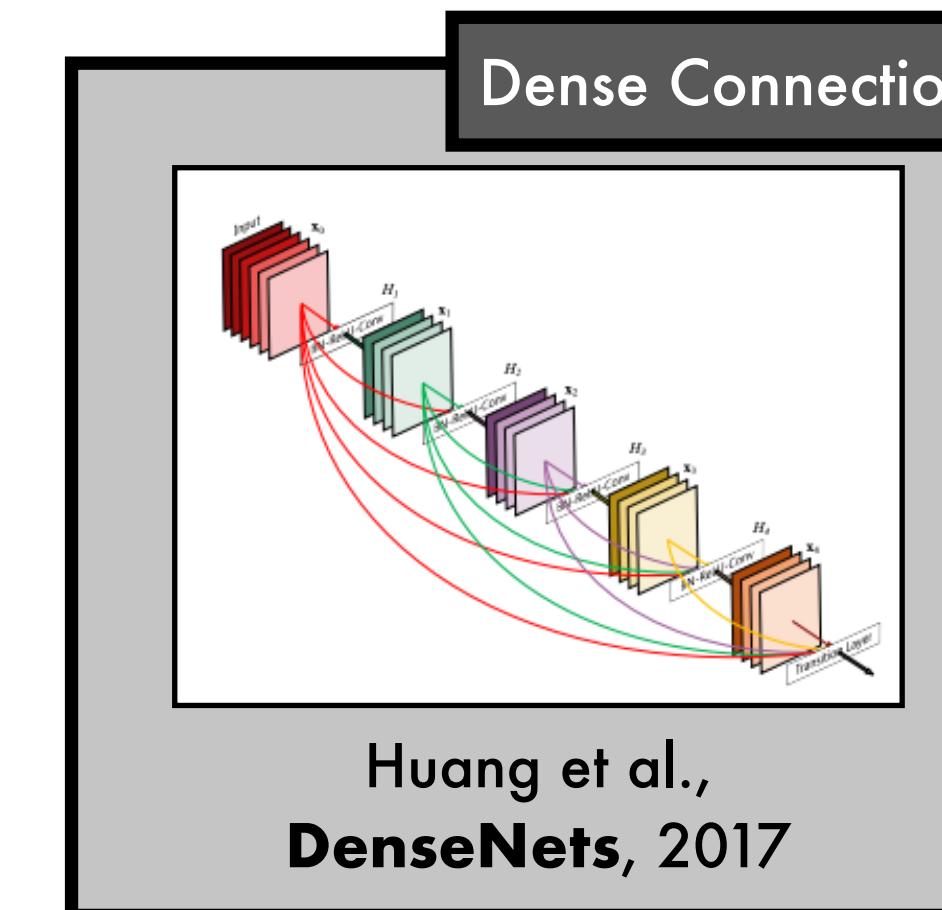
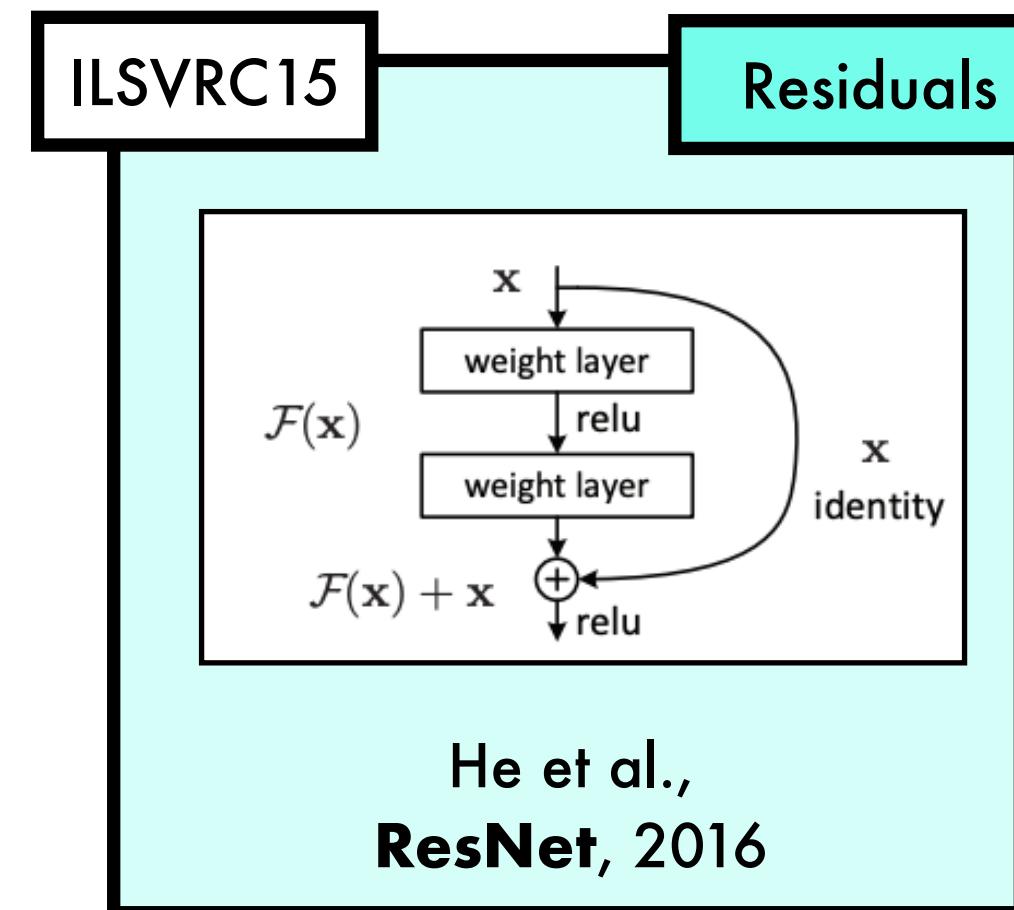
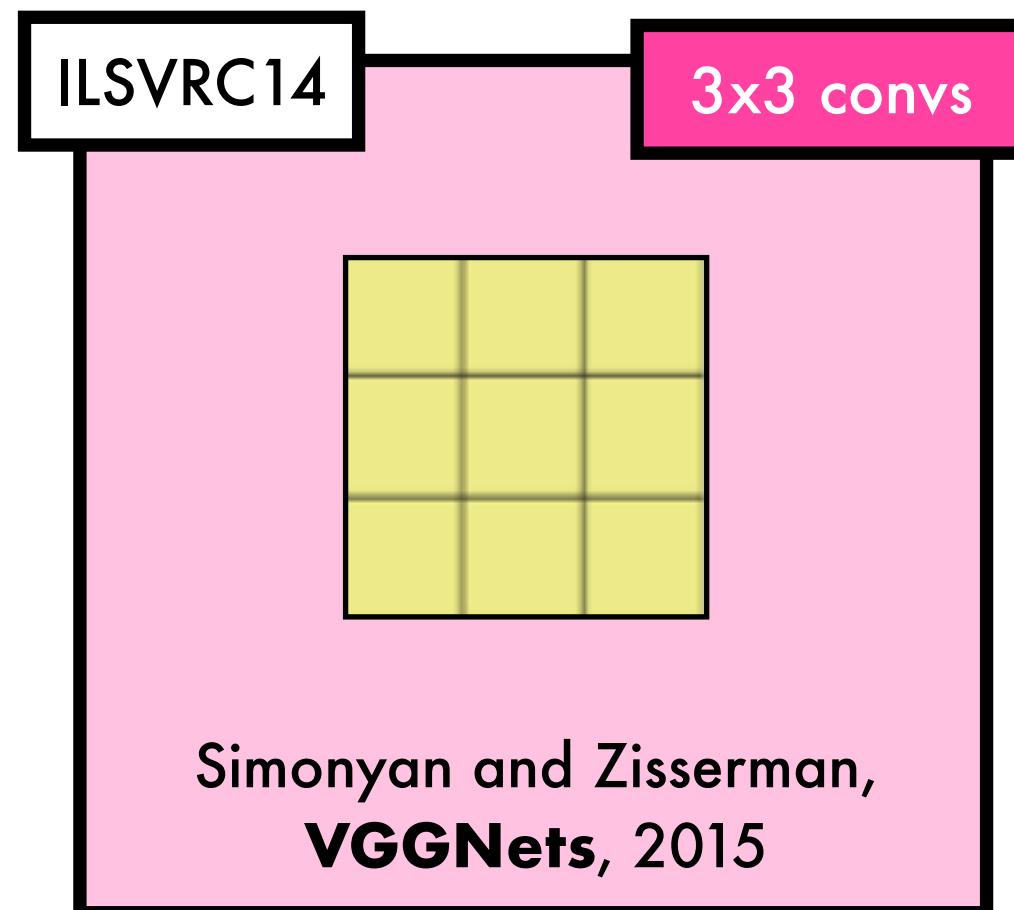
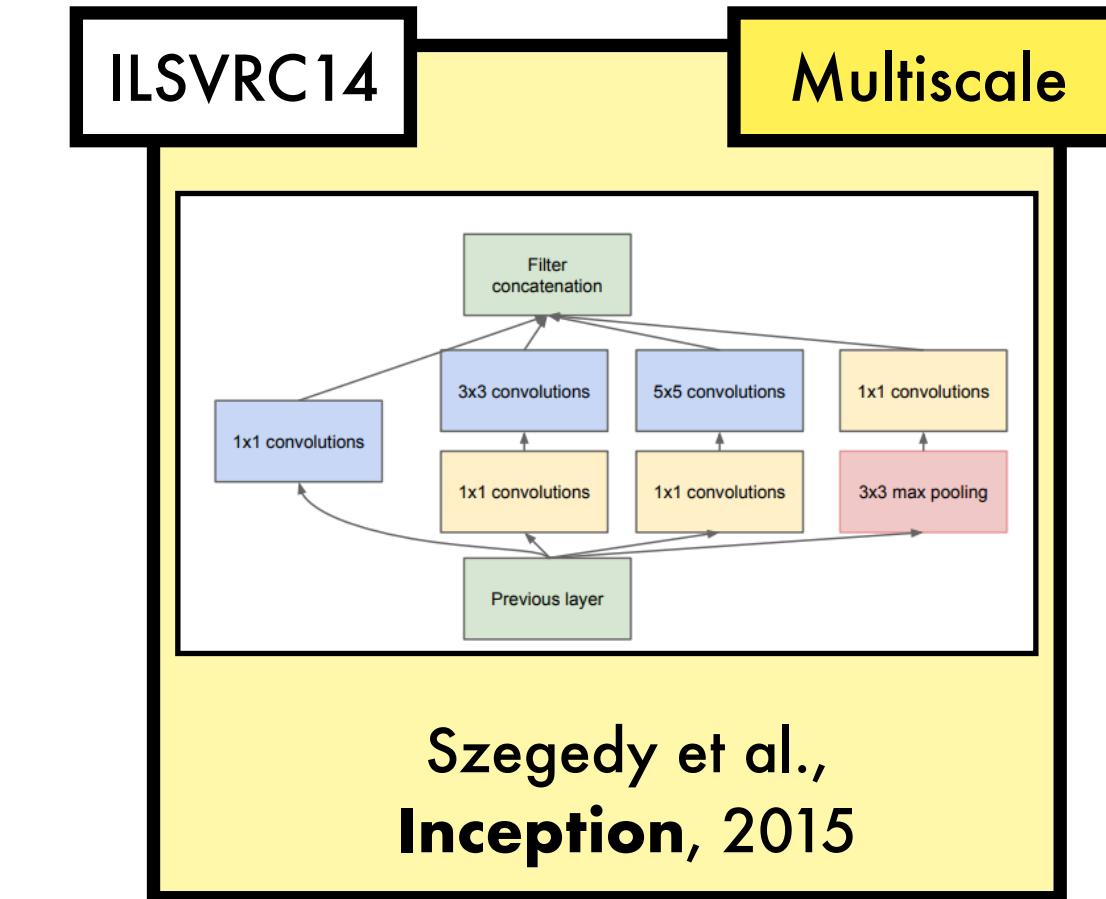
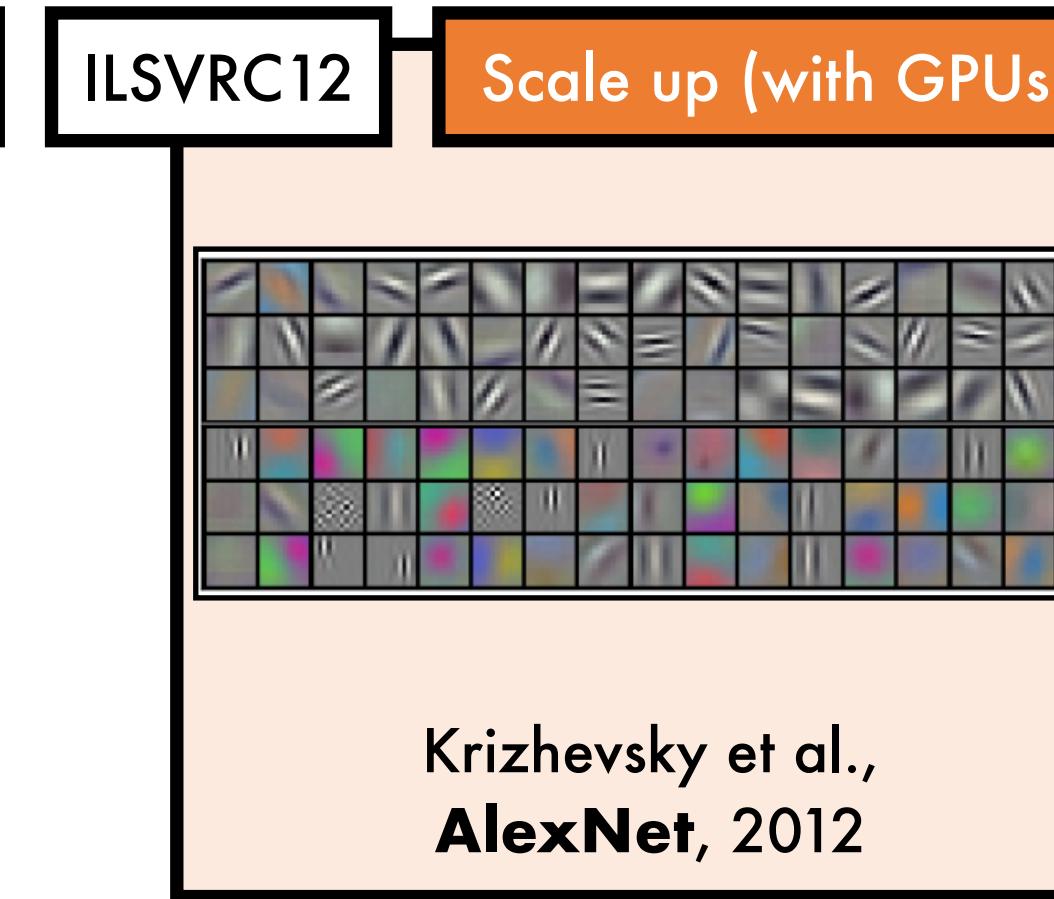
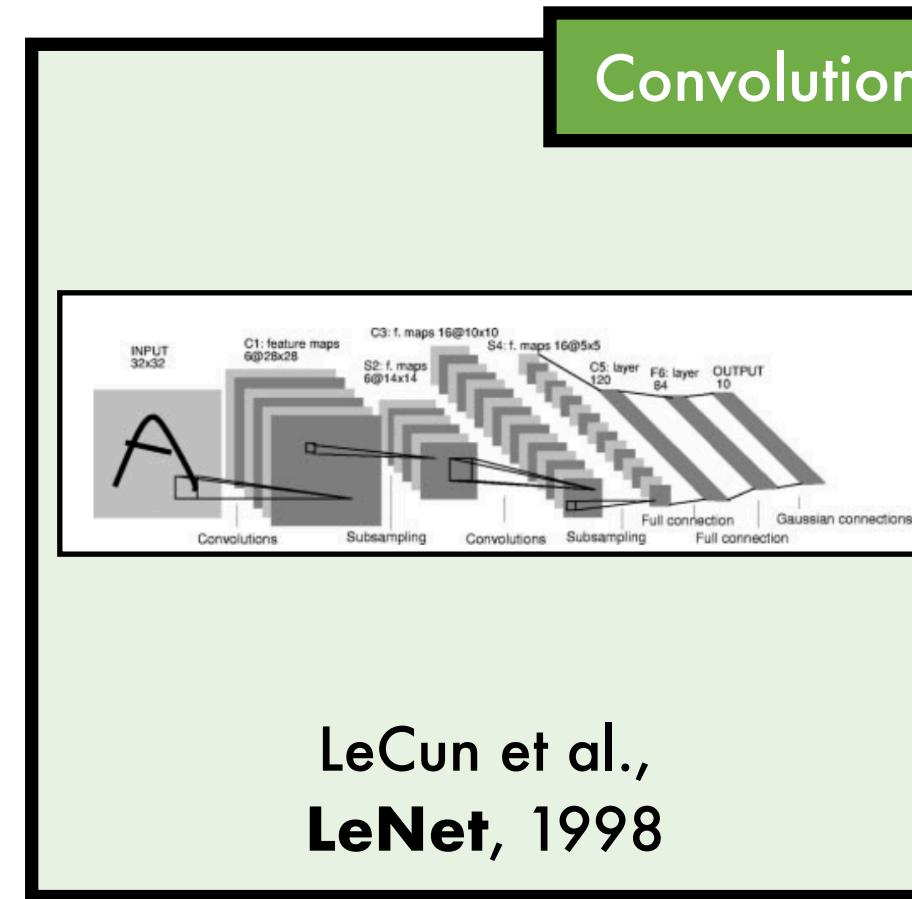
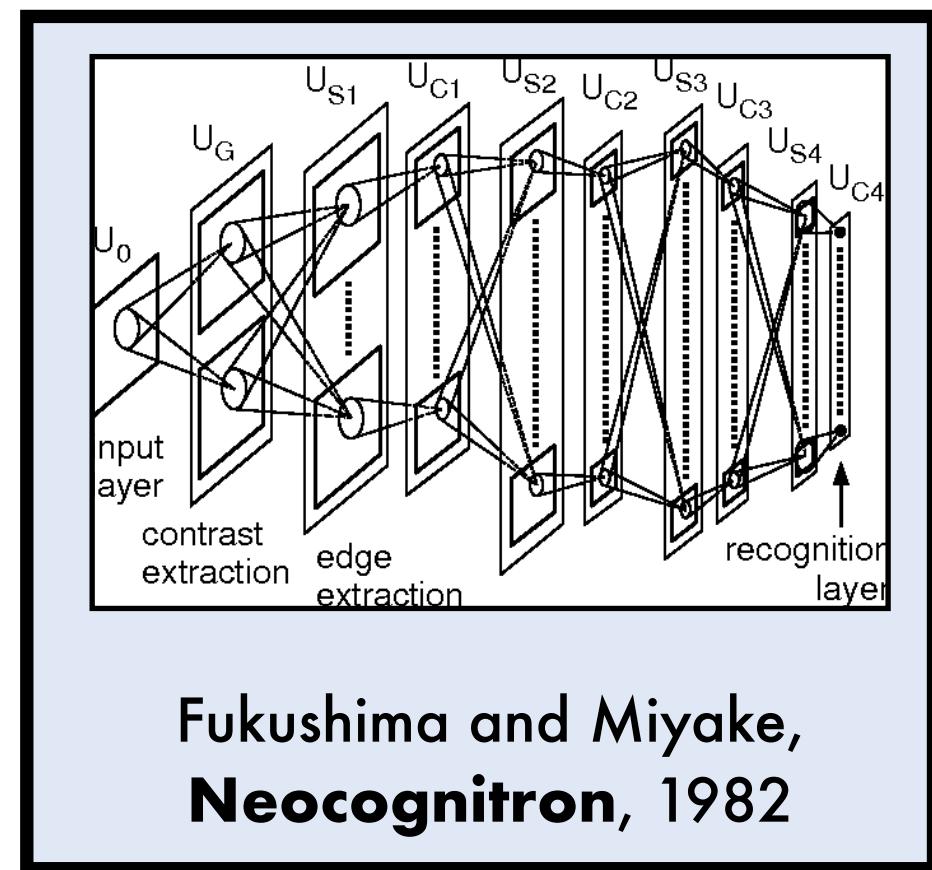
Time

For any given task, neural network design aims to produce architectures with:

Greater task-specific performance (e.g. accuracy)

Lower resource burden

Strategy 1: Neural Network Design by Hand



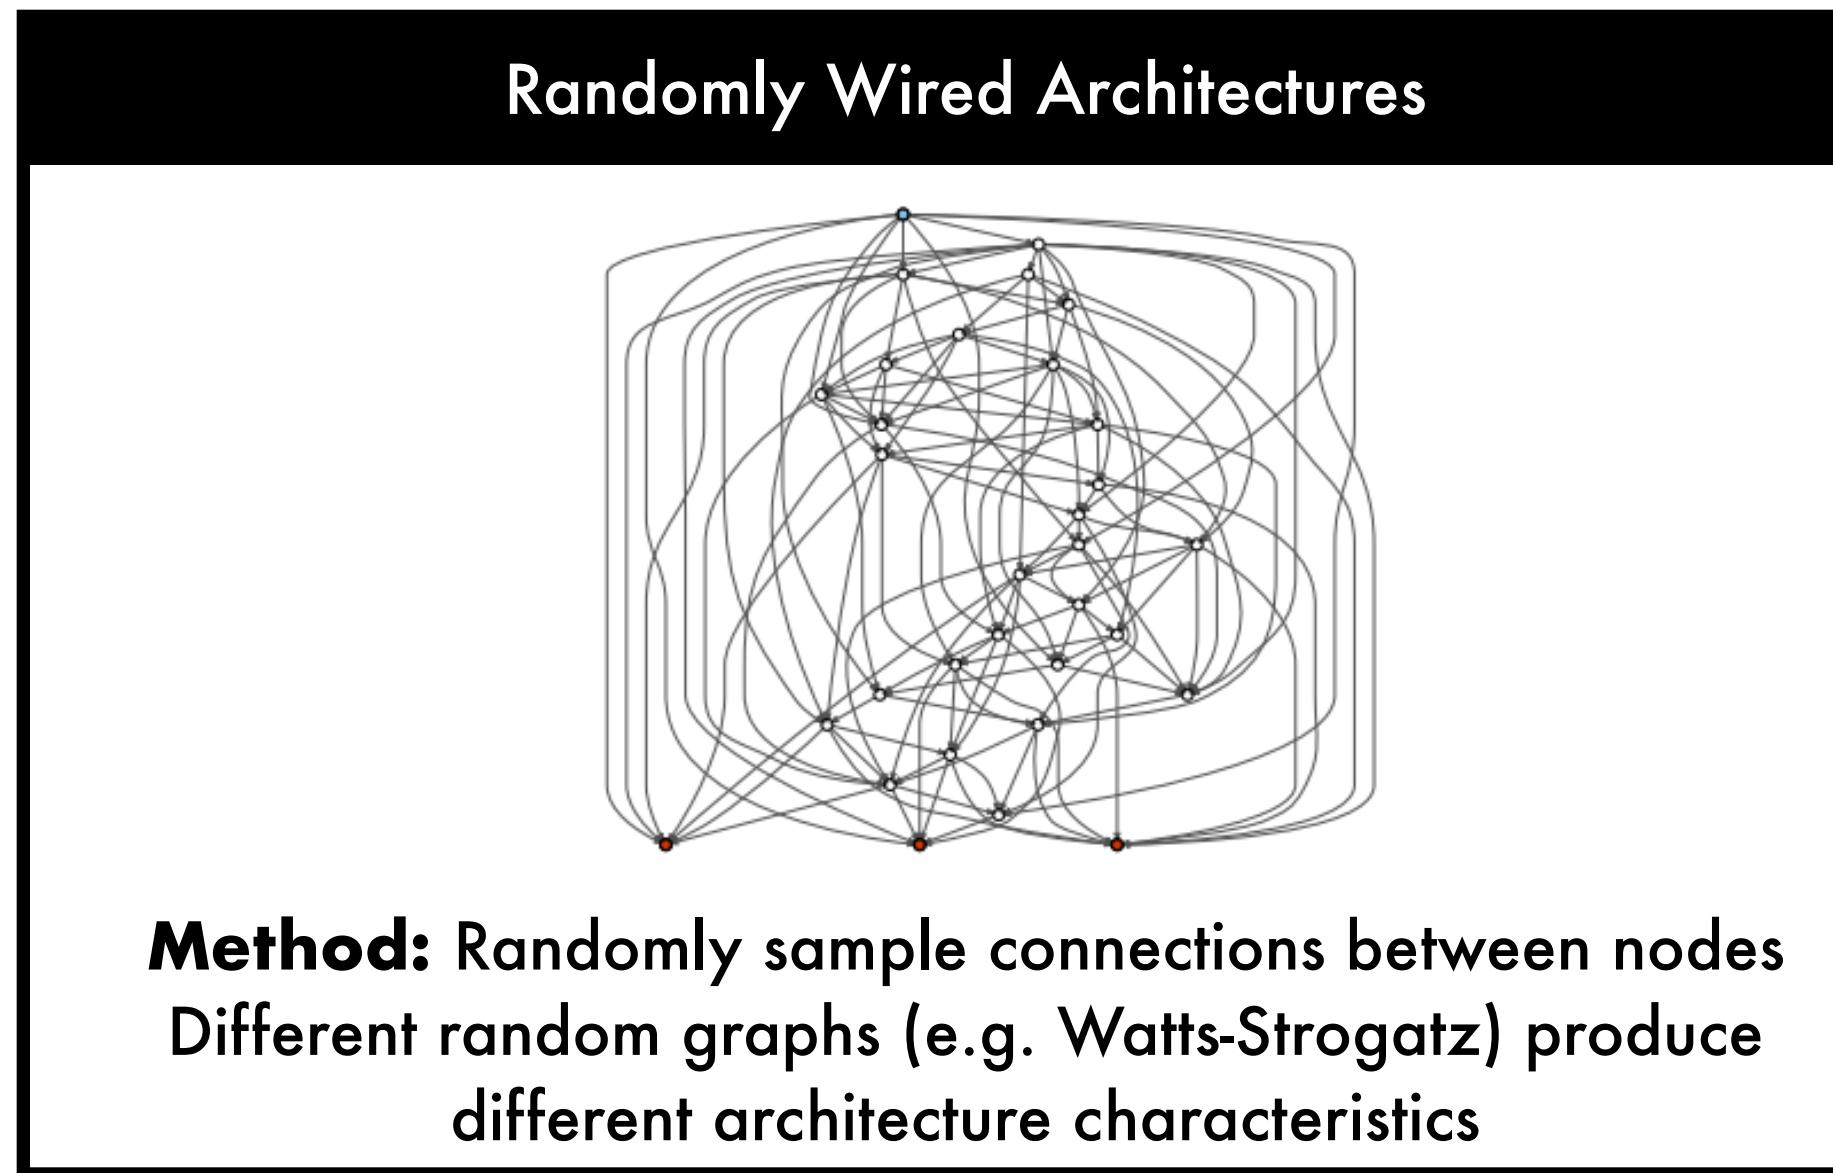
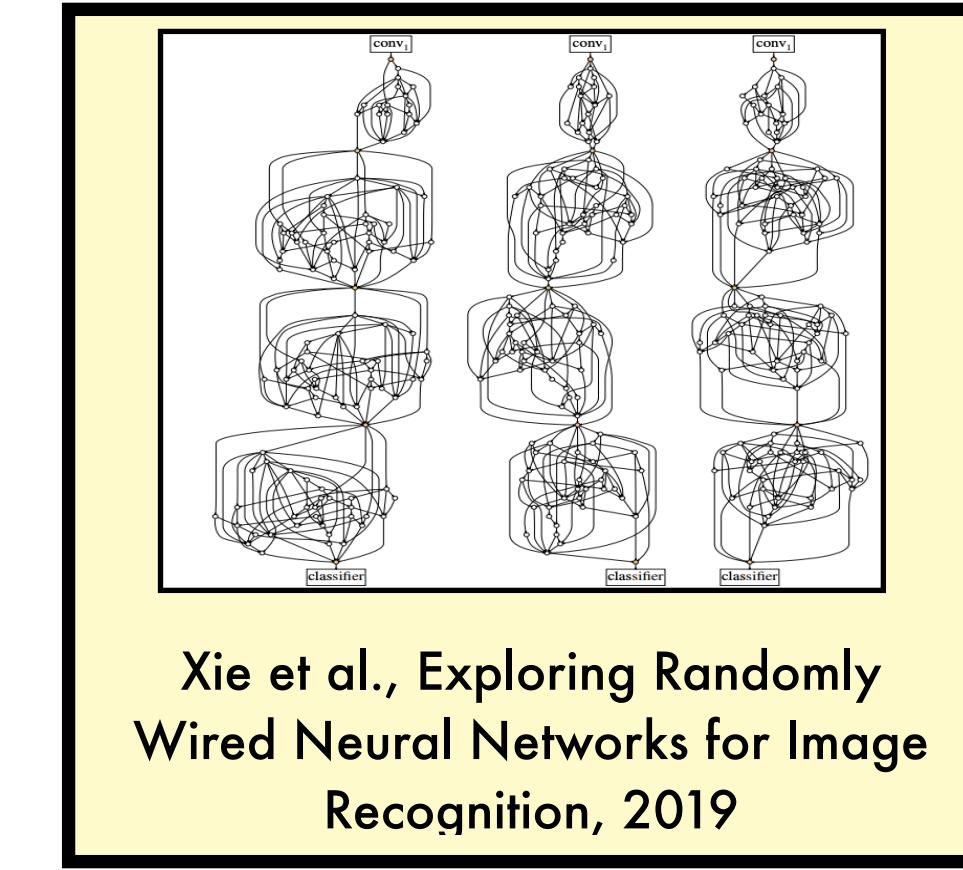
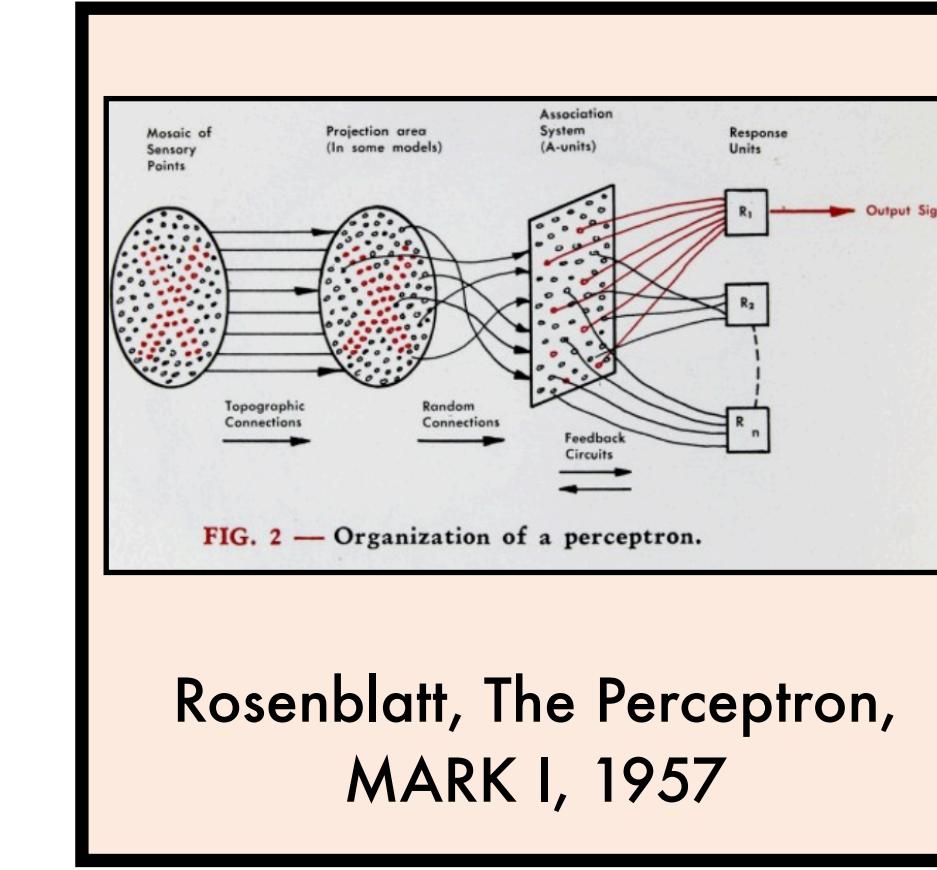
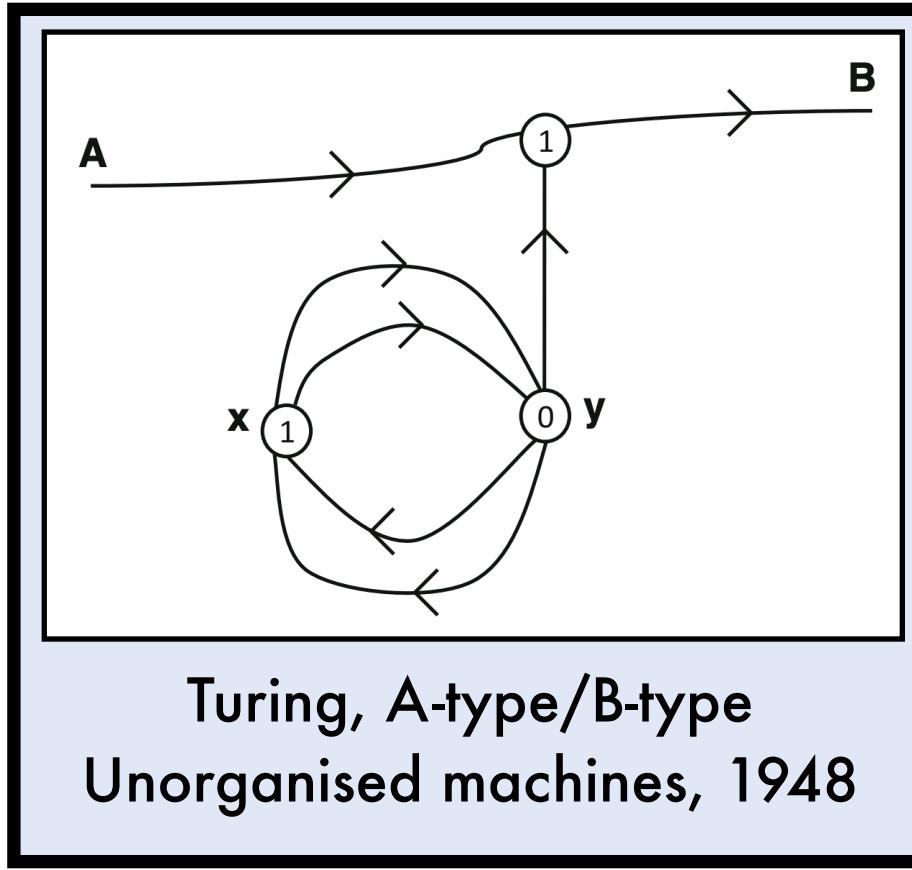
Aside: several of these architectures rose to prominence through strong performance on the ImageNet ILSVRC competition.

References:

- Fukushima and Miyake. "Neocognitron" CCNN, 1982
- LeCun, Y. et al. (1998). Gradient-based learning applied to document recognition. *IEEE*
- Krizhevsky, A et al. "Imagenet classification with deep CNNs." NeurIPS. 2012.
- Szegedy, C et al. (2015). Going deeper with convolutions. CVPR

- Simonyan et al., (2015). Very Deep Convolutional Networks for Large-Scale Image Recognition. ICLR
- He, Kaiming, et al. "Deep residual learning for image recognition." CVPR. 2016.
- Huang, Gao, et al. "Densely connected convolutional networks." CVPR 2017.
- Hu et al. "Squeeze-and-Excitation Networks." CVPR 2018
- Russakovsky et al. "ImageNet Large Scale Visual Recognition Challenge." IJCV 2015

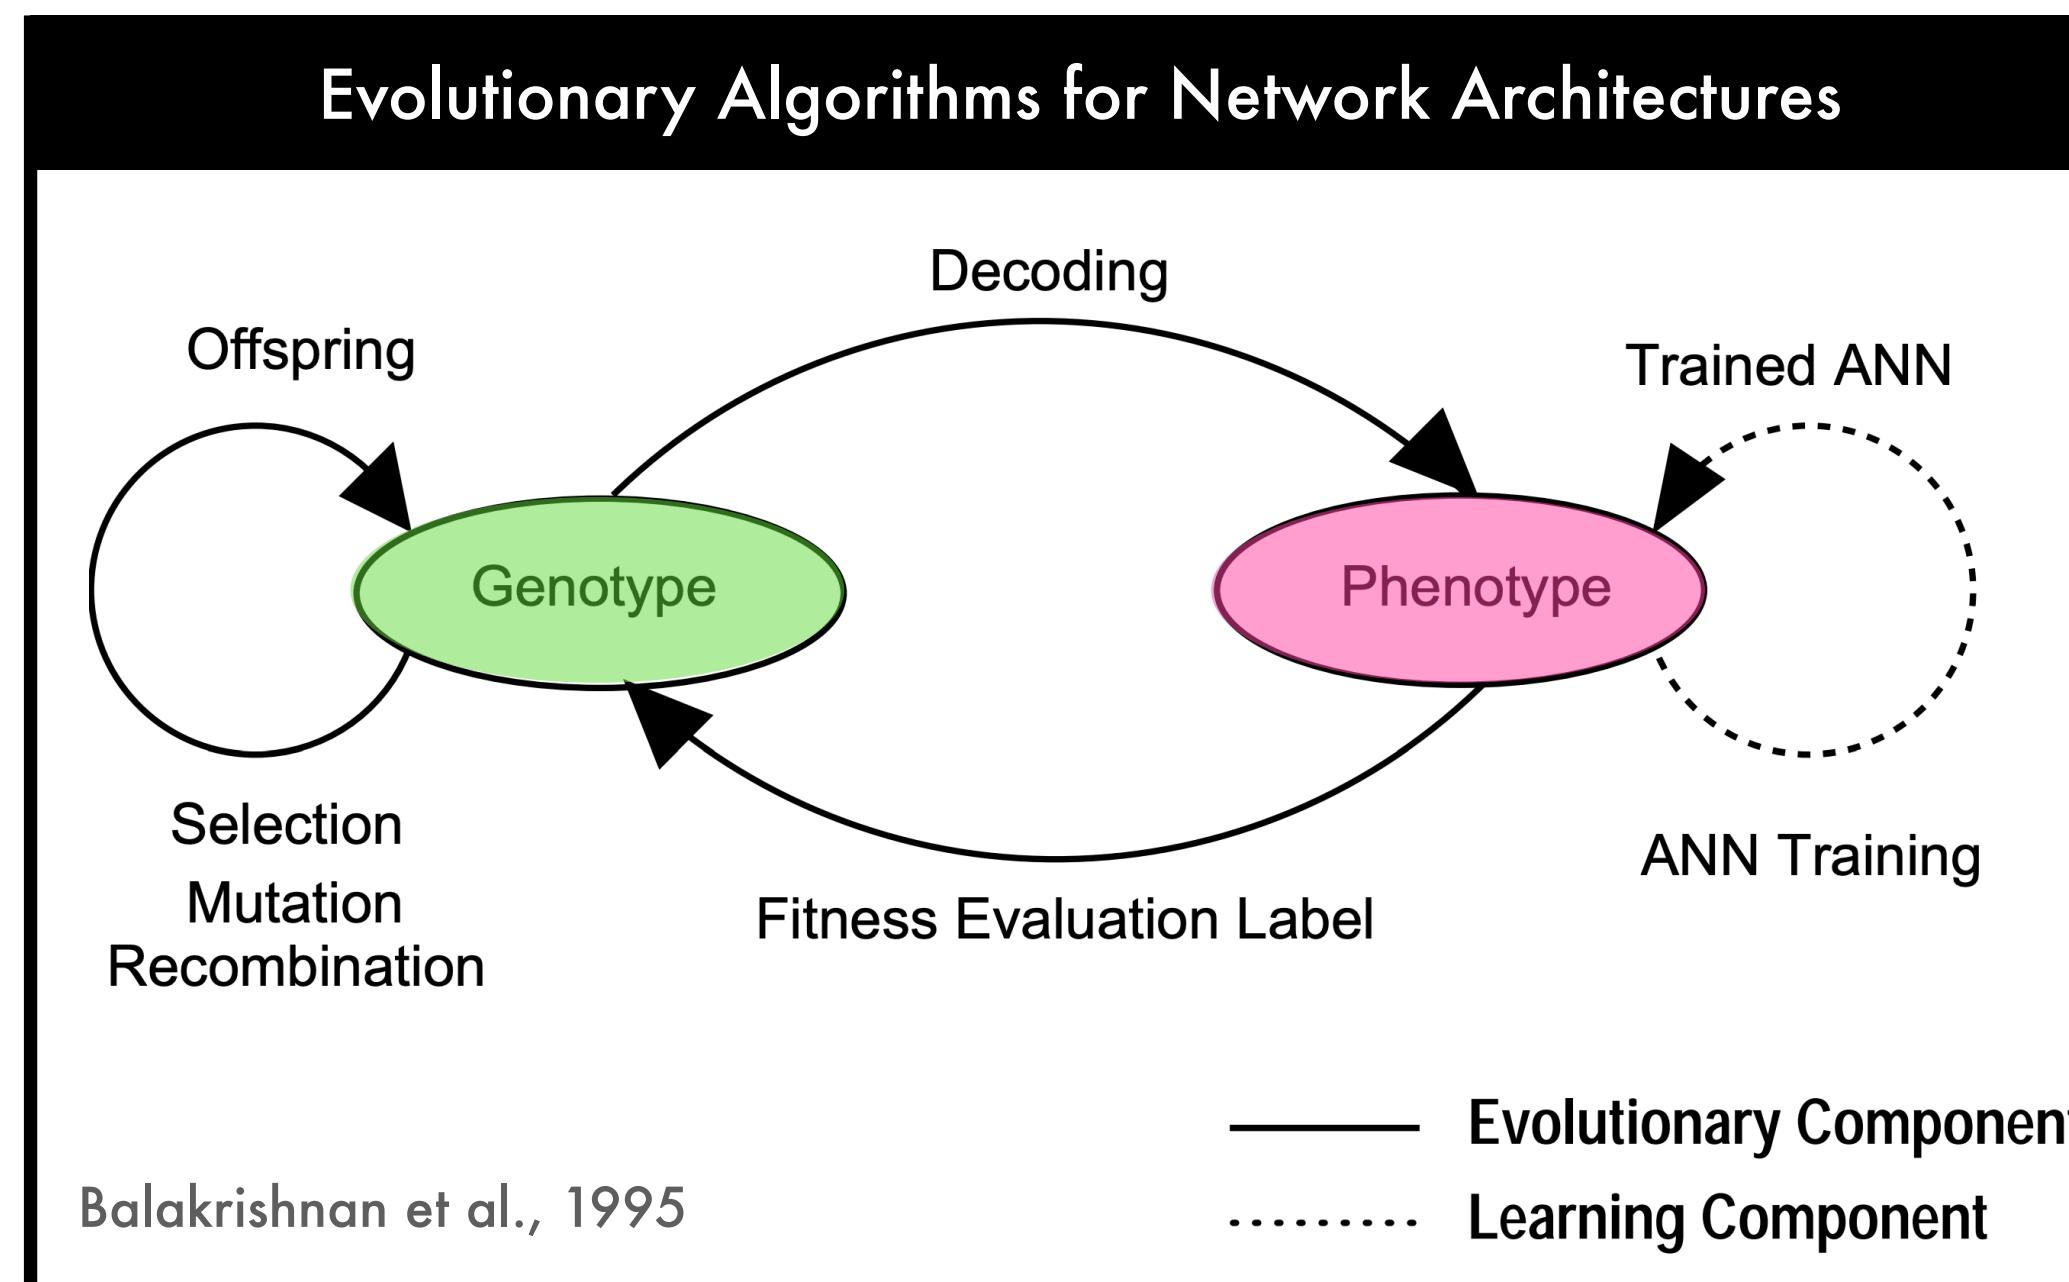
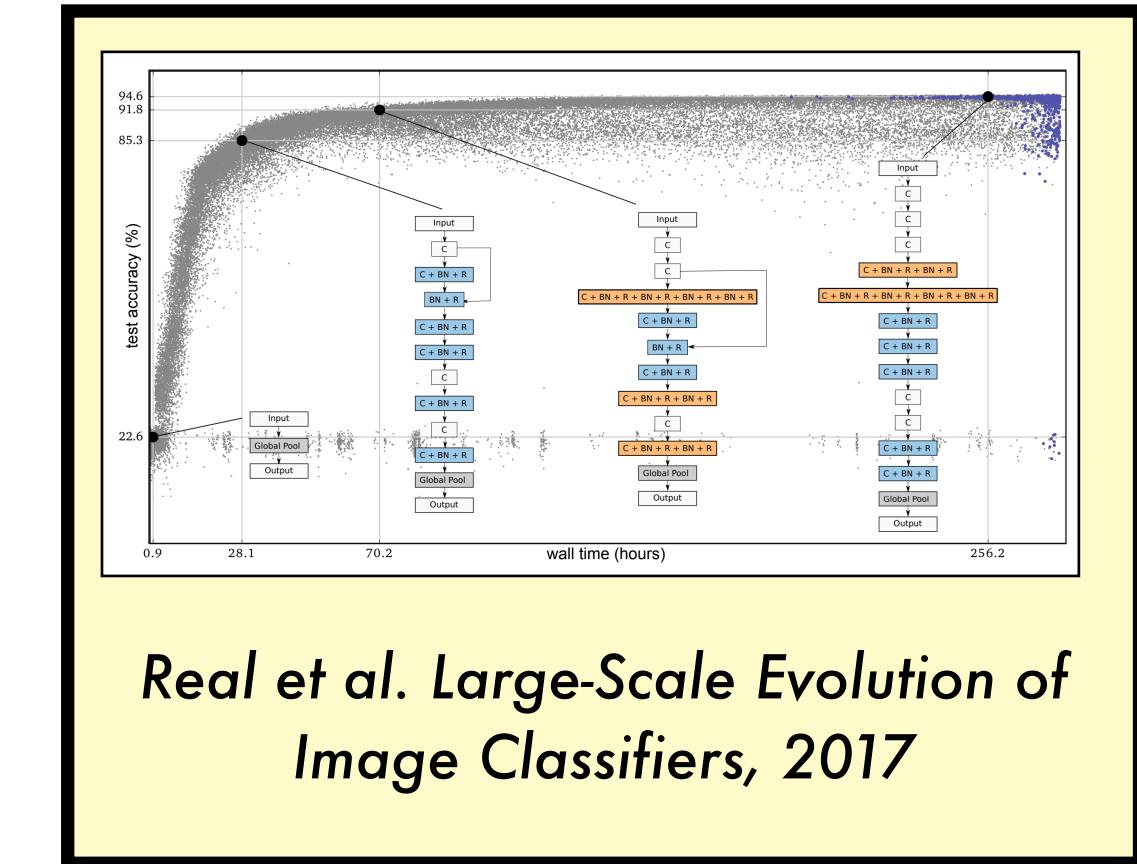
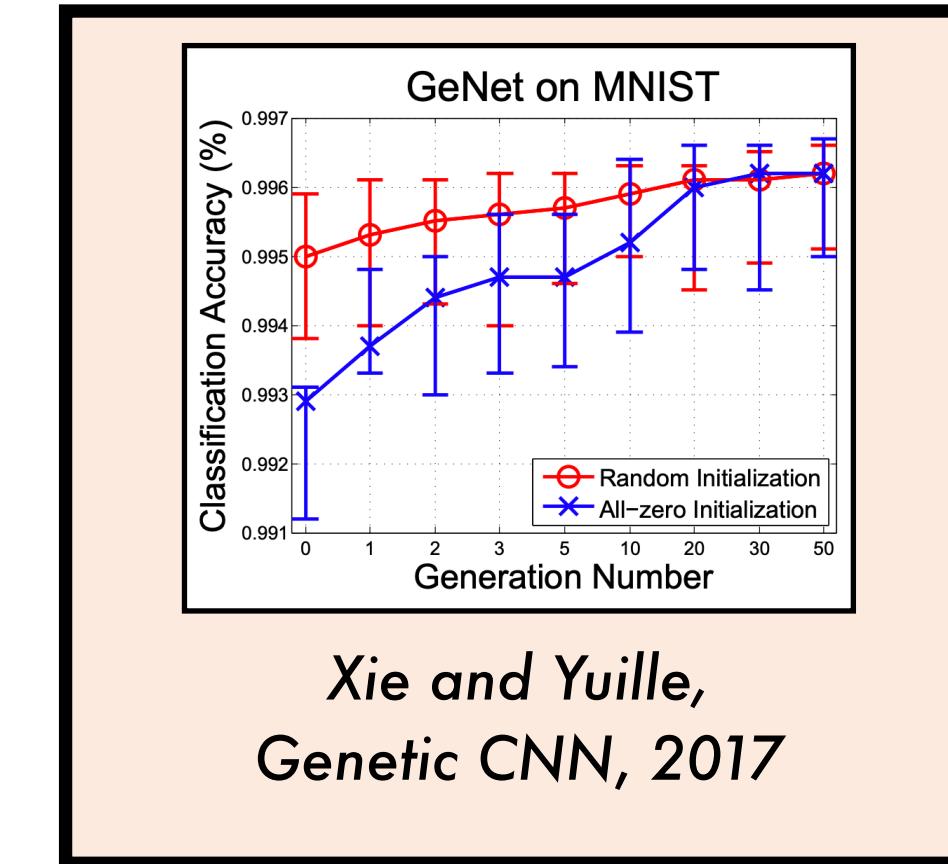
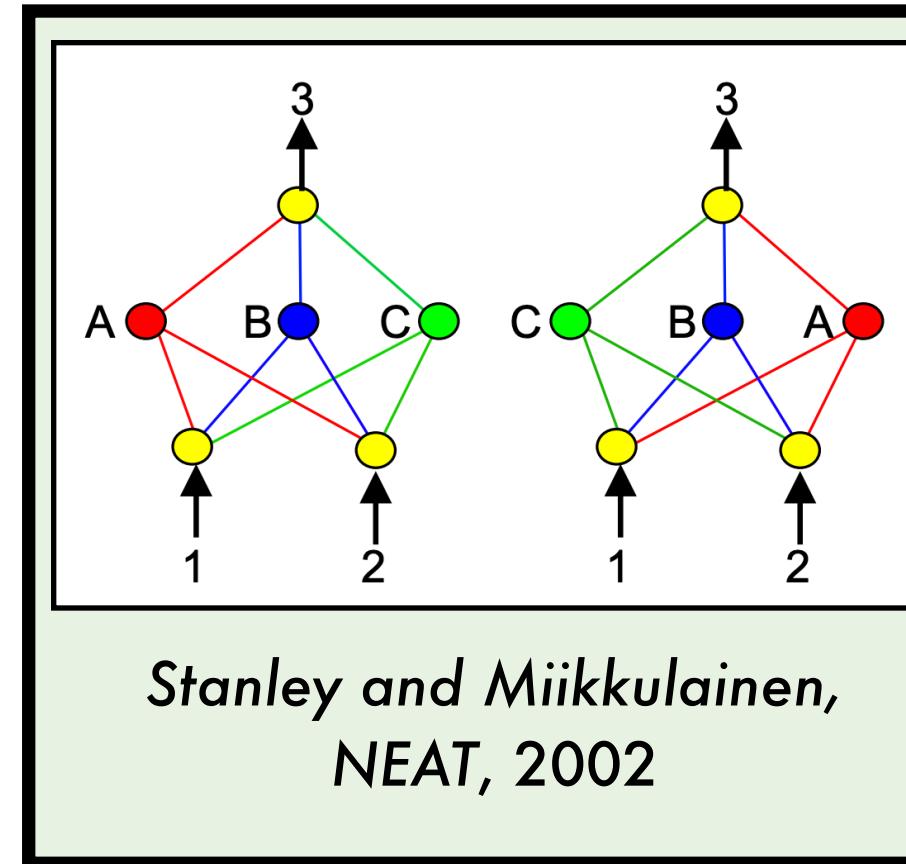
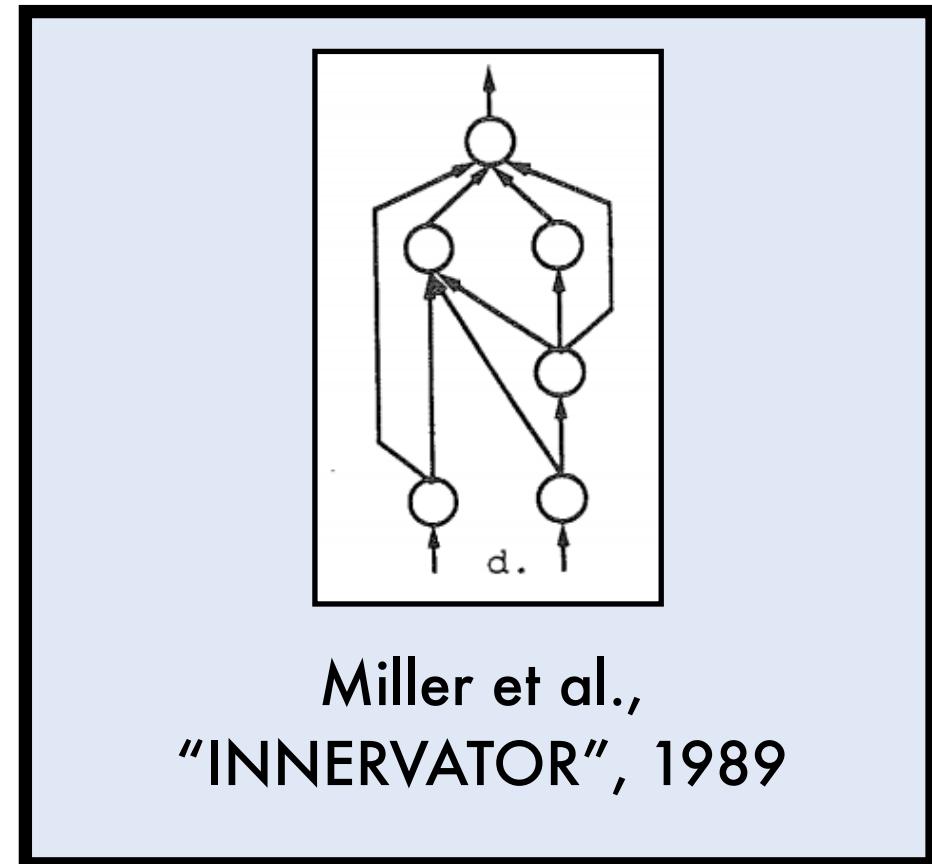
Strategy 2: Random Wiring



References:

- Watts and Strogatz. "Collective dynamics of 'small-world' networks." *Nature* 393 (1998): 440-442.
Turing, A. M. (1948). Intelligent machinery.
(Figure) Russell, S., & Norvig, P. (2002). *Artificial intelligence: a modern approach*.
Rosenblatt, F. (1957). *The perceptron, a perceiving and recognizing automaton Project Para*. Cornell Aeronautical Laboratory.
Xie, Saining, et al. "Exploring randomly wired neural networks for image recognition." CVPR. 2019.
Girshick, (2019) <https://neuralarchitects.org/slides/girshick-slides.pdf>

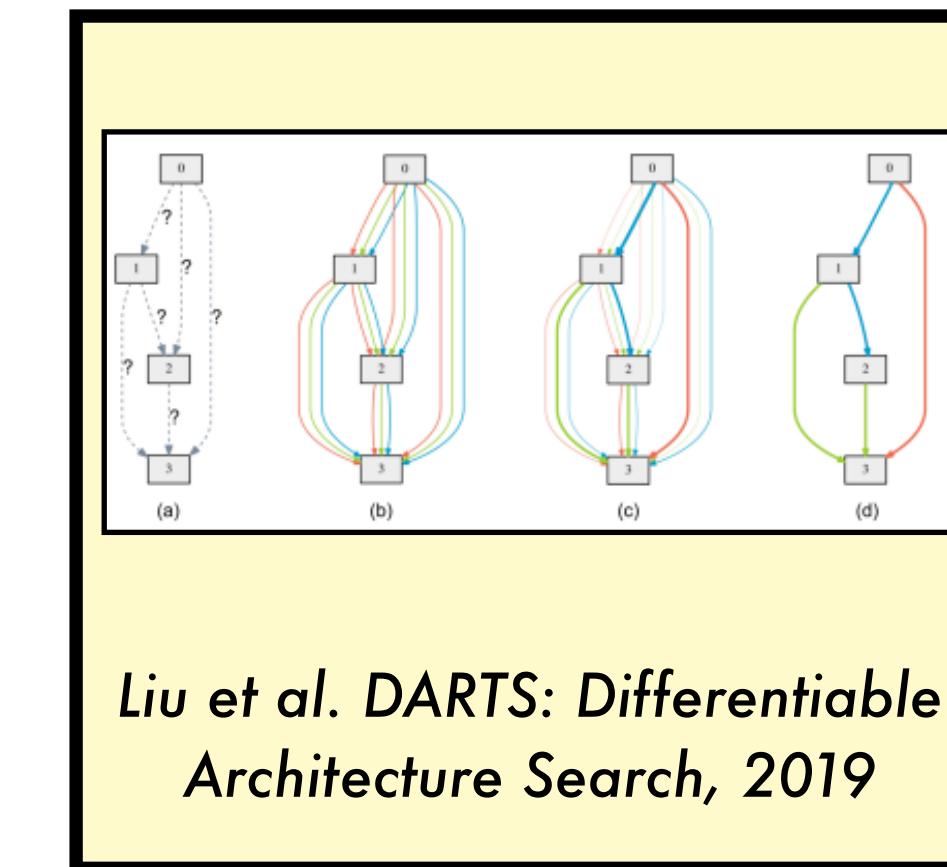
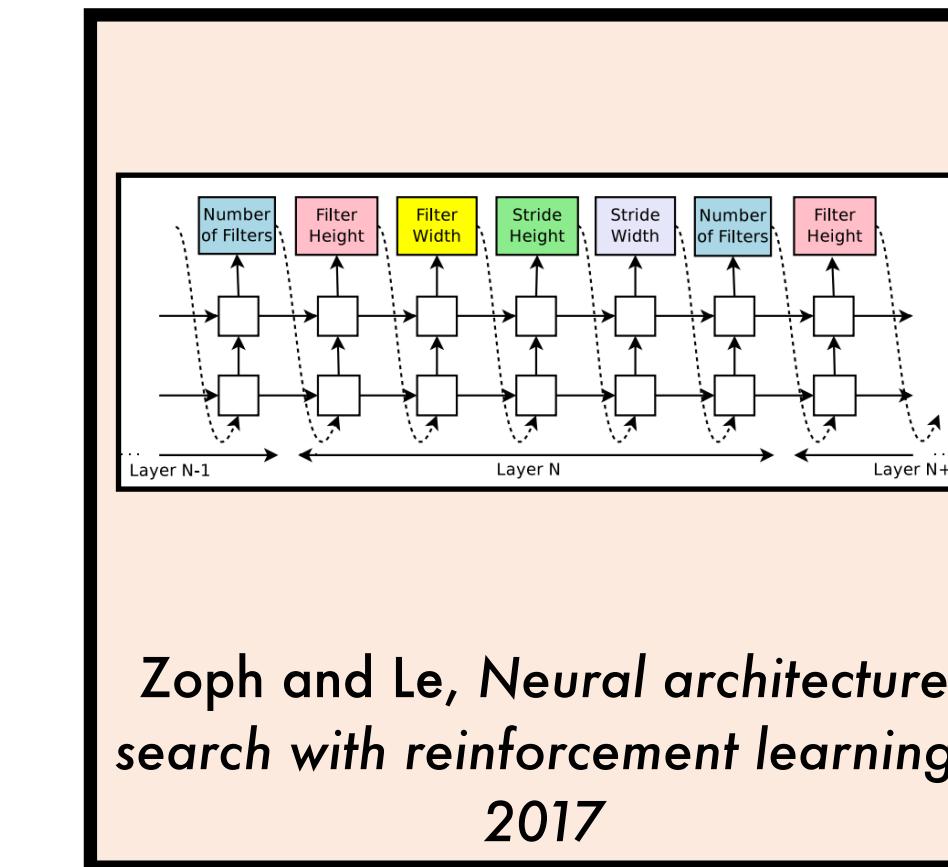
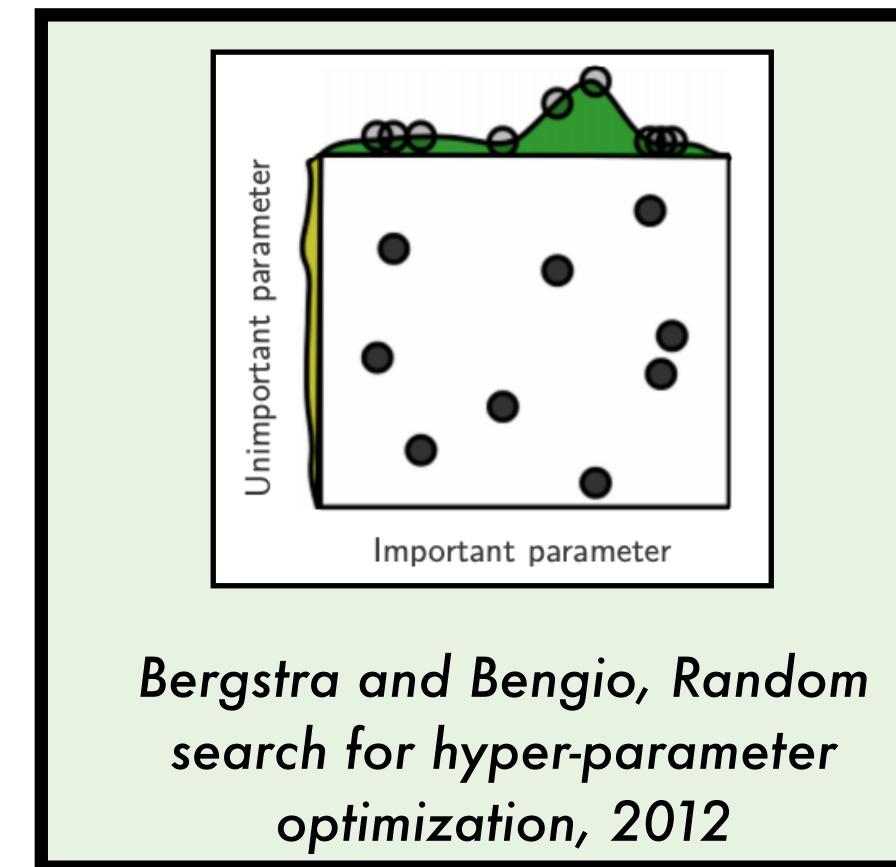
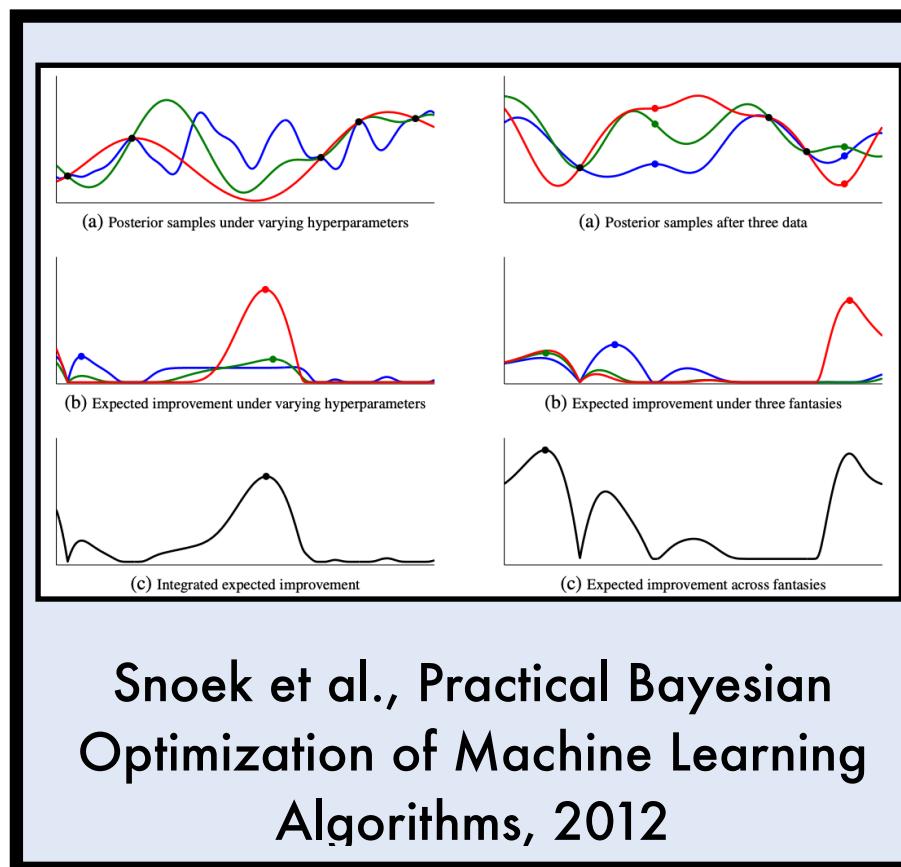
Strategy 3: Evolutionary Algorithms



References:

- Figure sourced from Evolutionary Design of Neural Architectures – A Preliminary Taxonomy and Guide to Literature, Balakrishnan et al., 1995
- Rechenberg I (1965) Cybernetic solution path of an experimental problem. Royal Aircraft Establishment
- J. Holland, Adaptation in natural and artificial systems, 1975
- P. M Todd. Evolutionary methods for connectionist architectures. Unpublished manuscript, 1988.
- Miller et al. Designing neural networks using genetic algorithms. In ICGA, 1989.
- Stanley et al. (2002). Evolving neural networks through augmenting topologies. *Evolutionary computation*.
- Bayer, Justin, et al. "Evolving memory cell structures for sequence learning." ICANN, 2009.
- Xie, Lingxi and Alan Loddon Yuille. "Genetic CNN." ICCV 2017
- R. Jozefowicz, et al.. "An empirical exploration of recurrent network architectures." ICML. 2015.
- Real, Esteban, et al. "Large-scale evolution of image classifiers." ICML 2017

Strategy 4: Neural Architecture Search



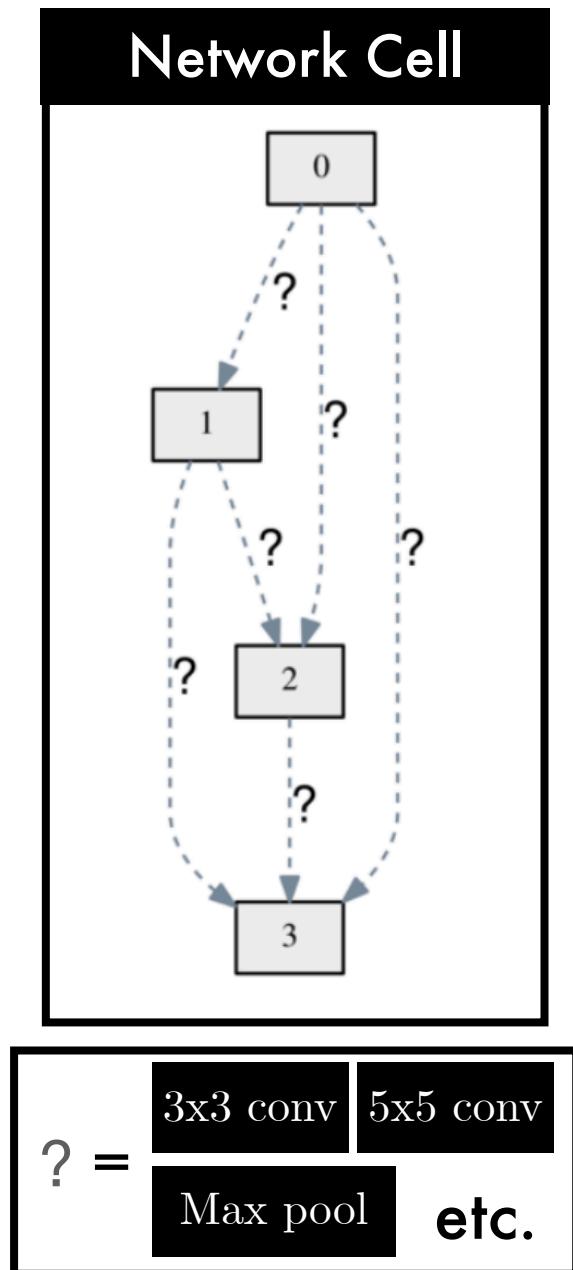
References:

- Snoek, Jasper, Hugo Larochelle, and Ryan P. Adams. "Practical bayesian optimization of machine learning algorithms." *Advances in neural information processing systems*. 2012.
- Bergstra, James, and Yoshua Bengio. "Random search for hyper-parameter optimization." *JMLR* (2012)
- Zoph, B., & Le, Q. V. (2016). Neural architecture search with reinforcement learning. *arXiv preprint arXiv:1611.01578*.
- Baker, Bowen, et al. "Accelerating neural architecture search using performance prediction." *arXiv preprint arXiv:1705.10823* (2017).
- Liu, Hanxiao, Karen Simonyan, and Yiming Yang. "Darts: Differentiable architecture search." *ICLR* 2019

DARTS: Differentiable Architecture Search

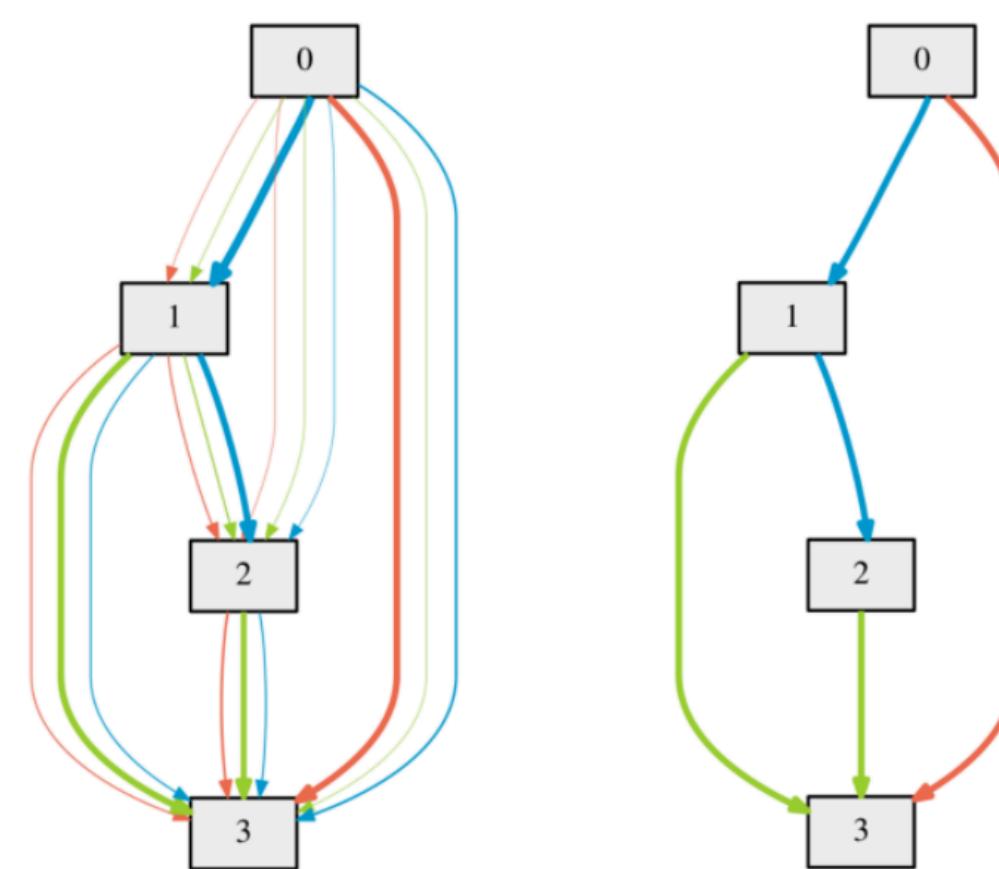
Challenge: architecture search is non-differentiable

Problem: Network performance (e.g. accuracy) does not change smoothly w.r.t architecture changes
 - we cannot use gradient-based optimisation :



DARTS solution: solve a **continuous relaxation** of the problem. To learn a cell:

- Place a mixture (weighted sum) of operations on each edge
- Jointly optimise network parameters and mixture probabilities
- Induce final architecture from mixing probabilities



Bilevel Optimisation

Each node can be computed from predecessors:

$$x^{(j)} = \sum_{i < j} o^{(i,j)}(x^{(i)}) \quad \text{operation from node } i \text{ to node } j$$

Relaxation: Consider mixtures of candidate operations, \mathcal{O} , via:

$$\bar{o}^{(i,j)}(x) = \sum_{o \in \mathcal{O}} \frac{\exp(\alpha_o^{(i,j)})}{\sum_{o' \in \mathcal{O}} \exp(\alpha_{o'}^{(i,j)})} o(x) \quad \text{operation weights}$$

The goal is then to learn $\alpha = \{\alpha^{(i,j)}\}$.

Let \mathcal{L}_{train} and \mathcal{L}_{val} denote training/validation loss.

Let w denote network parameters (e.g. convolution weights).

We'd like to solve a **bilevel optimisation** problem:

$$\min_{\alpha} \mathcal{L}_{val}(w^*(\alpha), \alpha) \quad \alpha \text{ is the upper-level variable}$$

$$\text{s.t. } w^*(\alpha) = \operatorname{argmin}_w \mathcal{L}_{train}(w, \alpha) \quad w \text{ is the lower-level variable}$$

Evaluating architecture gradients is prohibitively slow (the inner loop requires training a network) so we use an approximation:

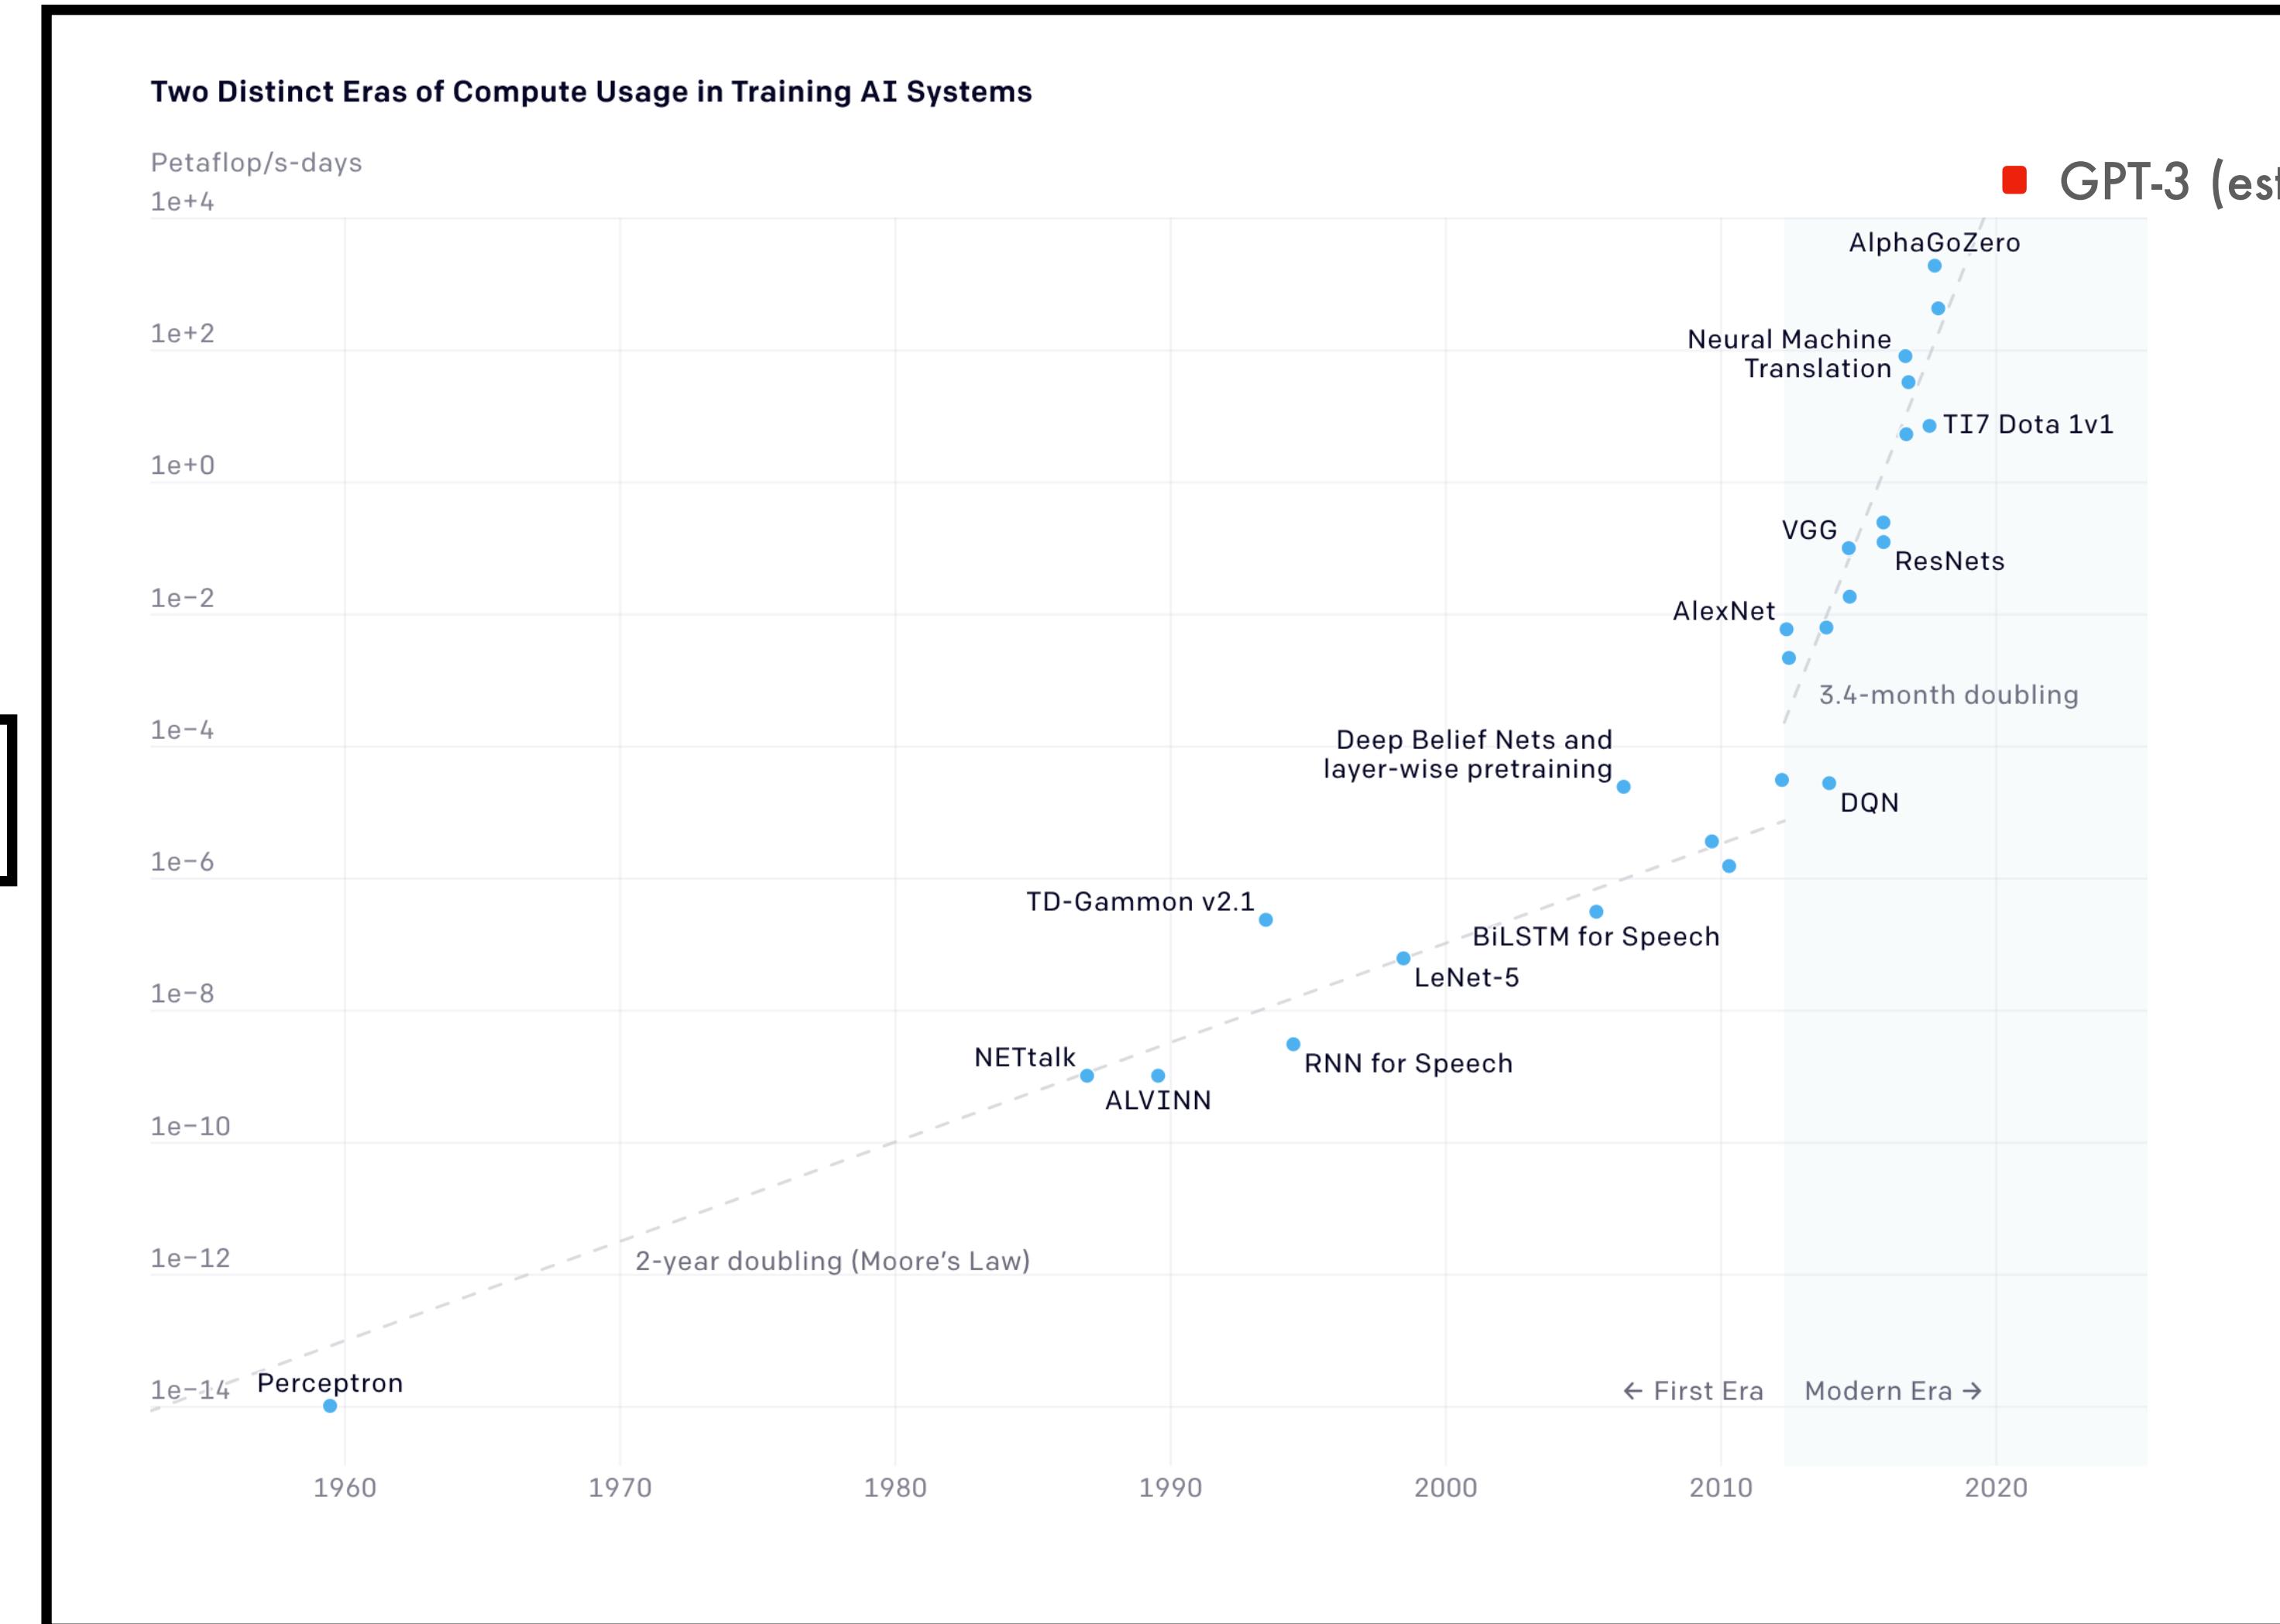
1 step of gradient descent

$$\nabla_{\alpha} \mathcal{L}_{val}(w^*(\alpha), \alpha) \approx \nabla_{\alpha} \mathcal{L}_{val}(w - \xi \nabla_w \mathcal{L}_{train}(w, \alpha), \alpha)$$

No formal convergence guarantees, but works in practice...

Scaling phenomena and the role of hardware

1 petaflop-day is approx.
8 V100 GPUs running for 1 day



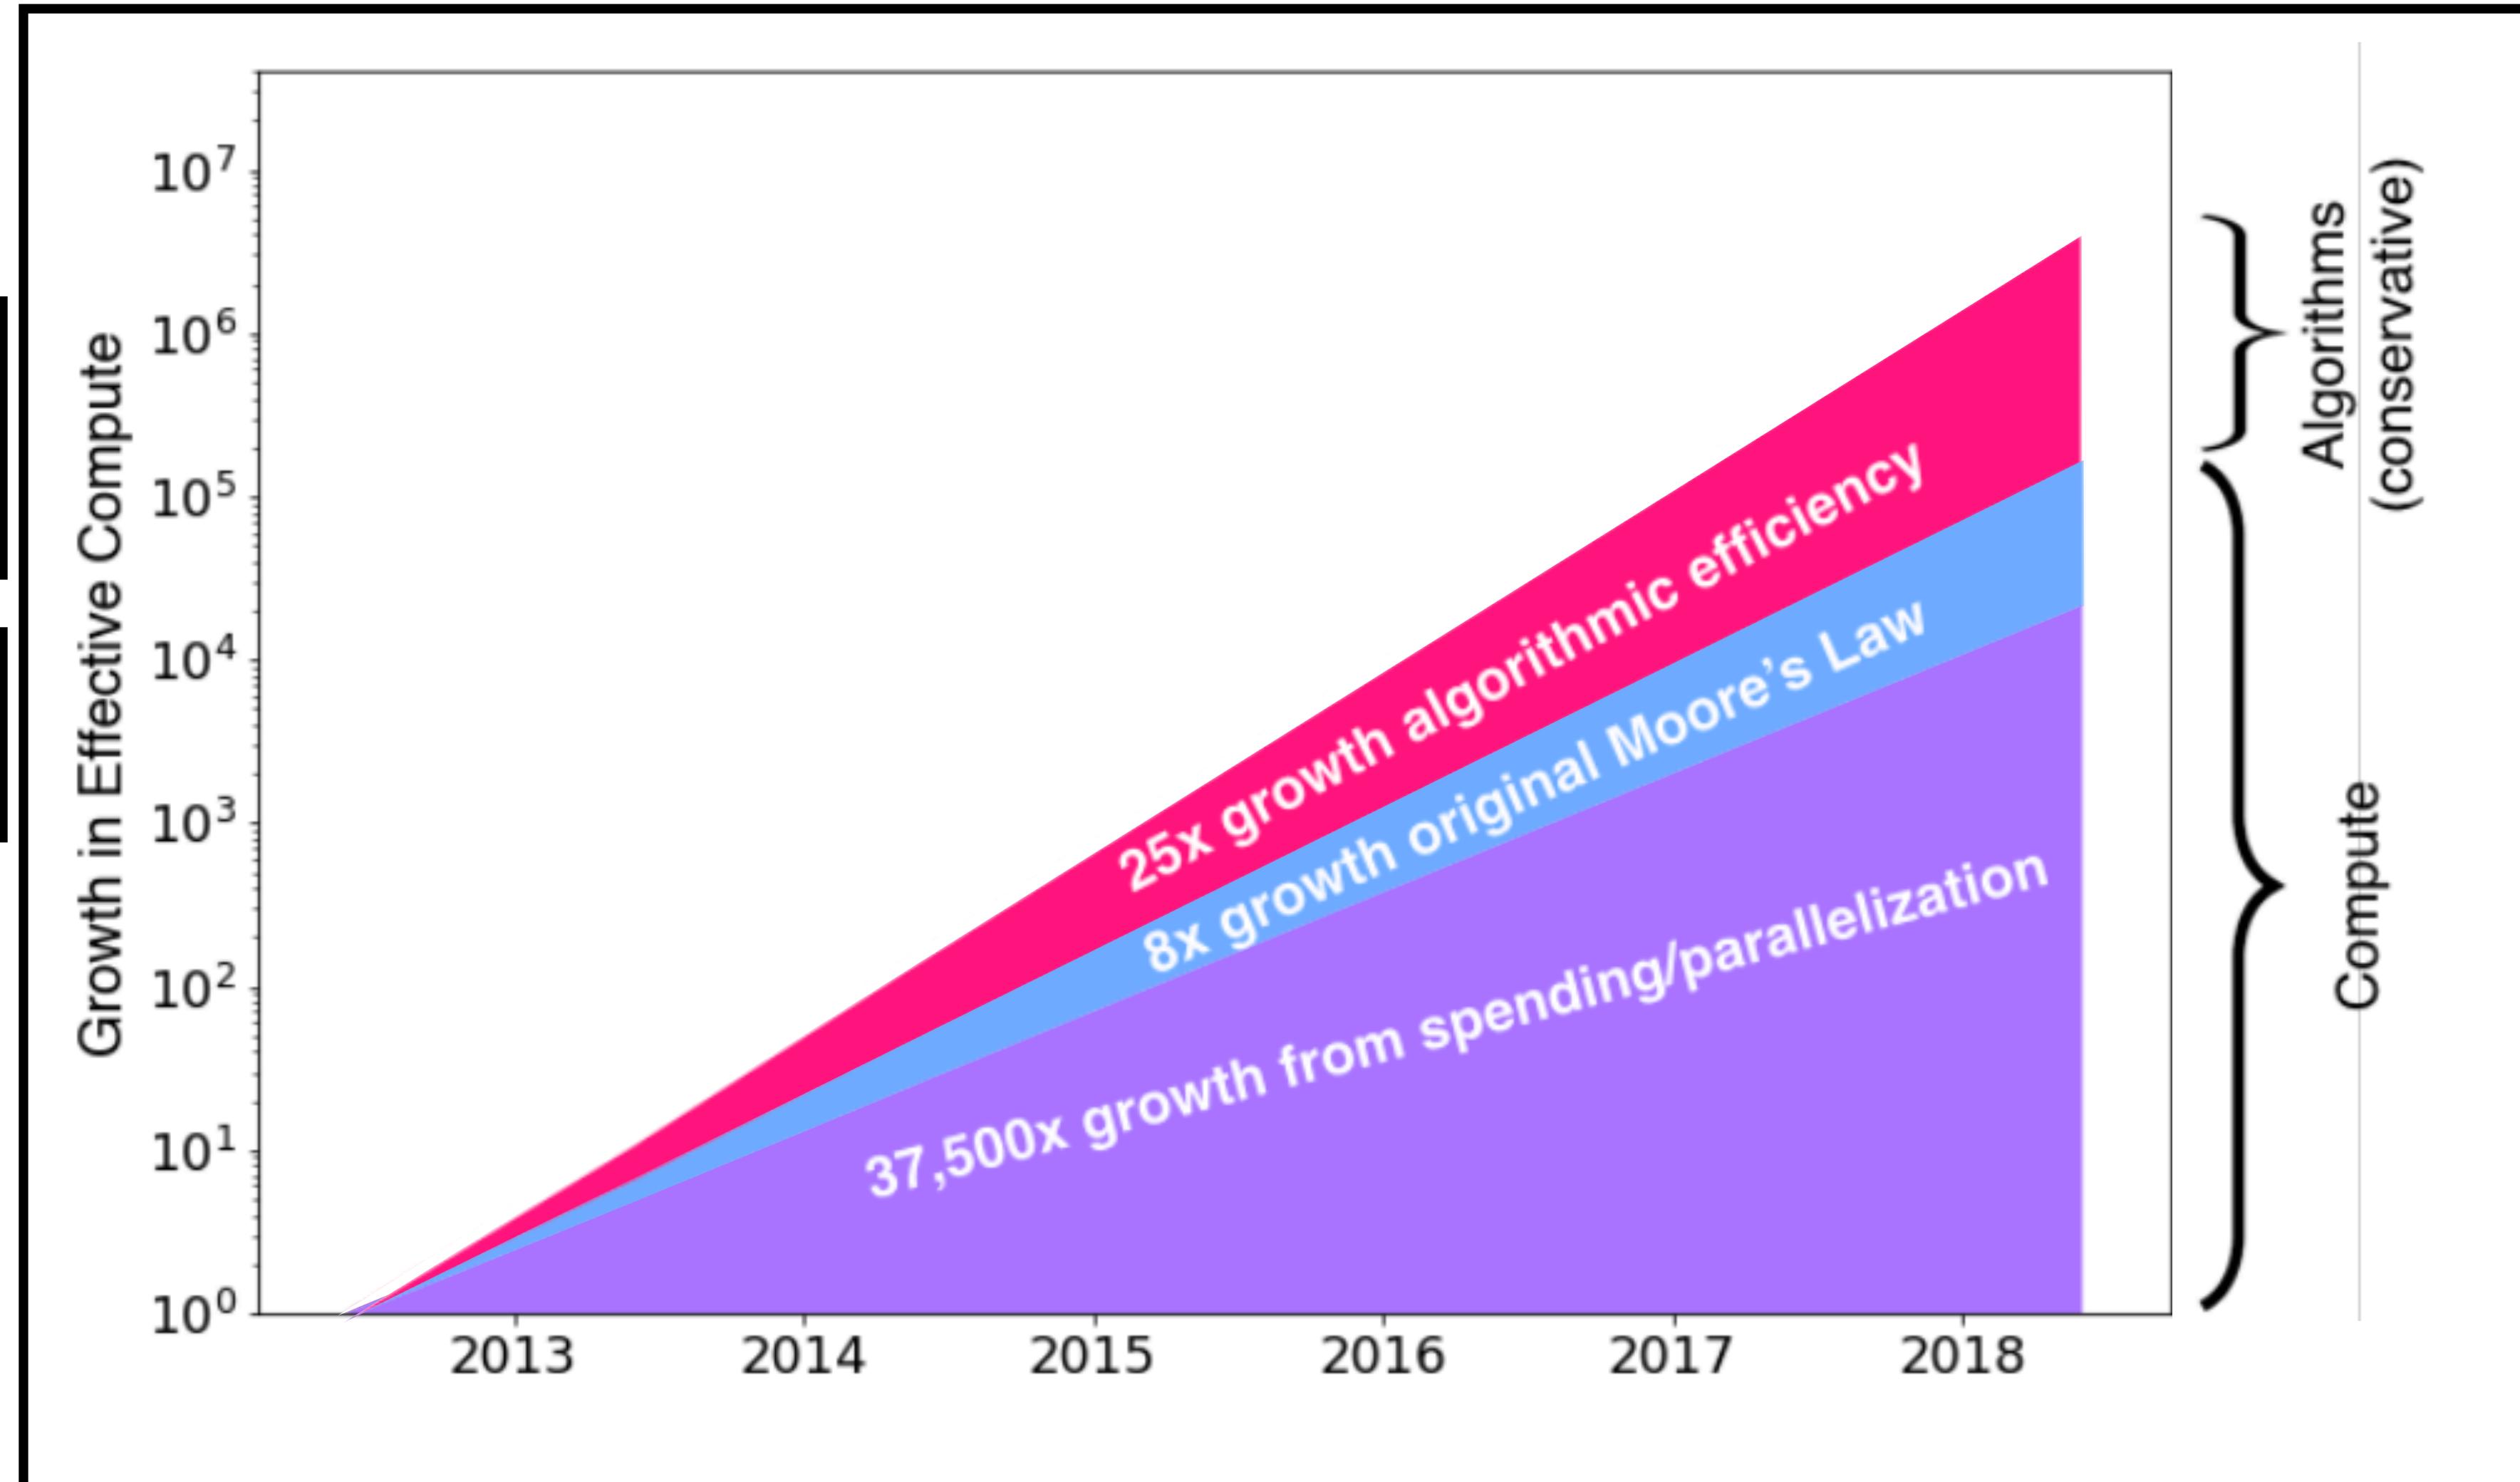
GPT-3 (175B parameters)
reportedly trained on a server
with several thousand GPUs

Megatron-Turing NLG 530B
(Nov, 2021) trained on 4K
A100 GPUs

What factors are enabling effective compute scaling?

Effective compute \approx FLOPs required to reach AlexNet-level ImageNet performance

Estimated cost of cloud compute for models like GPT-3: $\mathcal{O}(10\text{ Million})$ USD



Hernandez and Brown, "Measuring the Algorithmic Efficiency of Neural Networks." arXiv preprint arXiv:2005.04305 (2020).
<https://twitter.com/eturner303/status/1266264358771757057>

Scaling phenomena and the role of hardware

How important is scale for Deep Neural Networks?

Is it "just engineering", or something more fundamental?

Note: It is often challenging to analyse shifts from quantitative to qualitative differentiation.

Hierarchy of sciences

Is cell biology "*just*" applied molecular biology?

Is molecular biology "*just*" applied chemistry?

Is chemistry "*just*" applied many-body physics?

....

One science obeys the laws of the other.

But at each stage, new laws and concepts are necessary.

Qualitative vs Quantitative

FITZGERALD: The rich are different from us.

HEMINGWAY: Yes, they have more money.

"In almost all fields, a factor of ten means fundamentally new effects. If you increase magnification by a factor of 10 in Biology, you will see new things."

Hamming, Art of doing science and engineering, 1997

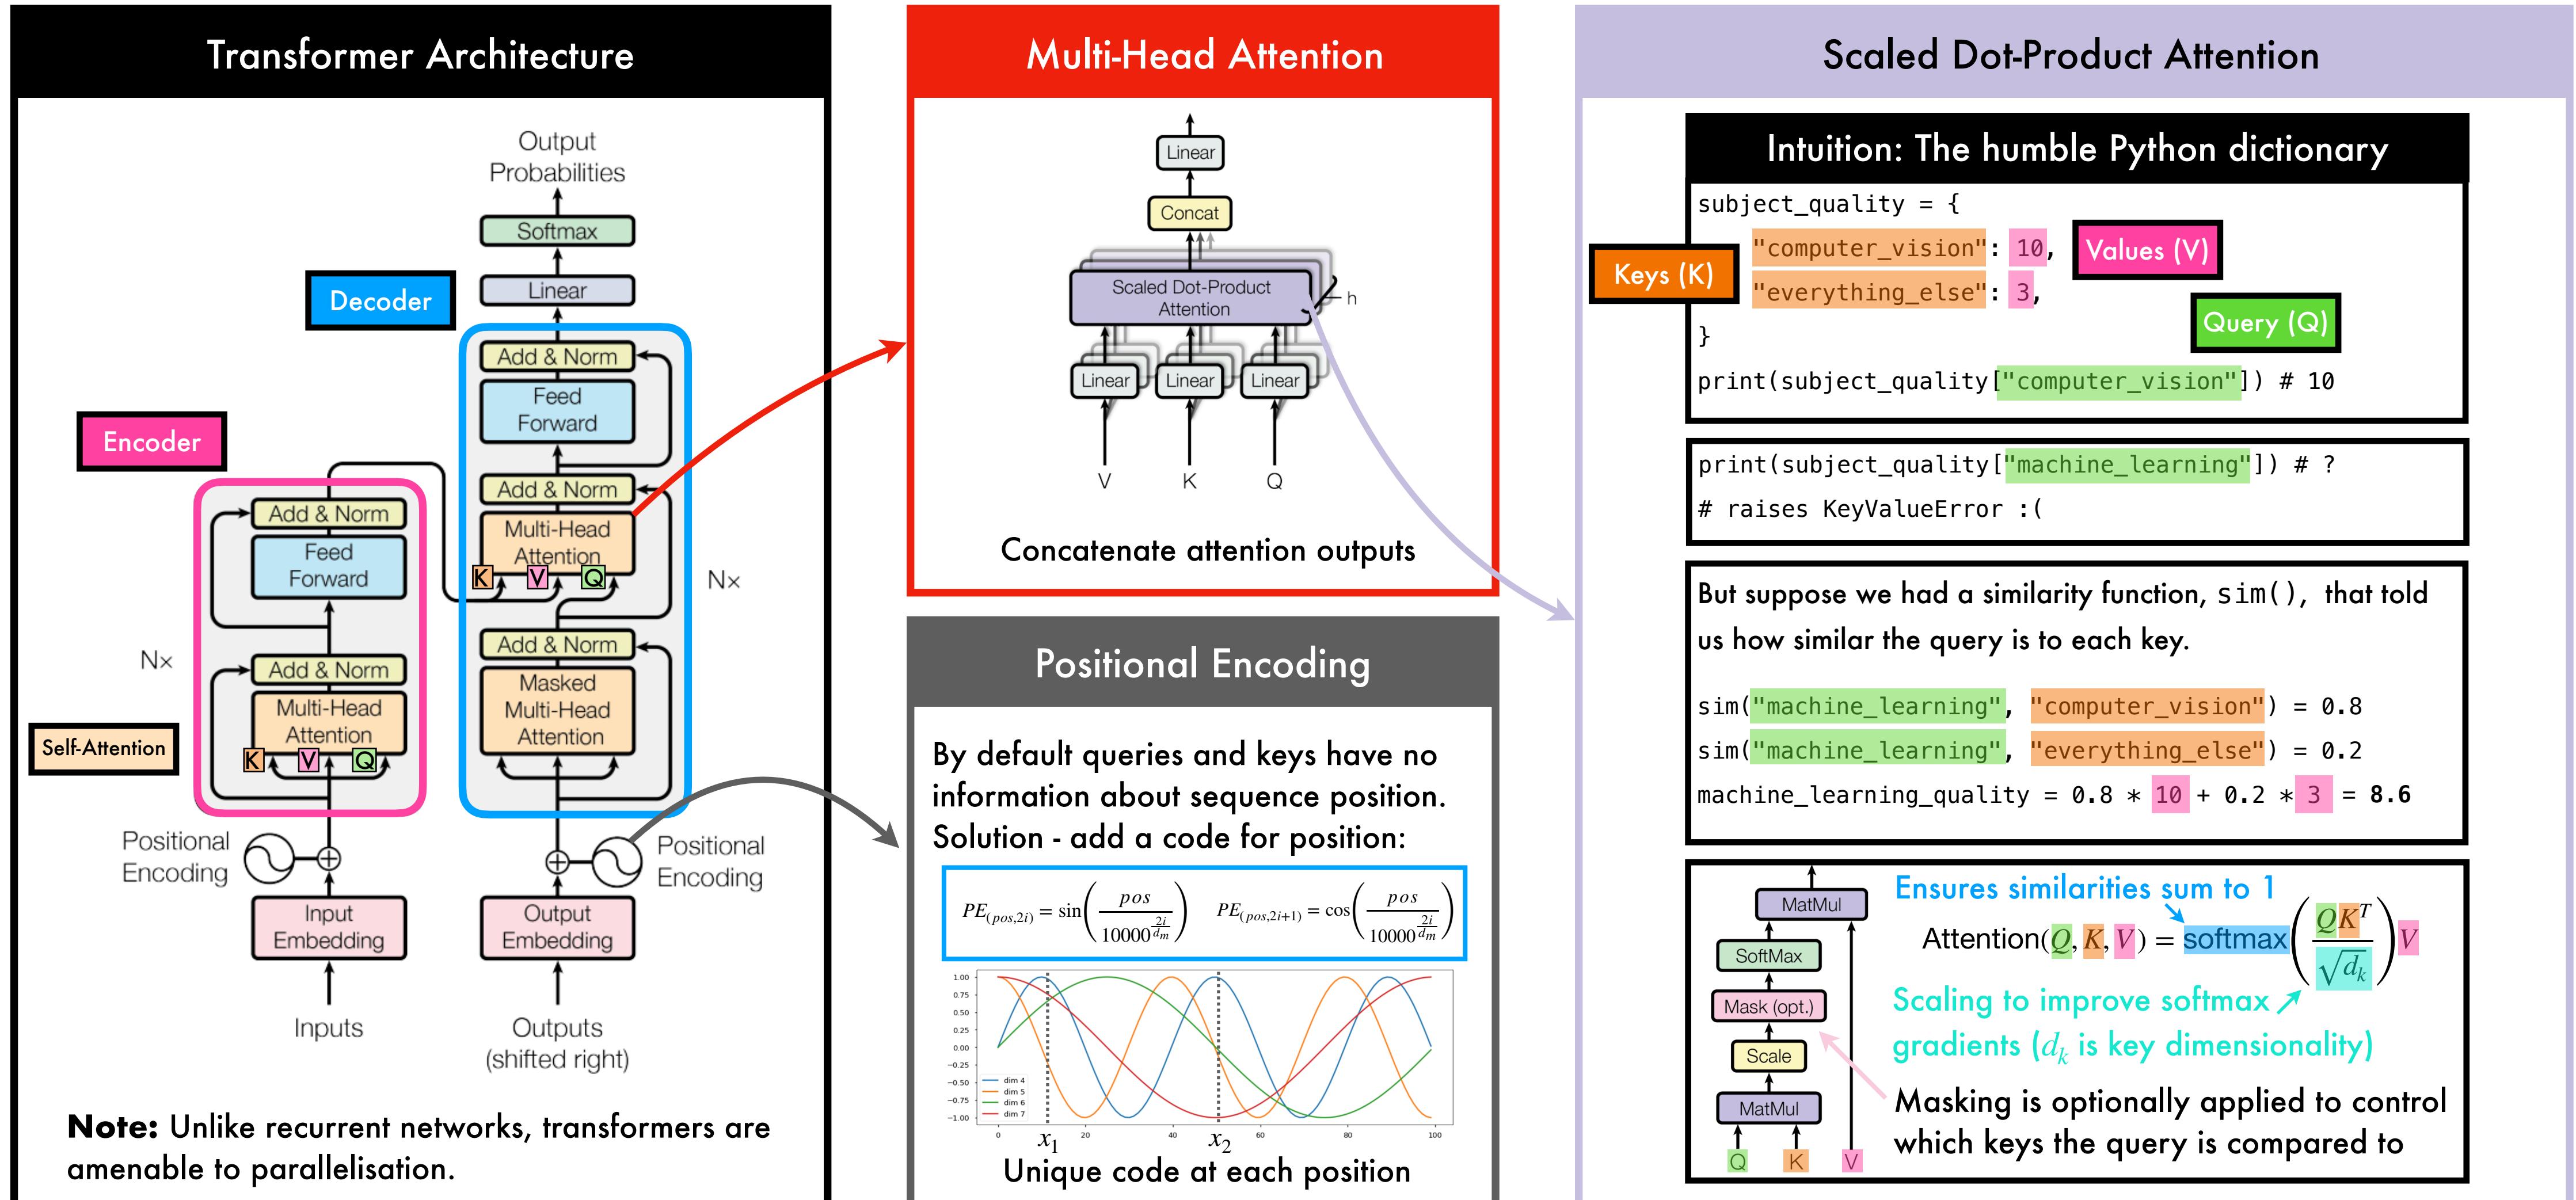
References/Footnotes:

P. Anderson, "More is different." Science 177 4047 (1972): 393-6

The "wisecrack" of Hemingway appears as a comment made by a character in one of his novels (<http://www.quotecounterquote.com/2009/11/rich-are-different-famous-quote.html>)

R. Hamming "The Art of Doing Science and Engineering: Learning to Learn." (1997)

The Transformer: a model that scales particularly well...

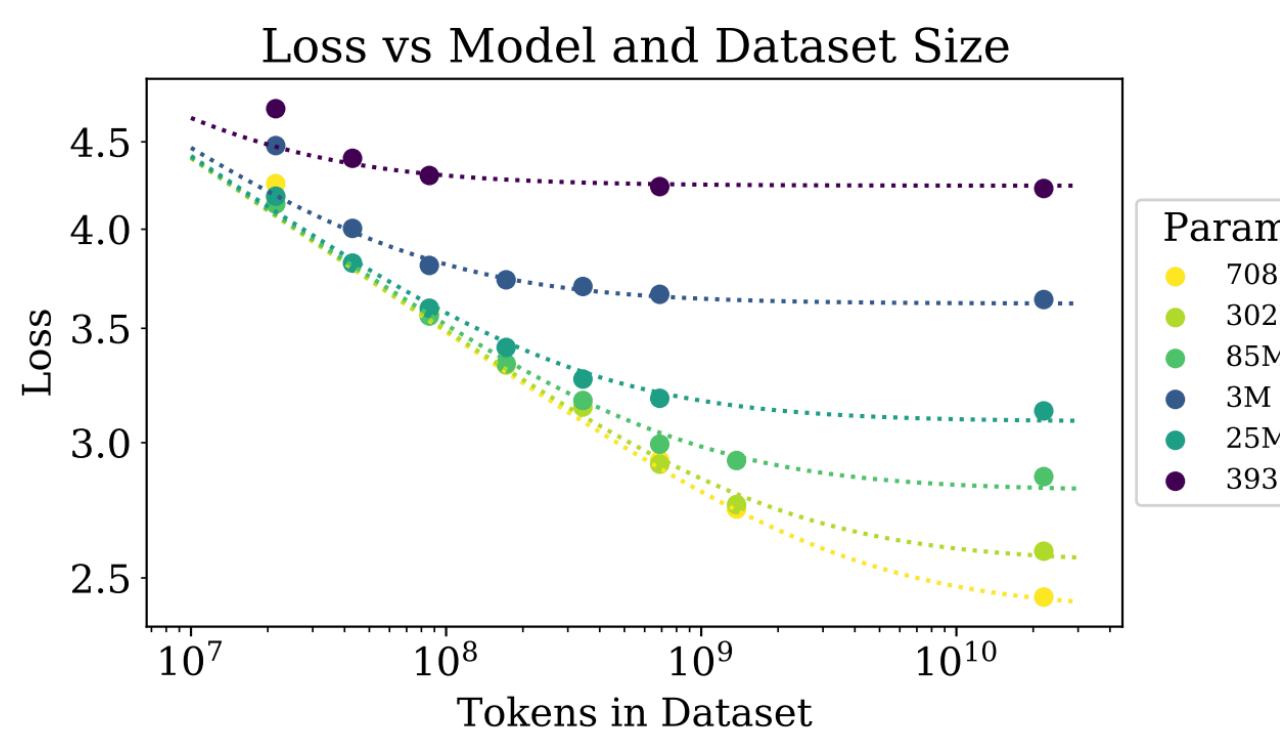


Transformer scaling laws for natural language

Predictable scaling

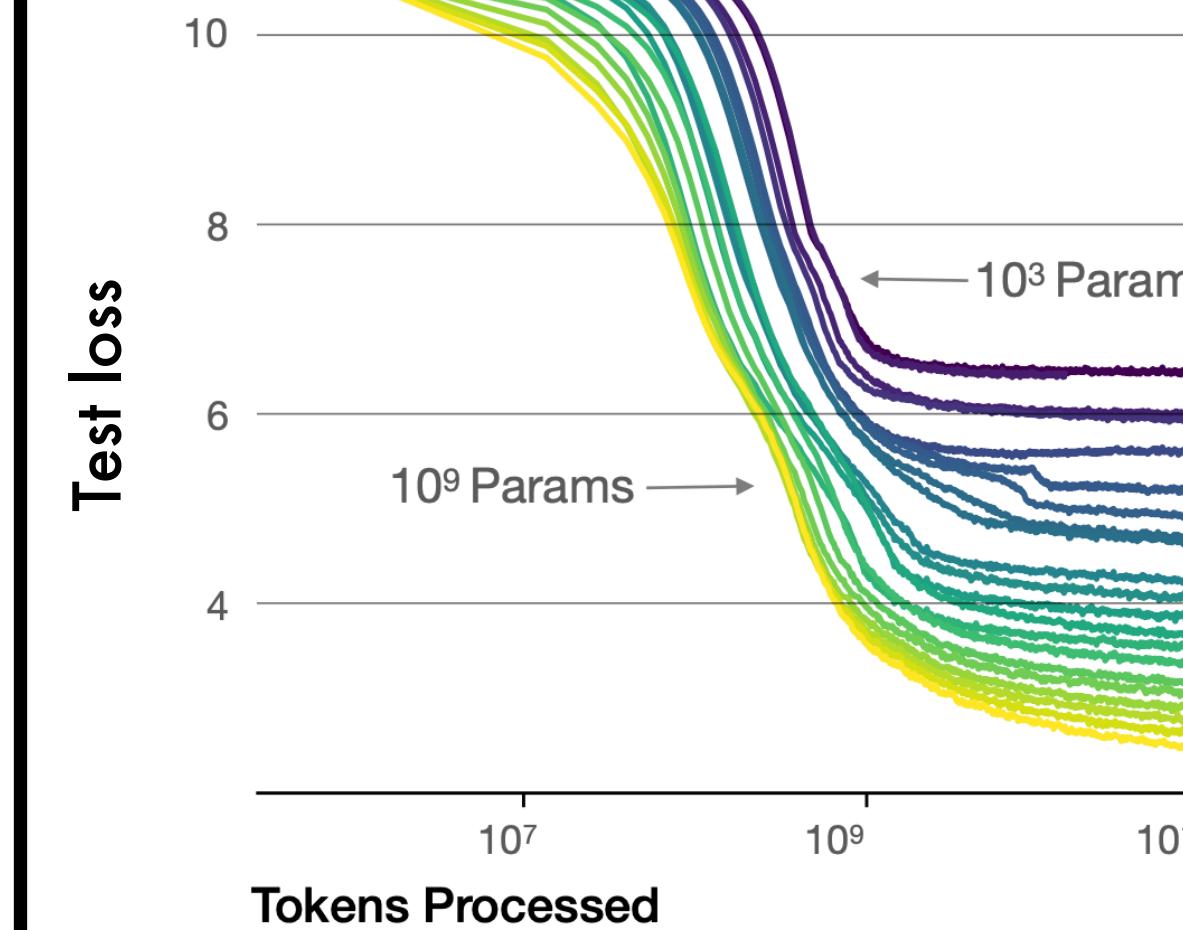
Transformer performance on language modelling tasks scales predictably as a *power law* with:

- Compute
- Training data size
- Model size

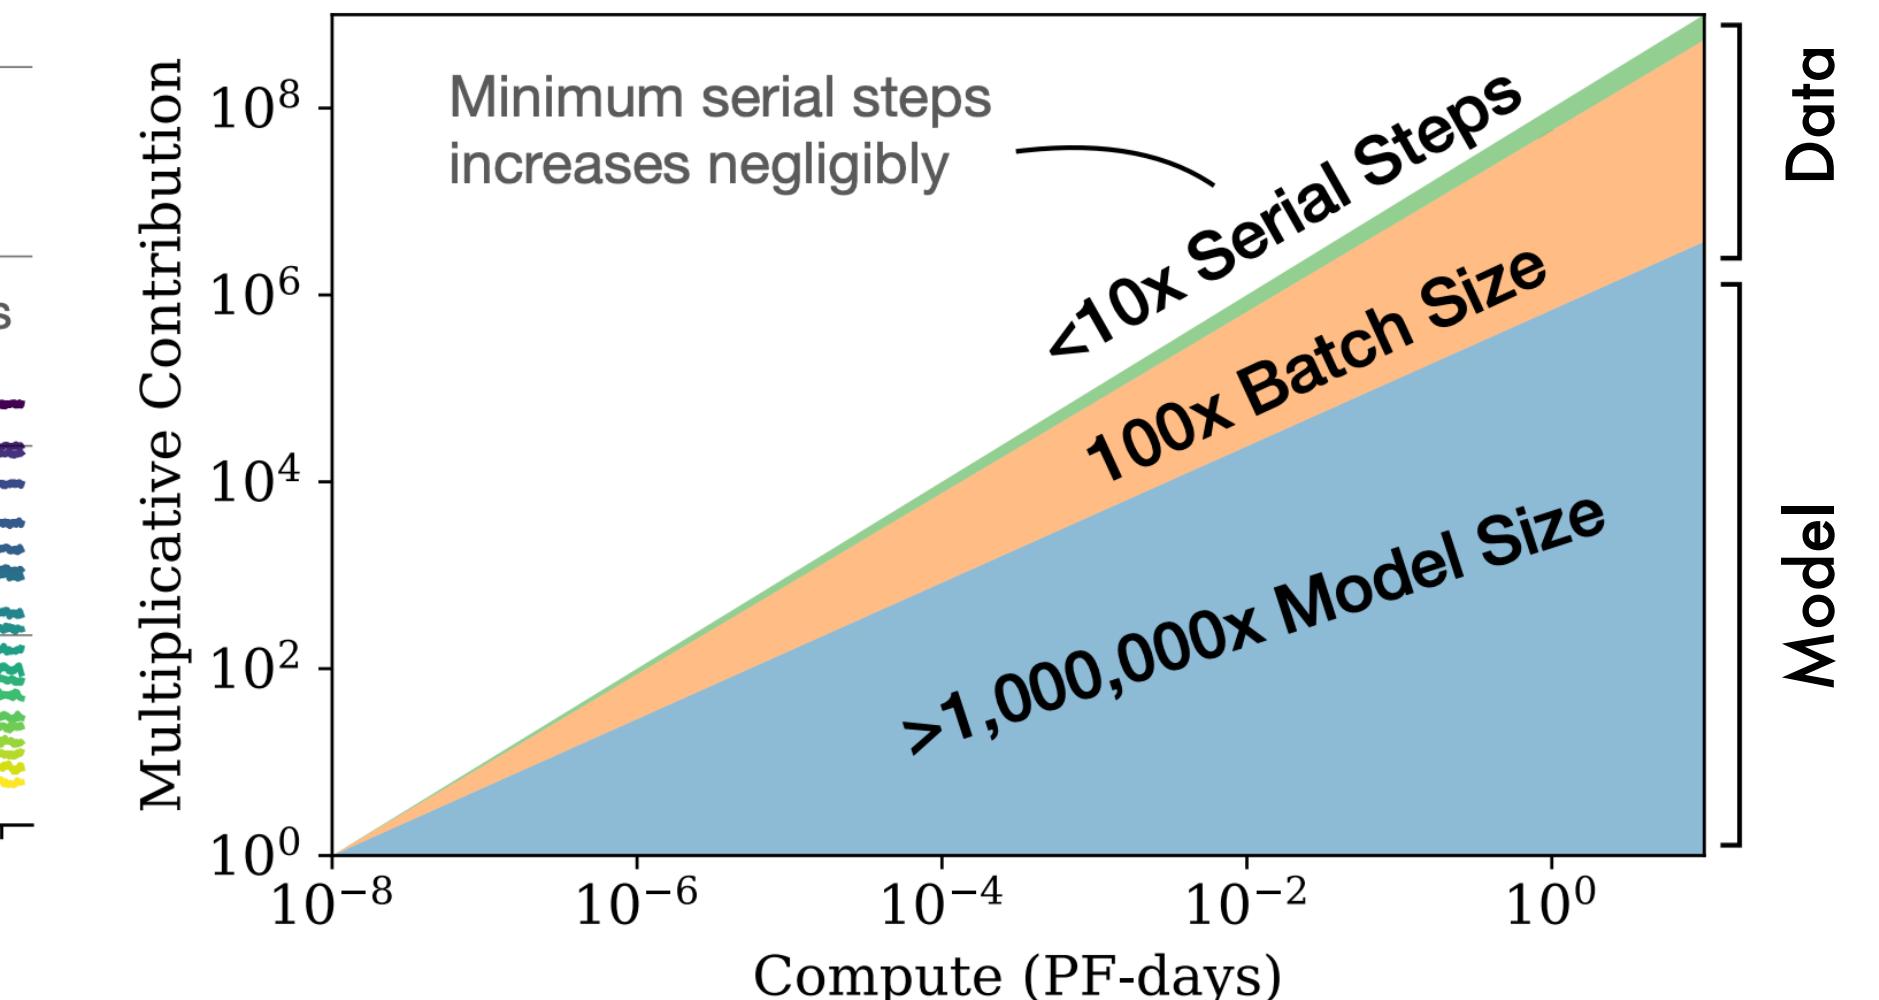


Some power laws were found that span more than seven orders of magnitude.

Intriguing characteristics

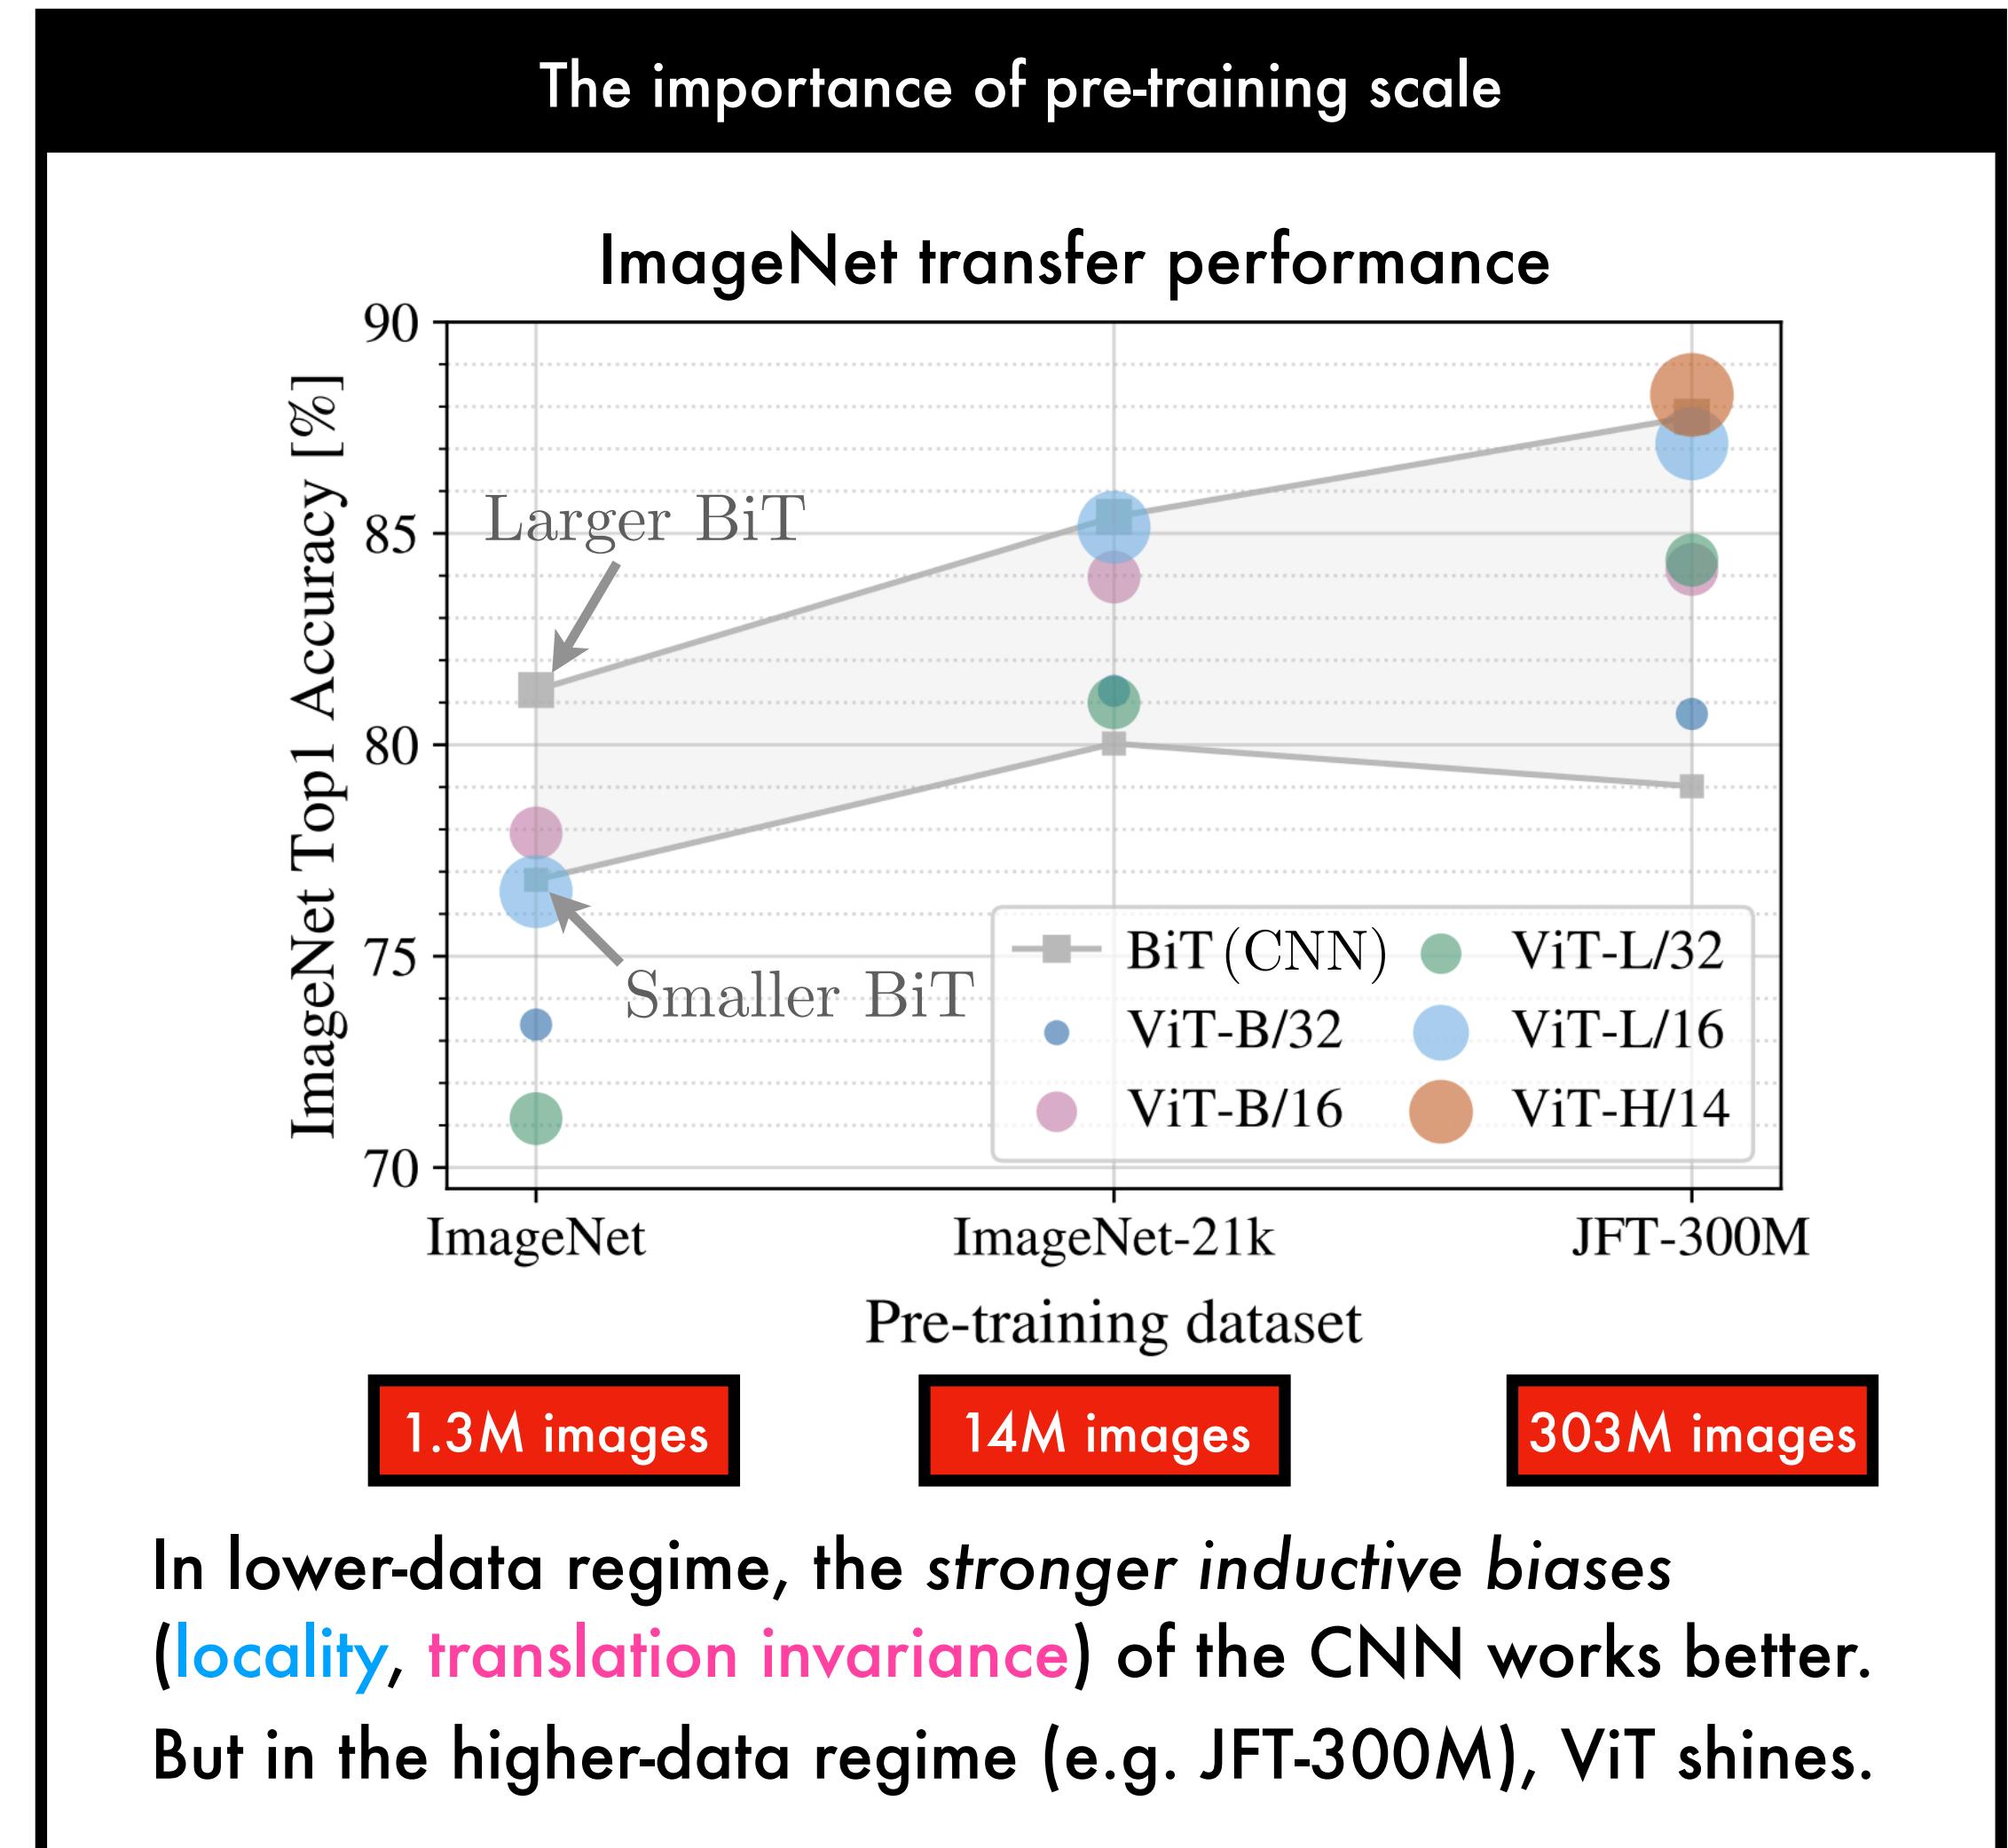
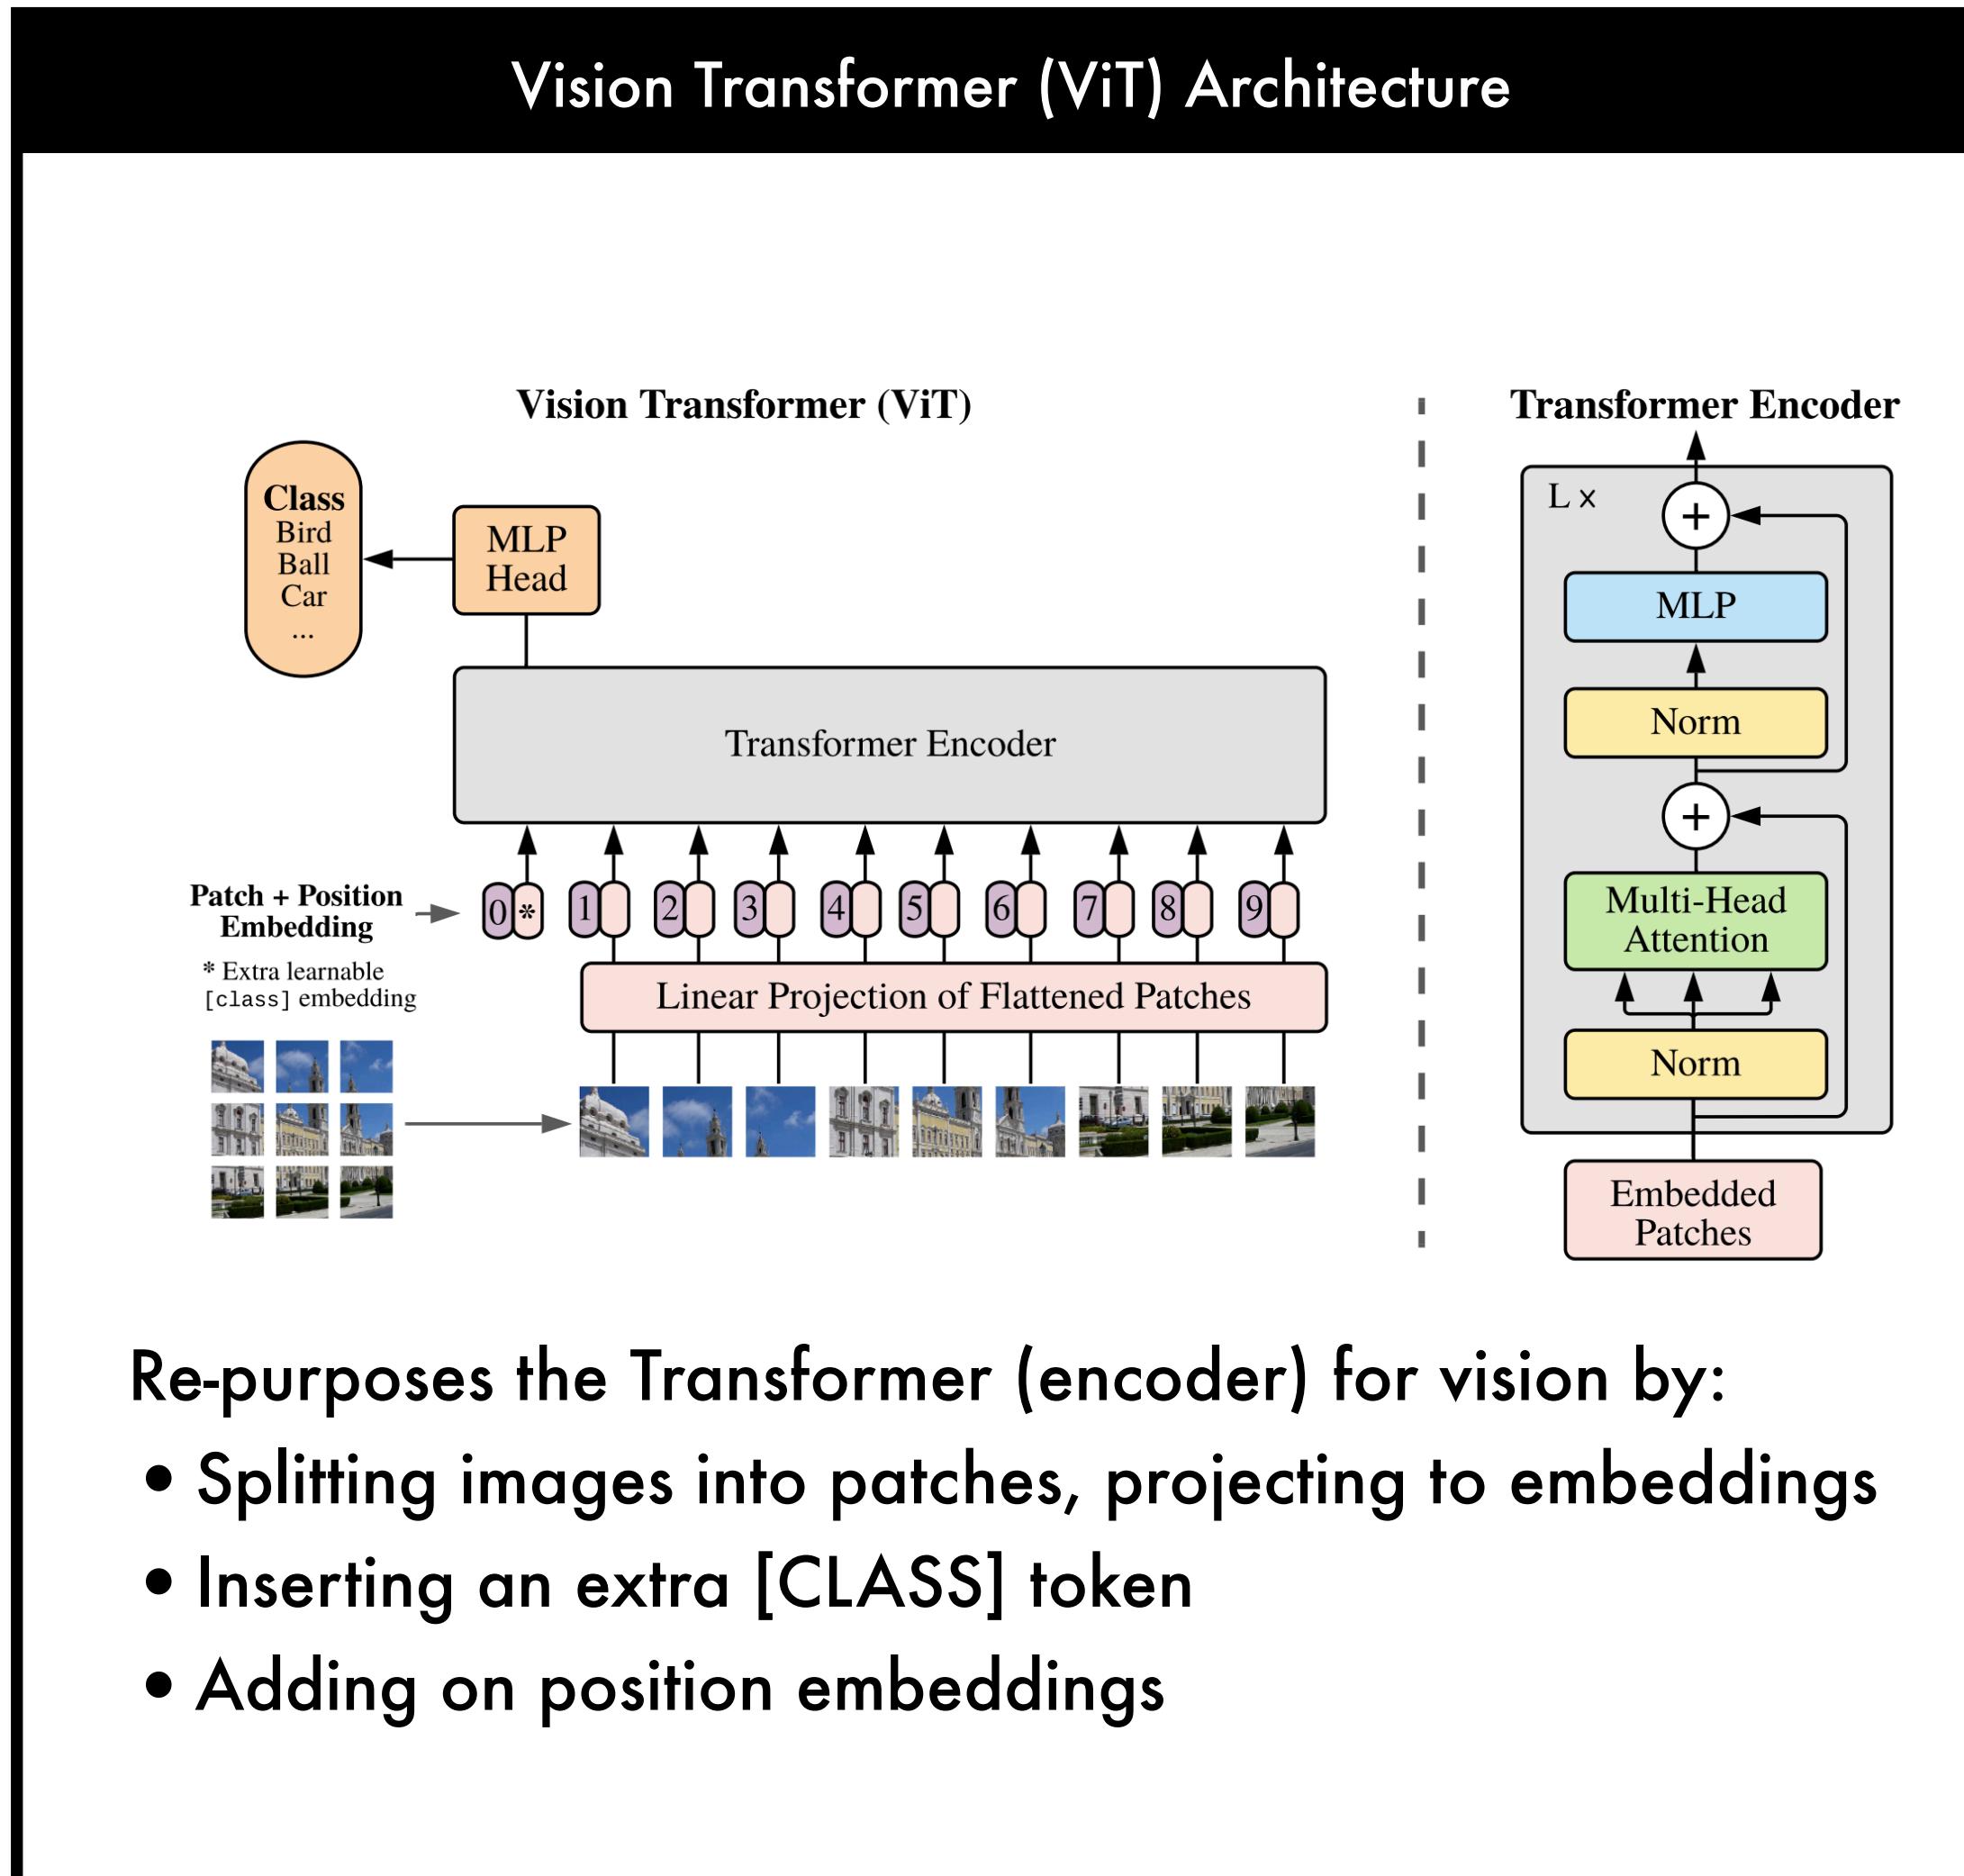


Larger models require **fewer samples** to reach the same performance.



If extra compute is available, allocate most towards increasing the **model size!**

Vision Transformer



Transformer Explosion

Historical context: non-local means

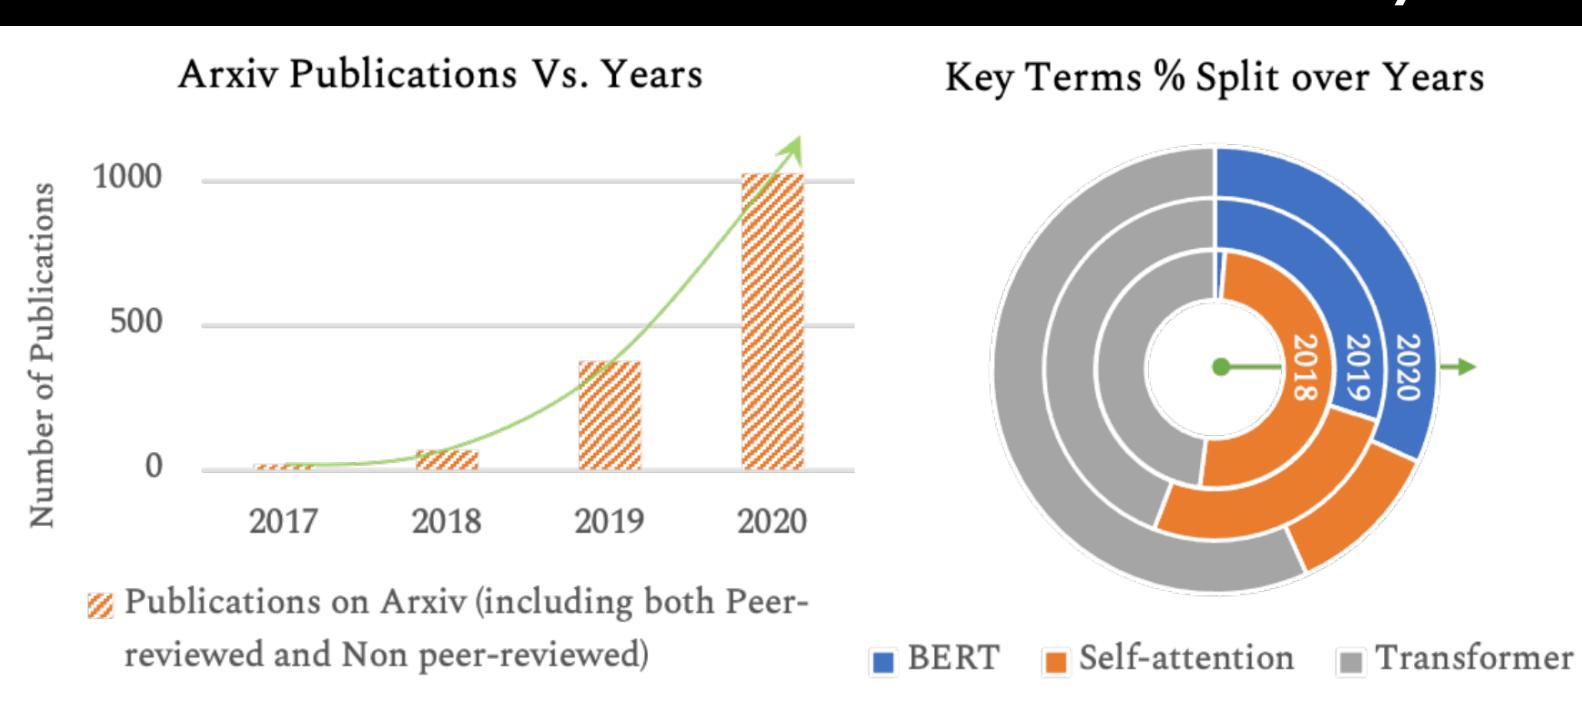
The "self-attention" operation has long been used in the image processing community for de-noising, under the name "non-local means":

$$NL[v](i) = \sum_{j \in I} w(i, j)v(j)$$

Here $v = \{v(i) | i \in I\}$ is a noisy image, and the weights $\{w(i, j)\}_j$ depend on the similarity between pixels i and j .

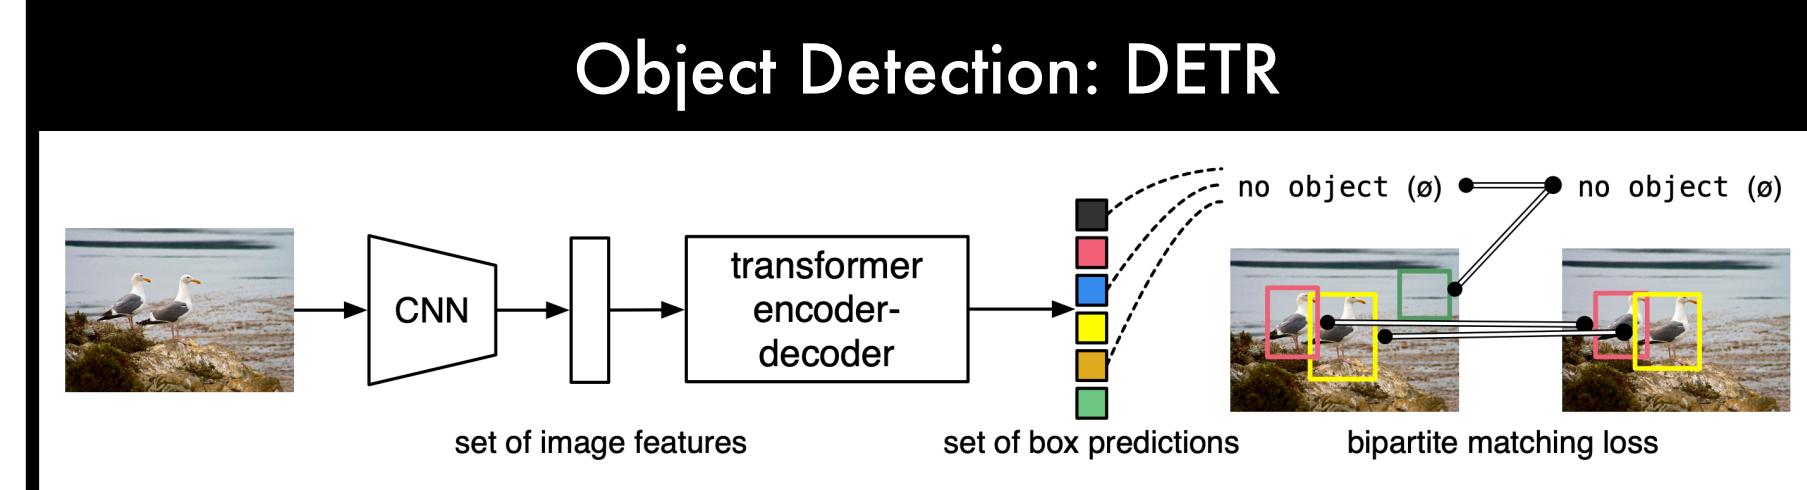
However, the broad applicability and value of this (highly flexible) operation has become clearer in recent years.

Research Interest in the Vision Community

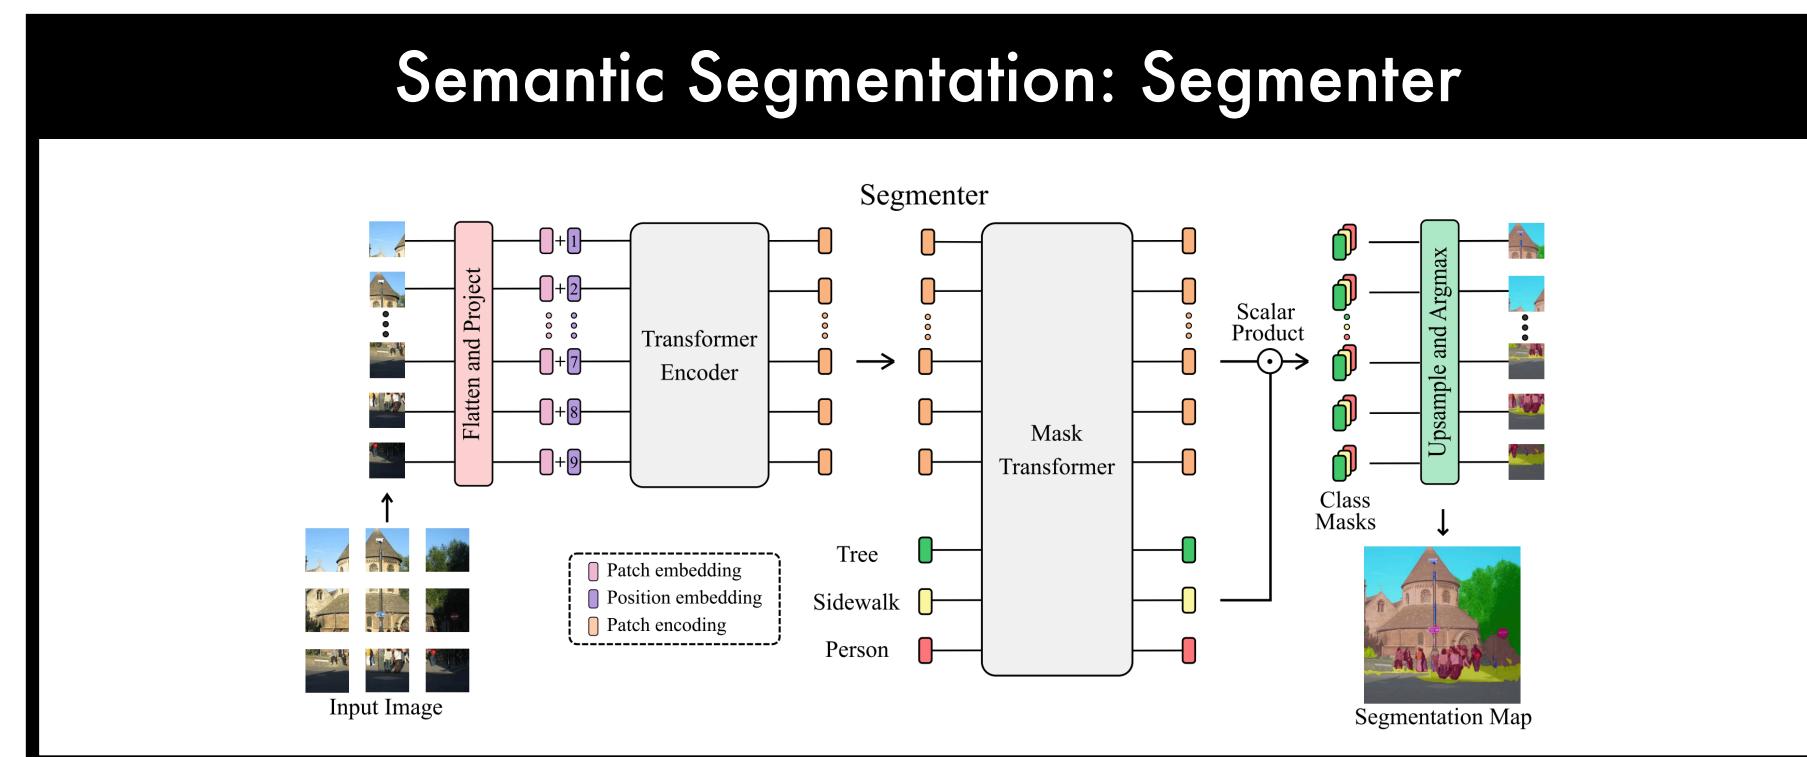


- Buades et al. "A non-local algorithm for image denoising." CVPR 2005
- Wang et al. "Non-local Neural Networks." CVPR 2018
- Khan et al. "Transformers in Vision: A Survey." ArXiv abs/2101.01169 (2021)

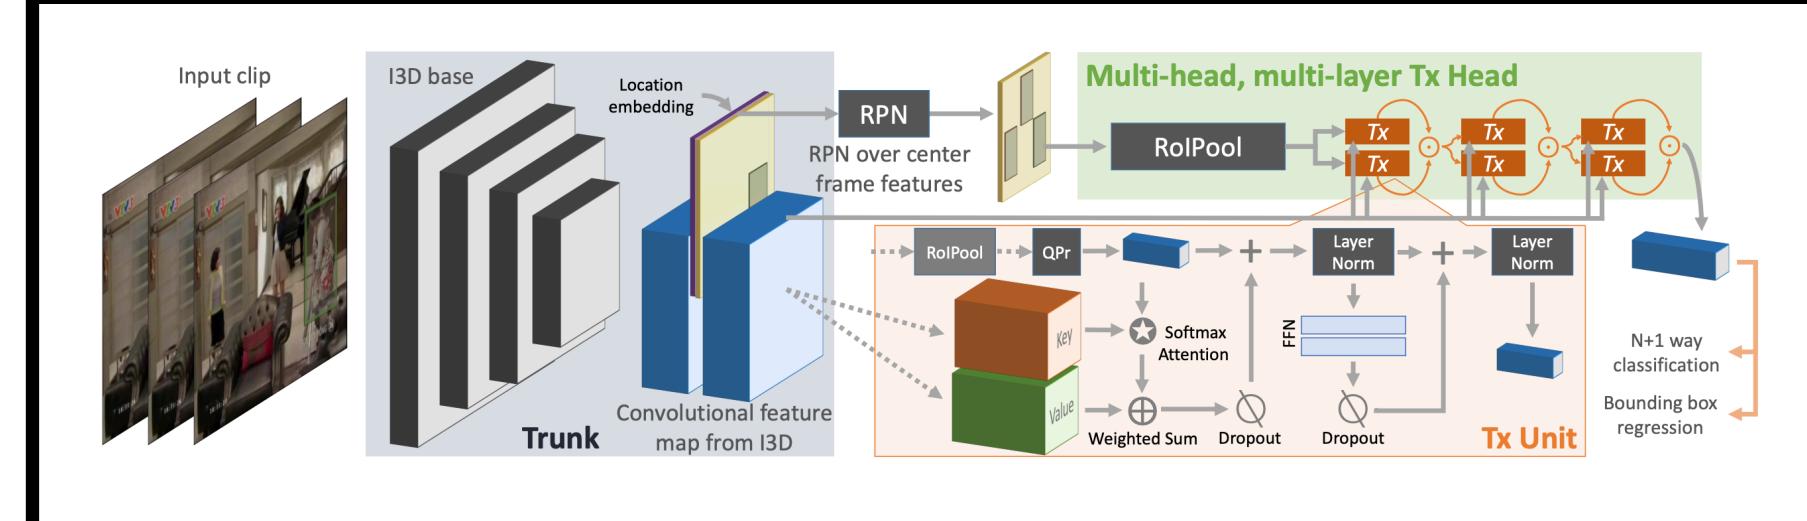
Object Detection: DETR



Semantic Segmentation: Segmenter



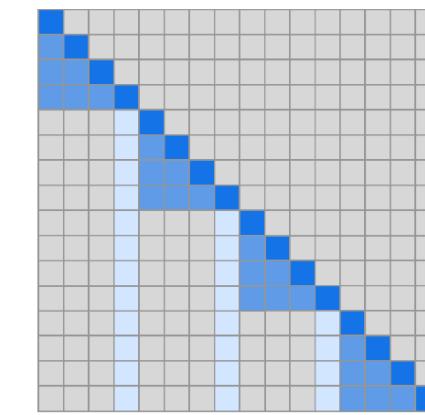
Action Recognition: Video Action Transformer Network



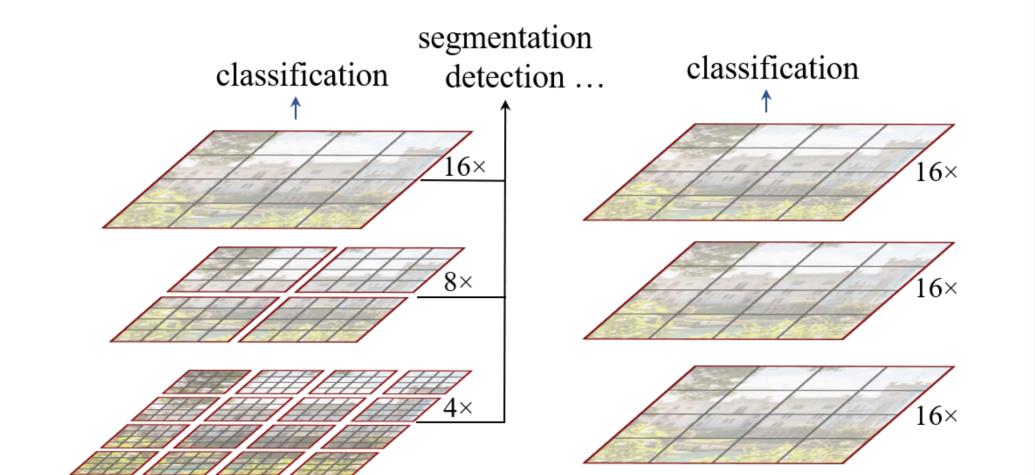
- Carion et al. "End-to-End Object Detection with Transformers." ECCV (2020)
- Strudel et al. "Segmenter: Transformer for Semantic Segmentation." ICCV 2021
- Girdhar et al. "Video Action Transformer Network." CVPR 2019

Computational tricks

Problem: self-attention has quadratic complexity in the input size (every element attends to every other element). Many solutions have been proposed, including:



The Sparse Transformer factors attention to reduce complexity to $\mathcal{O}(n\sqrt{n})$



The Swin Transformer achieves linear complexity by restricting self-attention to fixed regions (like a CNN....).

- Child et al. "Generating Long Sequences with Sparse Transformers." ArXiv abs/1904.10509 (2019)
- Liu et al. "Swin Transformer: Hierarchical Vision Transformer using Shifted Windows." ICCV 2021

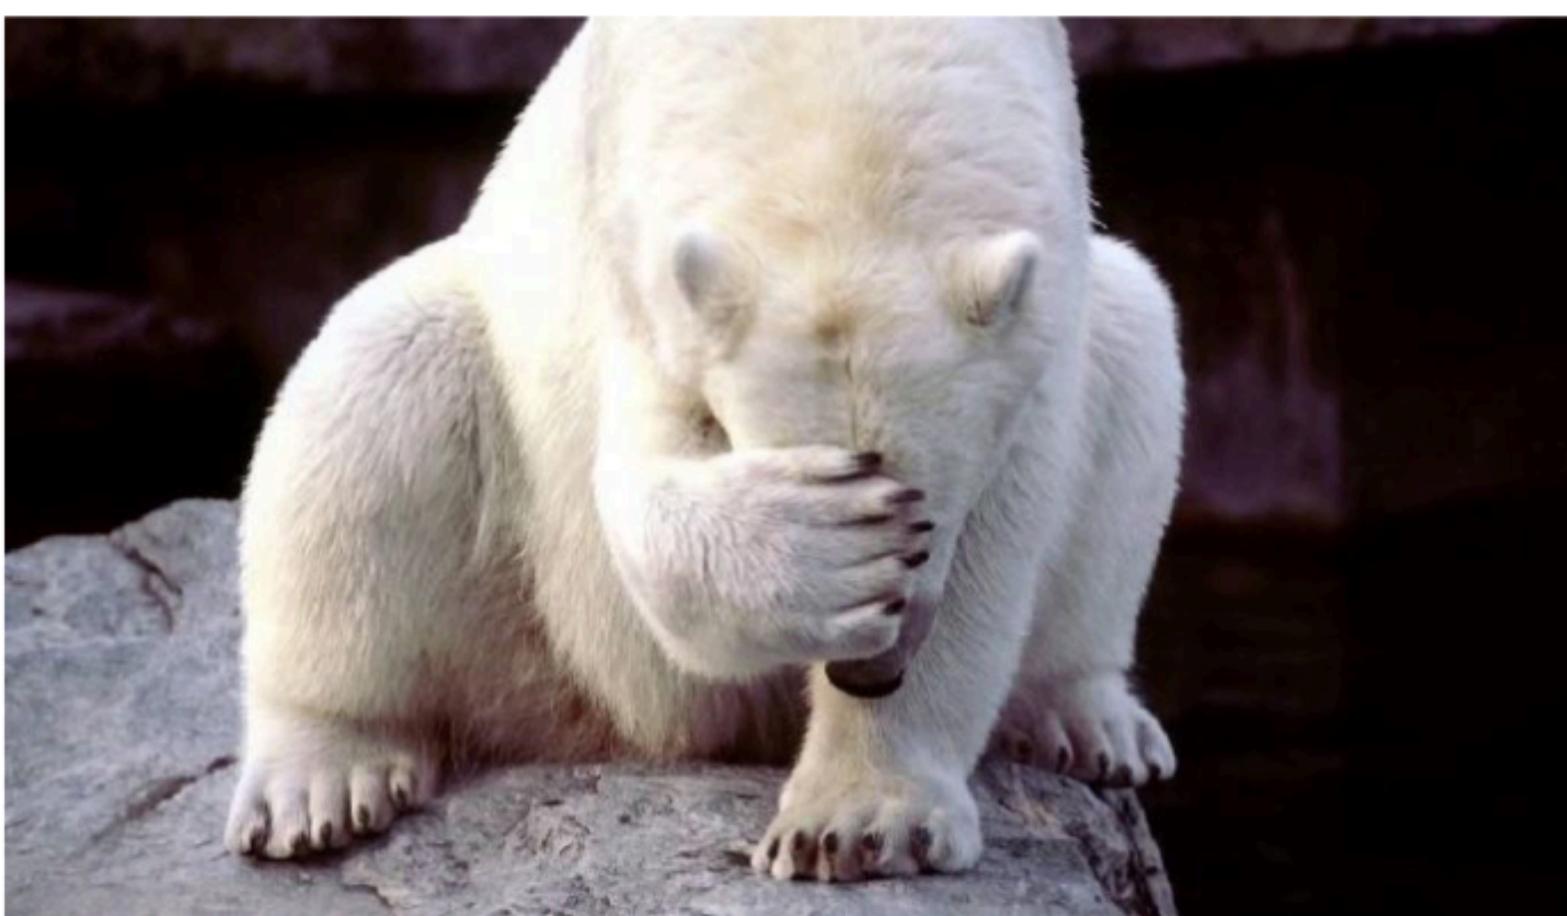
Neural Network Design and Energy Consumption

Deep Neural Networks are Energy Intensive

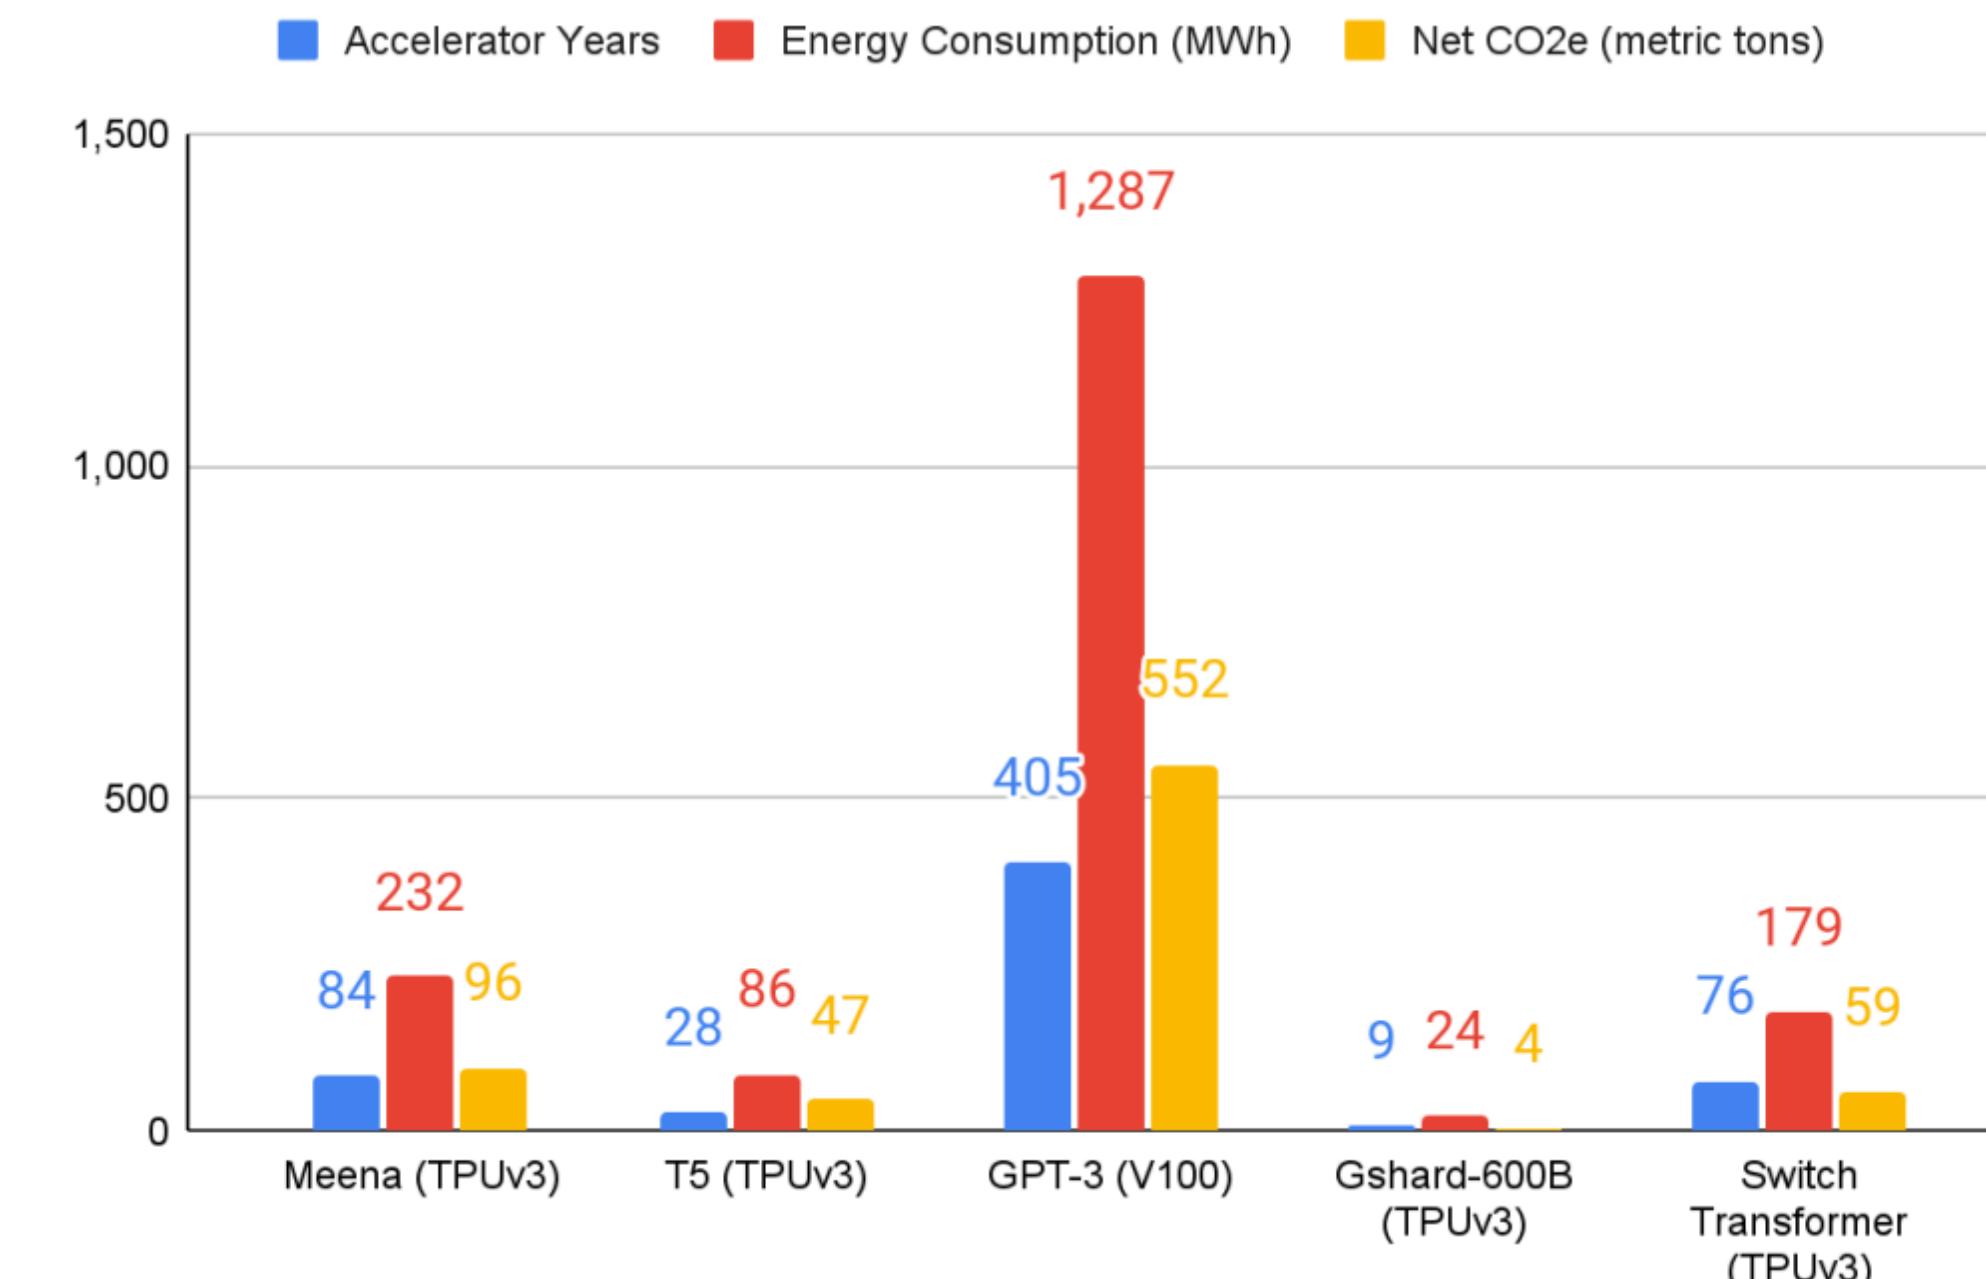
Consumption	CO ₂ e (lbs)
Air travel, 1 passenger, NY↔SF	1984
Human life, avg, 1 year	11,023
American life, avg, 1 year	36,156
Car, avg incl. fuel, 1 lifetime	126,000

Training one model (GPU)

NLP pipeline (parsing, SRL)	39
w/ tuning & experimentation	78,468
Transformer (big)	192
w/ neural architecture search	626,155



Transformers represent many of the biggest models



Reasons for optimism:

- There are significant opportunities for grid efficiency: training is not time-sensitive (can be scheduled to maximise peak renewable energy times)
- Fusion is only 30 years away....

Strubell, Emma et al. "Energy and Policy Considerations for Deep Learning in NLP." ArXiv abs/1906.02243 (2019)

Image credit: <https://www.desktopbackground.org/wallpaper/white-bear-put-hand-on-head-wild-animal-wallpaper-jpg-492933>

Patterson et al. "Carbon Emissions and Large Neural Network Training." ArXiv abs/2104.10350 (2021)

End of Lecture 15