

Overview of Natural Language Processing

Part II & ACS L90

Lecture 5: Dependency Parsing

Andrew Caines

based on slides by Ann Copestake, Simone Teufel, Paula Buttery, Weiwei Sun

Department of Computer Science and Technology
University of Cambridge


Michaelmas 2021/22

Part 1: Background

Recall lecture 4: limitations of CFGs

Natural languages are provably **non-context-free**.

Natural languages as mildly context-sensitive languages



An alternative: model relations between word tokens

- rather than directly modelling phrase structure: model (labelled) relations between *word tokens*
- NB: why *word tokens*?

Aside: what's a word (token)?

- Well, what's a word?
- “Count the words in this sentence.” =6 (orthography)
- “That's easy!” 2?
- tokenized →1:That 2:'s 3:easy 4:!
- phon. “Why you gotta be a time-waster?”
- →1:Why 2:you 3:got 4:ta 5:be 6:a 7:timewaster 8:?
- For syntactic parsing: morpho-syntax over phonology/orthography ... (plus semantics?)
- MWEs: “It's a dog eat dog world.”
- →1:It 2:'s 3:a 4:dog_eat_dog 5:world 6:.
- Agglutination (root&affixes): Turkish *gelememiş* “apparently he couldn't come”
- Polysynthesis (diff.word.classes): Yupik (Alaska-Siberia)
tuntussuqatarniksaitengqiggtuq “He had not yet said again that he was going to hunt reindeer”
- Tokenization as a fundamental task in NLP (En pre-trained tokenizers);
nb complexity in Chinese, Japanese, Arabic, Hebrew, etc

An alternative: dependency grammar

- (un)Labelled relations between *word tokens*
- Binary directed relations from **heads** to **dependents**
- Dependency structure as a directed graph: $G = (V, A)$
- A set of vertices V (word tokens in a string: *written sentence, transcribed spoken utterance, tweet, etc*)
- A set of ordered pairs of vertices A , capturing relations between elements in V
- Different linguistic formalisms with different constraints, but for computational approaches, G should be acyclic and have:
 - ① a single designated root node with no incoming arcs
 - ② vertices (word tokens) with exactly 1 incoming arc each (except the root node which has 0)
 - ③ a unique path from the root node to each vertex

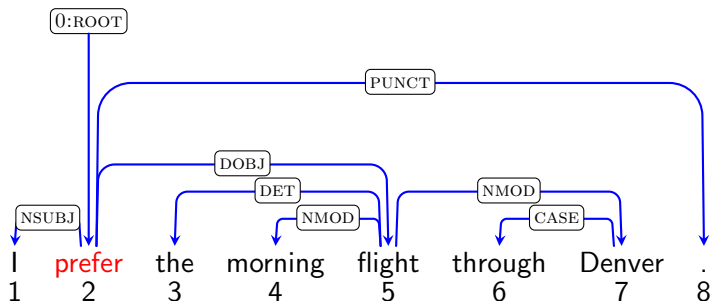
An alternative: dependency grammar

- First: tokenize string $x \rightarrow$ indexed word tokens \rightarrow implicit **root** token at location 0
- Identify head–dependent token pairs
- (optionally/usually) With labels from a fixed inventory, e.g. root, nsubj, dobj, det, nmod
- Phrase structure not directly modelled, explicit structuring of nested constituents dropped (might infer it)
- Instead dependency parse to find relations between lexical items (& the root)

An example: dependency structure

I prefer the morning flight through Denver .

An example: dependency structure



Remember:

- 1 a single designated root node with no incoming arcs
- 2 vertices (word tokens) with exactly 1 incoming arc each (except the root node)
- 3 a unique path from the root node to each vertex

Dependency relations

- ROOT
- NSUBJ: nominal subject
- (D)OBJ: direct object
- DET: determiner
- NMOD: nominal modifier
- CASE: case marker
- PUNCT
- And more, and alternatives

Dependency structure & phrase structure side-by-side

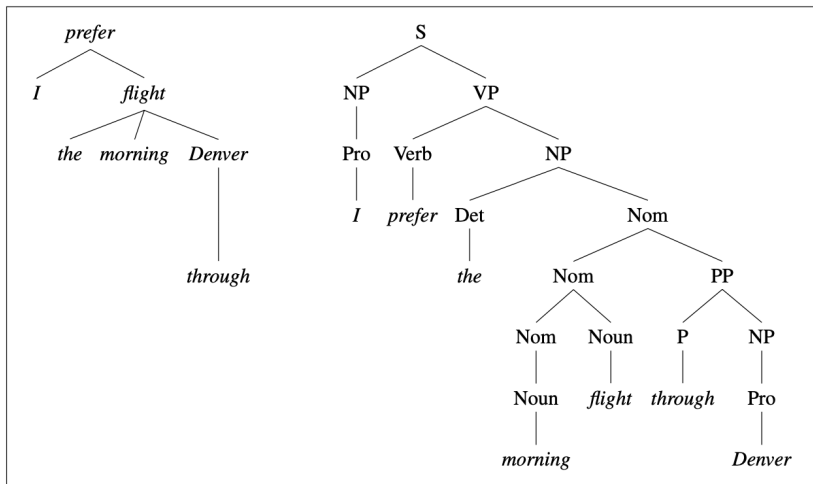



Figure 14.1 Dependency and constituent analyses for *I prefer the morning flight through Denver*.



(Formal Models of Language: weak equivalence between dep.gramms & CFGs, can derive same set of strings, systematic mapping b/w trees)

Dependency structure (contd)




- XCOMP: open clausal complement (without its own subject; subject usually the object or subject of the higher clause)
- MARK: marker to indicate subordinate clauses, e.g. *that, if, whether, to* (might be omitted)

Dependency structure (contd)




- CCOMP: clausal complement
- MARK: marker (might be omitted)

Dependency structure (contd)



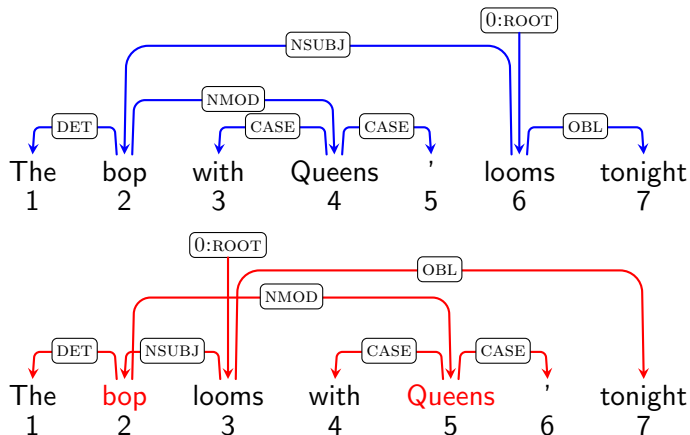
- But semantics: *John* is also the AGENT of *go* (control) – *not captured by dependency analysis* because no phrase structure. And this kind of relation is systematic.
- *He wants to sleep in class.*
- *He promises her not to sleep in class.*
- Also raising: *She seems to be reading a book.*

Projectivity



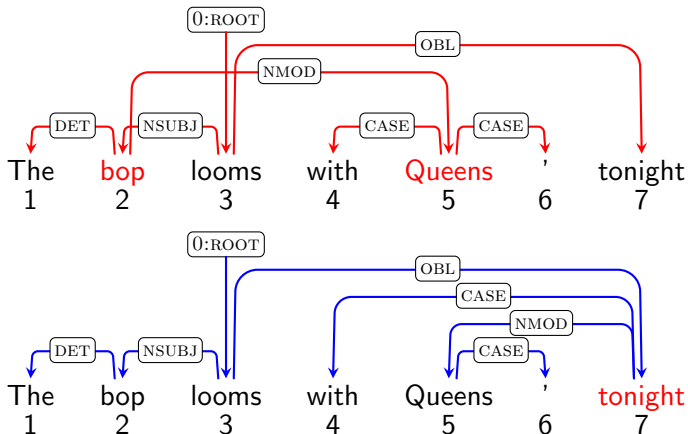
- Dependency trees can be projective or non-projective: tree is projective if every arc is projective
- Arc from head to dependent is projective if there is a path from the head to every token between the head and dependent in the string (can check this visually)
- OBL: oblique nominal ('tonight' treated as a temporal nominal; cf. 'later' an adverbial with ADVMOD)

Projectivity



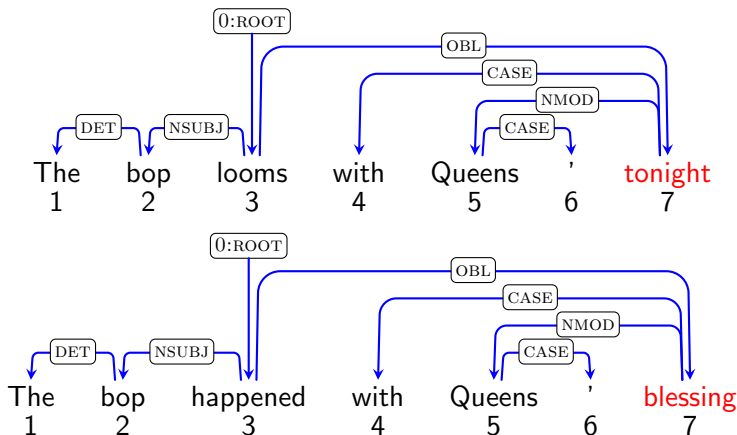
- Dependency trees can be projective or *non-projective*
- Arc from head to dependent is projective if there is a path from the head to every token between the head and dependent in the string

Projectivity



- Dependency trees can be projective or *non-projective*
- But in practice: most widely used dependency parsers produce projective trees only (for others see J&M 3rd edn §14.5)

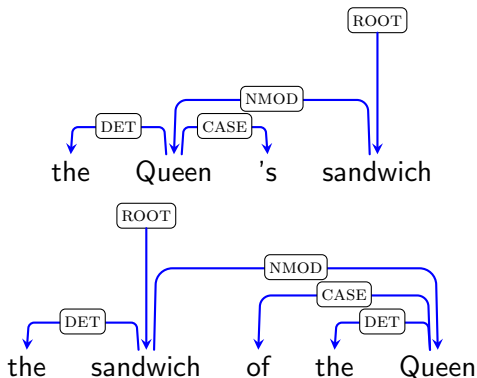
Projectivity



- Dependency trees can be projective or *non-projective*
- But in practice: parsers tend to be projective only (!linguistic alarm)



The CASE relation

- The CASE relation: prepositions, postpositions, clitic markers
- Parallel analyses for possessive alternation:



The CASE relation

- The CASE relation: prepositions, postpositions, clitic markers
- Cross-linguistic consistency:



Word order

- English word order: usually subject verb object (SVO)
- 'who did what to whom': *The dog bites the man* \neq *The man bites the dog* (word order matters)
- Passives are still SVO (morpho-syntactically, cf. semantics): *The man was bitten by the dog*
- Can 'front' the object in topicalization: *The man, the dog bites ... He's having a rough day*
- Other languages may be more flexible with word order

Word order

- e.g. German: *Der Hund beißt den Mann* = *Den Mann beißt der Hund*.
'The dog bites the man'
- e.g. Russian (example from Emily Bender, 2013):

Chelovek ukusil sobaku

man.NOM.SG.M bite.PAST.PFV.SG.M dog.ACC.SG.F

'the man bit the dog'

All word orders possible with same meaning (in different discourse contexts):

Chelovek ukusil sobaku

Chelovek sobaku ukusil

Ukusil chelovek sobaku

Ukusil sobaku chelovek

Sobaku chelovek ukusil

Sobaku ukusil chelovek

Word order

- Due to WO variability: rules like $S \rightarrow NP VP$ do not work in all languages (rule explosion)
- Could allow discontinuous constituency: e.g. **Sobaku** **chelovek** **ukusil** ('dog man bit') with a split VP
- But parsing discontinuities is not straightforward
- \rightarrow dependency structures
- Binary relations between words

(questions?) → Part 2: Parsing

Transition-based dependency parsing

A **transition system** for parsing is a quadruple $S = (C, T, c_s, C_t)$, where

- ① C is a set of **configurations**, each of which represents a **parser state**.
- ② T is a set of **transitions**, each of which represents a **parsing action**,
- ③ c_s **initializes** S by mapping a sentence x to a particular configuration,
- ④ $C_t \subseteq C$ is a set of **terminal** configurations.

Transition-based dependency parsing

A stack-based **configuration** (parser state) for a string $x = w_0, w_1, \dots, w_n$ (V) as $c = (\sigma, \beta, \mathcal{A})$, where

- 1 σ is a stack of tokens with which to build the parse;
- 2 β is a buffer of tokens yet to be parsed;
- 3 \mathcal{A} is a set of relations arcs representing a dependency tree

Transition-based dependency parsing

```
DEPPARSE( $x = (w_0, w_1, \dots, w_n)$ )  
1   $c \leftarrow \{[\text{ROOT}], [w_0, w_1, \dots, w_n], []\}$   
2  while  $c \notin C_t$   
3       $t \leftarrow \text{ORACLE}(c)$ ; choose transition operator  
4       $c \leftarrow \text{APPLY}(t, c)$ ; apply it, creating a new state  
5  return  $c$ 
```

- Initial configuration c_s : ROOT in stack, tokens of x in buffer, empty set of relations in A
- Actions based on current *configuration*: {stack, buffer, relations so far} and consulting an *oracle* for correct operation, which is applied to produce a new configuration
- \rightarrow End goal: buffer empty and only ROOT on stack, A as a set of relations representing the parse for input string x , such that $G = (V, A)$ is a dependency graph for x


Transition-based dependency parsing

- The **arc standard** approach (Covington 2001, Nivre 2003)
- Available transitions/actions: SHIFT, LEFT-ARC, RIGHT-ARC
- At each step: either SHIFT or link word tokens with LEFT-ARC, RIGHT-ARC (analogies with intuitive tree drawing)
- SHIFT: pop word token from front of buffer to top of stack (\approx postpone dealing with the current token, store it for later on)
- LEFT-ARC: draw head-dependent arc between token at the top of the stack and the second item on the stack (cannot be ROOT); remove second item from stack (\approx assign current token as head of a previously seen token)
- RIGHT-ARC: draw head-dependent arc from second token on the stack to the first one; remove top item from stack (\approx assign previously seen token as head of current token)
- Can parameterize the actions to produce labels: LEFT-ARC_{subj}, etc

Transition-based dependency parsing

- Metaphor from shift-reduce parsing (analysis of programming languages): left-arc and right-arc are REDUCE operations, combining elements on the stack
- *Greedy*: Oracle provides a single choice at each step, parser proceeds, no further exploration or opportunity to back-track, single parse at the end
- Complexity is linear to length of the sentence: a single left-to-right pass through x
- Parsing as a “sequence of transitions through the space of possible configurations” (J&M)
- (more complex algorithms in J&M 3rd edn ch.14, and L95)
- → walk-through...

Parsing example



Shift

Left-Arc_{nsubj}

Left-Arc_{advmod}

Left-Arc_{aux}

Left-Arc_{det}

Right-Arc_{obj}

Stack

[ROOT] [ROOT, John] [ROOT, John, is] [ROOT, is] [ROOT, is, carefully] [ROOT, is, carefully, checking] [ROOT, is, carefully, checking, the] [ROOT, is, carefully, checking, the, machine]

Buffer

[John, is, ..., the, machine] [is, carefully, checking, the, machine] [carefully, checking, the, machine] [carefully, checking, the, machine]

Transition-based dependency parsing

- How to learn to take the correct actions for each parse configuration?
- Feature-based classification: surrounding words, lemmas, PoS-tags, morph features, relations already assigned (and concatenations) – associated with the correct action to take
- e.g. s_0^w , s_1^w , b_0^w , s_0^l , s_1^l , s_0^t , b_0^t , s_0^m , b_0^m , r_0 , r_1 , etc
- In training: choose LEFT-ARC if it produces a correct head-dependent relation given the reference parse and current configuration
- Or choose RIGHT-ARC if it produces a correct head-dependent relation given the reference parse and dependents have already been assigned (*prefer* the *flight* through *Denver*)
- Else choose SHIFT
- Train with a classifier such as multinomial logistic regression, support vector machines, etc
- More recently: word (token) embeddings and neural networks (next lecture)

Training data

- Training data: hand-labelled treebanks (a corpus sub-type)
- Laborious and costly to put together (or, automatically derived from phrase-structure annotation, can be lossy)
- Best-known are the annotated *WSJ* texts in the Penn Treebank
- and *Universal Dependencies*: a long-running, multi-institution project
- Collecting and releasing treebanks for many languages (v2.9, gt. 100 at present)
- Concretely proposed in publications by Ryan McDonald et al (ACL 2013) then Marie-Catherine de Marneffe et al (LREC 2014); the ideas pre-date this, have evolved since
- Website with more info and open datasets: universaldependencies.org
- 5th Universal Dependencies Workshop 2021 (ACL Anthology)
- Often with a shared task: cross-lingual parsing

CoNLL-X treebank format

```
# sent_id = 1
# text = They buy and sell books.
1  They    they    PRON    PRP    Case=Nom|Number=Plur      2  nsubj    2:nsubj|4:nsubj    _
2  buy     buy     VERB    VBP    Number=Plur|Person=3|Tense=Pres 0  root      0:root            _
3  and     and     CONJ    CC      _                      4  cc        4:cc              _
4  sell    sell    VERB    VBP    Number=Plur|Person=3|Tense=Pres 2  conj      0:root|2:conj     _
5  books   book    NOUN    NNS     Number=Plur              2  obj       2:obj|4:obj       SpaceAfter=No
6  .       .       PUNCT    .       _                      2  punct     2:punct           _

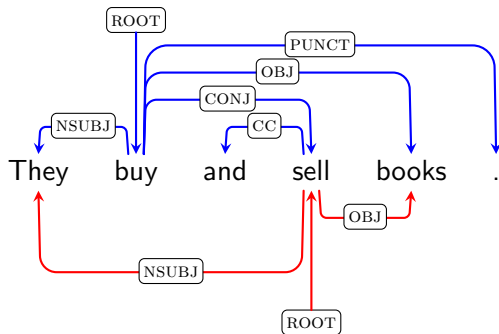

# sent_id = 2
# text = I have no clue.
1  I       I       PRON    PRP    Case=Nom|Number=Sing|Person=1    2  nsubj    _ _
2  have    have    VERB    VBP    Number=Sing|Person=1|Tense=Pres 0  root     _ _
3  no      no      DET     DT     PronType=Neg                    4  det      _ _
4  clue    clue    NOUN    NN     Number=Sing                   2  obj      _ SpaceAfter=No
5  .       .       PUNCT    .       _                      2  punct    _ _

# sent_id = panc0.s4
# text = तत् यथानुश्रूयते।
# translit = tat yathānuśrūyate.
# text_fr = Voilà ce qui nous est parvenu par la tradition orale.
# text_en = This is what is heard.
1  तत्      तद्      DET      _      Case=Nom|...|PronType=Dem    3  nsubj    _ Translit=tat|LTranslit=tad|Gloss=it
2-3 यथानुश्रूयते _      _      _      _      _      _      _ SpaceAfter=No
2  यथा     यथा     ADV      _      PronType=Rel                 3  advmod   _ Translit=yathā|LTranslit=yathā|Gloss=how
3  अनुश्रूयते अनु-श्रु VERB      _      Mood=Ind|...|Voice=Pass      0  root     _ Translit=anuśrūyate|LTranslit=anu-śru|Gloss=
4  |       |       PUNCT    _      _                          3  punct    _ Translit=.|LTranslit=.|Gloss=.
```

ID, form, lemma, UPOS, XPOS, morph, head, dep-rel, enhanced deps, misc.
and # commented info

Enhanced dependencies

- As well as 'basic' UDs: intended for shallow analyses and downstream NLP tasks
- 'Enhanced' UDs: for optional, deeper syntactic analyses



Enhanced dependencies

- For optional, deeper syntactic analyses
- To make implicit relations explicit (e.g. coordination example), and make finer-grained distinctions within large catch-all classes (e.g. nominal modifiers NMOD, oblique nominals OBL)
- Other examples include:
 - control/raising: from embedded verb to subj ('*Jo* wants to *buy* a book', '*Jo* seems to be *reading*'; !non-proj)
 - relative clauses: rel.pronouns attached to antecedents as well as the main predicate in the relative clause ('the *boy* ← *who* flew'; !cyclic)
 - case information: to disambiguate semantic roles ('they went *to* dinner *after* work', OBL:TO, ADVCL:AFTER)
- Enhanced UD's were a feature of the UD shared task this year: you could look at the published papers to see how participants adapted to them

More transition-based dependency parsing

- Can adapt or introduce other actions: SHIFT, LEFT-ARC, RIGHT-ARC
- REDUCE in arc-eager: pop the stack (left & right-arcs b/w stack and buffer)
- Can assign rightward relations sooner, rather than holding them on the stack, reduces late-assignment errors
- EDIT: to delete previously proposed arcs
- e.g. can deal with spoken disfluencies: false starts and repetition
- Training on the annotated treebank in the Switchboard Corpus of telephone conversations
- As with the Redshift non-monotonic (revise previous decisions), arc-eager (see J&M) parser by Honnibal & Johnson (CoNLL 2013, TACL 2014)
- Honnibal → spaCy: industry-ready dependency parsing
- Previously Redshift, and now neural

Transition-based dependency parsing in the wild

DEPENDENCY PARSING SYSTEM	UAS	LAS
spaCy RoBERTa (2020)	95.1	93.7
Mrini et al. (2019)	97.4	96.3
Zhou and Zhao (2019)	97.2	95.7

- Penn Treebank test set (*WSJ*)
- UAS: unlabelled attachment scores
- LAS: labelled attachment scores

Multilingual state-of-the-art dependency parsing

Model	LAS	MLAS	BLEX	Paper / Source
Stanford (Qi et al.)	74.16	62.08	65.28	Universal Dependency Parsing from Scratch
UDPipe Future (Straka)	73.11	61.25	64.49	UDPipe 2.0 Prototype at CoNLL 2018 UD Shared Task
HIT-SCIR (Che et al.)	75.84	59.78	65.33	Towards Better UD Parsing: Deep Contextualized Word Embeddings, Ensemble, and Treebank Concatenation

- from: NLP Progress <http://nlpprogress.com>
- CoNLL 2018 Shared Task: 82 test sets, 57 languages
- MLAS: morphology-aware labeled attachment score (incl PoS-tags & morph feats)
- BLEX: bi-lexical dependency score (dep.rels & lemmas)

Summary

- Pros include: argument relations to the fore, enables cross-lingual work (use of morph feats, flex word order, UDs), fast and practical
- Cons include: linguistically shallow, some semantic relations not captured, catch-all classes, parses of convenience (because projectivity constraint)
- Greedy algorithm can go wrong, but usually reasonable accuracy (cp human language processing?)
- No notion of grammaticality (so robust to typos)
- Fast (linear time) and practical: downstream applications such as question answering, coreference resolution, information extraction – approximate semantic relations between words via morpho-syntax
- What's your motivation? (NLP / CL)
- More on dependency parsing in Part IB/II Formal Models of Language & ACS L95

NLP lectures

- 1 Overview lecture
- 2 Morphology & FSTs
- 3 PoS-tagging & log-linear models
- 4 Phrase structure & structured prediction
- 5 Dependency parsing
(Now could be a good time to consolidate: morpho-syntactic foundations)
- 6 Neural Networks

Further reading

- Ann's lecture notes.
<https://www.cl.cam.ac.uk/teaching/2122/NLP/materials.html>
- Chapter 14 'Dependency Parsing' by Jurafsky & Martin *Speech and Language Processing* 3rd edition (in prep).
<https://web.stanford.edu/~jurafsky/slp3/14.pdf>
- de Marneffe et al, 'Universal Dependencies', in *Computational Linguistics* https://doi.org/10.1162/coli_a_00402
- Universal Dependencies events
<https://universaldependencies.org/events.html>
- IWPT 2021 Shared Task on Parsing into Enhanced Universal Dependencies <https://universaldependencies.org/iwpt21>
and Proceedings <https://aclanthology.org/volumes/2021.iwpt-1>